

# VISVESVARAYA TECHNOLOGICAL UNIVERSITY

“JnanaSangama”, Belgaum -590014, Karnataka.



## LAB REPORT on

## Big Data Analytics(23CS6PCBDA)

*Submitted by*

**Maria Sayeema(1BM22CS151)**

*in partial fulfillment for the award of the degree of*  
**BACHELOR OF ENGINEERING**  
*in*  
**COMPUTER SCIENCE AND ENGINEERING**



**B.M.S. COLLEGE OF ENGINEERING**  
(Autonomous Institution under VTU)  
**BENGALURU-560019**  
**Feb-2024 to July-2024**

**B. M. S. College of Engineering,**  
**Bull Temple Road, Bangalore 560019**  
(Affiliated To Visvesvaraya Technological University, Belgaum)  
**Department of Computer Science and Engineering**



**CERTIFICATE**

This is to certify that the Lab work entitled **Big Data Analytics** carried out by **Maria Sayeema(1BM22CS151)**, who is bonafide student of **B. M. S. College of Engineering**. It is in partial fulfillment for the award of **Bachelor of Engineering in Computer Science and Engineering** of the Visvesvaraya Technological University, Belgaum during the year 2024. The Lab report has been approved as it satisfies the academic requirements in respect of a **Course Big Data Analytics(23CS6PCBDA)** work prescribed for the said degree.

Prof.Amrutha B  
Assistant Professor  
Department of CSE  
BMSCE, Bengaluru

**Dr. Kavitha Sooda**  
Professor and Head  
Department of CSE  
BMSCE, Bengaluru

## Index Sheet

Sl. No.	Experiment Title	Page No.
1	MongoDB- CRUD Operations Demonstration (Practice and Self Study)	4
2	Neo4j	22
3	Perform the following DB operations using Cassandra. a) Create a keyspace by name Employee b) Create a column family by name Employee-Info with attributes Emp_Id Primary Key, Emp_Name, Designation, Date_of_Joining, Salary, Dept_Name c) Insert the values into the table in batch d) Update Employee name and Department of Emp-Id 121 e) Sort the details of Employee records based on salary f) Alter the schema of the table Employee_Info to add a column Projects which stores a set of Projects done by the corresponding Employee. g) Update the altered table to add project names. h) Create a TTL of 15 seconds to display the values of Employees.	27
4	Perform the following DB operations using Cassandra. a) Create a keyspace by name Library b) Create a column family by name Library-Info with attributes Stud_Id Primary Key, Counter_value of type Counter, Stud_Name, Book-Name, Book-Id, Date_of_issue c) Insert the values into the table in batch d) Display the details of the table created and increase the value of the counter e) Write a query to show that a student with id 112 has taken a book "BDA" 2 times. f) Export the created column to a csv file	30

	g) Import a given csv dataset from local file system into Cassandra column family	
5	Execution of HDFS Commands for interaction with Hadoop Environment. (Minimum 10 commands to be executed)	37
6	Implement Wordcount program on Hadoop framework	41
7	From the following link extract the weather data <a href="https://github.com/tomwhite/hadoop-book/tree/master/input/ncdc/all">https://github.com/tomwhite/hadoop-book/tree/master/input/ncdc/all</a> a) Create a MapReduce program to find average temperature for each year from NCDC data set. b) find the mean max temperature for every month.	48
8	Write a Scala program to print numbers from 1 to 100 using for loop.	55
9	Using RDD and FlatMap count how many times each word appears in a file and write out a list of words whose count is strictly greater than 4 using Spark.	57
10	Write a simple streaming program in Spark to receive text data streams on a particular port, perform basic text cleaning (like white space removal, stop words removal, lemmatization, etc.), and print the cleaned text on the screen. (Open Ended Question).	58

### Course Outcome

<b>CO1</b>	Apply the concept of NoSQL, Hadoop or Spark for a given task
<b>CO2</b>	Analyse big data analytics mechanisms that can be applied to obtain solution for a given problem.
<b>CO3</b>	Design and implement solutions using data analytics mechanisms for a given problem.

## Experiment - 1

### MongoDB- CRUD Operations Demonstration

```
Activities Terminal Mar 11 14:30
mongosh mongodb+srv://<credentials>@cluster0.tud3r.mongodb.net/

benecasp@benecasp-HP-Elite-Tower-880-G5-Desktop-PC: $ mongosh 'mongodb+srv://cluster0.tud3r.mongodb.net/' --apiVersion 1 --username mariacs22 --password 8Ja9u6zrEjzBlot
Current Mongosh Log ID: 67c7fab58790ab6f31c0e7d
Connecting to:
  mongodb+srv://<credentials>@cluster0.tud3r.mongodb.net/?appName=mongosh+2.2.1
Using MongoDB:
  8.0.5 (API Version 1)
Using Mongosh:
  2.2.1
mongosh 2.4.2 is available for download: https://www.mongodb.com/try/download/shell
For mongosh info see: https://docs.mongodb.com/mongosh-shell/

Atlas atlas-vvj19i-shard-0 [primary] test>
Atlas atlas-vvj19i-shard-0 [primary] test> use myDB
switched to db myDB
Atlas atlas-vvj19i-shard-0 [primary] myDB> create.collection("student1");
MongoServerError: create is not defined
Atlas atlas-vvj19i-shard-0 [primary] myDB> db.createCollection("Student1");
Uncaught:
SyntaxError: Unexpected token, expected ":", "}" (1:29)
> 1 | db.createCollection("Student1");
    |                               ^
    |
    |
Atlas atlas-vvj19i-shard-0 [primary] myDB> db.createCollection("Student1");
{ ok: 1 }
Atlas atlas-vvj19i-shard-0 [primary] myDB> db.Student1.insert({
...   _id: 1,
...   StudName: "MichelleJacinth",
...   Grade: "VII",
...   Hobbies: "InternetSurfing"
... });
DeprecationWarning: Collection.insert() is deprecated. Use insertOne, insertMany, or bulkWrite.
{ acknowledged: true, insertedIds: { '0': 1 } }
Atlas atlas-vvj19i-shard-0 [primary] myDB> db.Student1.update(
...   { _id: 1, StudName: "Ayanashele", Grade: "VII" },
...   { $set: { Hobbies: "Skating" } },
...   { upsert: true }
... );
DeprecationWarning: Collection.update() is deprecated. Use updateOne, updateMany, or bulkWrite.
{ acknowledged: true,
  insertedId: 1,
  matchedCount: 0,
  modifiedCount: 0,
  upsertedCount: 1 }
Atlas atlas-vvj19i-shard-0 [primary] myDB> db.Student1.find();
[ { _id: 1,
  StudName: "MichelleJacinth",
  Grade: "VII",
  Hobbies: "InternetSurfing" },
  { _id: 1, Grade: "VII", StudName: "Ayanashele", Hobbies: "Skating" } ]
```

```
Activities Terminal Mar 11 14:31 mongosh mongodb+srv://<credentials>@cluster0.tud3r.mongodb.net/

{ "_id": 3, Grade: 'VII', StudName: 'AryanDavid', Hobbies: 'Skating' }
}
Atlas atlas-vvj19l-shard-0 [primary] myDB> db.Student1.update(
... { '_id': 3, StudName: 'AryanDavid', Grade: 'VII' },
... { $set: { Hobbies: 'Chess' } }
... );
{
  acknowledged: true,
  insertedId: null,
  matchedCount: 1,
  modifiedCount: 1,
  upsertedCount: 0
}
Atlas atlas-vvj19l-shard-0 [primary] myDB> ds.Student1.find();
ReferenceError: ds is not defined
Atlas atlas-vvj19l-shard-0 [primary] myDB> db.Student1.find();
[
  {
    "_id": 1,
    StudName: 'MichelleJacinth',
    Grade: 'VII',
    Hobbies: 'InternetSurfing'
  },
  {
    "_id": 3, Grade: 'VII', StudName: 'AryanDavid', Hobbies: 'Chess' }
]
Atlas atlas-vvj19l-shard-0 [primary] myDB> db.Student1.find({StudName: 'Aryan David'});
Atlas atlas-vvj19l-shard-0 [primary] myDB> db.Student1.find({StudName: 'Aryan David'});
Atlas atlas-vvj19l-shard-0 [primary] myDB> db.Student1.find({StudName: 'AryanDavid'});
[ { "_id": 3, Grade: 'VII', StudName: 'AryanDavid', Hobbies: 'Chess' } ]
Atlas atlas-vvj19l-shard-0 [primary] myDB> db.Student1.find({}, {StudName: 1, GrGrade: 1, _id: 0});
[
  { StudName: 'MichelleJacinth', Grade: 'VII' },
  { Grade: 'VII', StudName: 'AryanDavid' }
]
Atlas atlas-vvj19l-shard-0 [primary] myDB> db.Student1.find({}, {StudName: 1, Grade: 1, _id: 0});
[
  { StudName: 'MichelleJacinth', Grade: 'VII' },
  { Grade: 'VII', StudName: 'AryanDavid' }
]
Atlas atlas-vvj19l-shard-0 [primary] myDB> db.Student1.find({Grade: { $eq: 'VII' }}).pretty();
[
  {
    "_id": 1,
    StudName: 'MichelleJacinth',
    Grade: 'VII',
    Hobbies: 'InternetSurfing'
  },
  {
    "_id": 3, Grade: 'VII', StudName: 'AryanDavid', Hobbies: 'Chess' }
]
Atlas atlas-vvj19l-shard-0 [primary] myDB> db.Student1.find({Hobbies: { $in: ['Chess', 'Skating'] }}).pretty();
[ { "_id": 3, Grade: 'VII', StudName: 'AryanDavid', Hobbies: 'Chess' } ]
Atlas atlas-vvj19l-shard-0 [primary] myDB> db.Student1.find({StudName: /A/}).pretty();
[
  {
    "_id": 1,
    StudName: 'MichelleJacinth',
    Grade: 'VII',
    Hobbies: 'InternetSurfing'
  },
  {
    "_id": 3, Grade: 'VII', StudName: 'AryanDavid', Hobbies: 'Chess' }
]
Atlas atlas-vvj19l-shard-0 [primary] myDB> db.Student1.find({Hobbies: { $in: ['Chess', 'Skating'] }}).pretty();
[ { "_id": 3, Grade: 'VII', StudName: 'AryanDavid', Hobbies: 'Chess' } ]
Atlas atlas-vvj19l-shard-0 [primary] myDB> db.Student1.find({StudName: /A/}).pretty();
[
  {
    "_id": 1,
    StudName: 'MichelleJacinth',
    Grade: 'VII',
    Hobbies: 'InternetSurfing'
  },
  {
    "_id": 3, Grade: 'VII', StudName: 'AryanDavid', Hobbies: 'Chess' }
]
Atlas atlas-vvj19l-shard-0 [primary] myDB> db.Student1.count();
DeprecationWarning: Collection.count() is deprecated. Use countDocuments or estimatedDocumentCount.
Atlas atlas-vvj19l-shard-0 [primary] myDB> db.Student1.find().sort({StudName: -1}).pretty();
[
  {
    "_id": 1,
    StudName: 'MichelleJacinth',
    Grade: 'VII',
    Hobbies: 'InternetSurfing'
  },
  {
    "_id": 3, Grade: 'VII', StudName: 'AryanDavid', Hobbies: 'Chess' }
]
Atlas atlas-vvj19l-shard-0 [primary] myDB> █
```

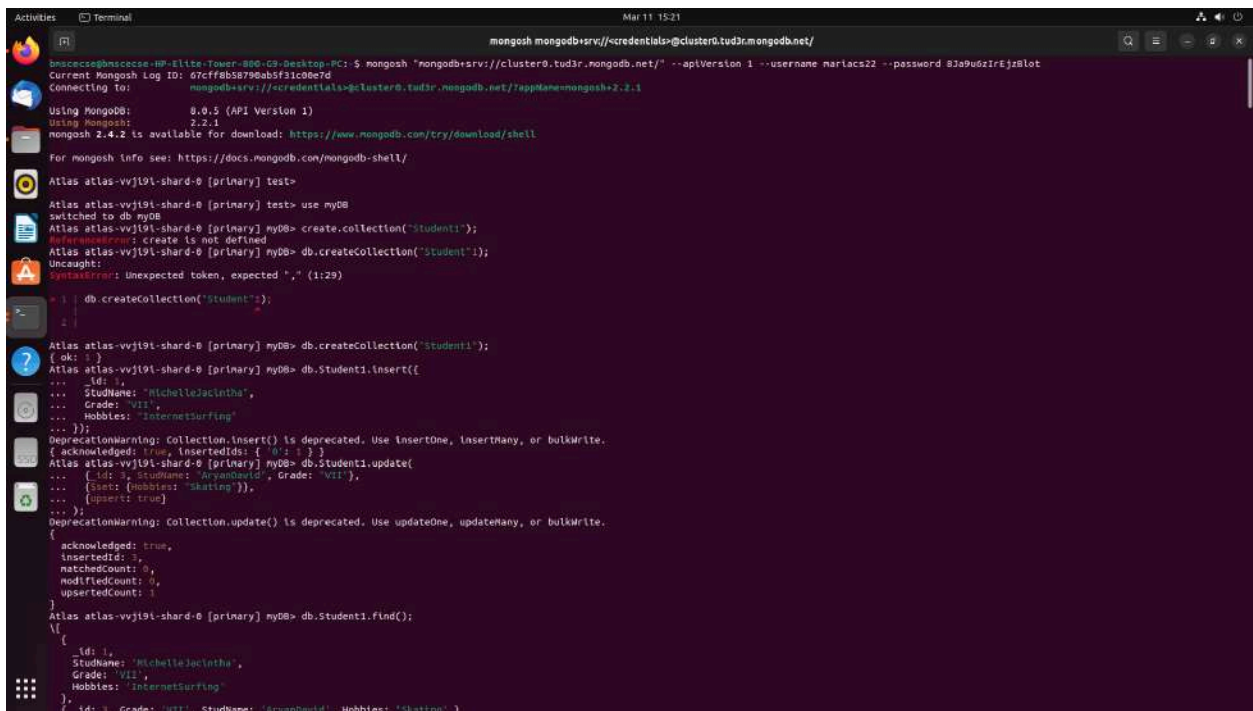
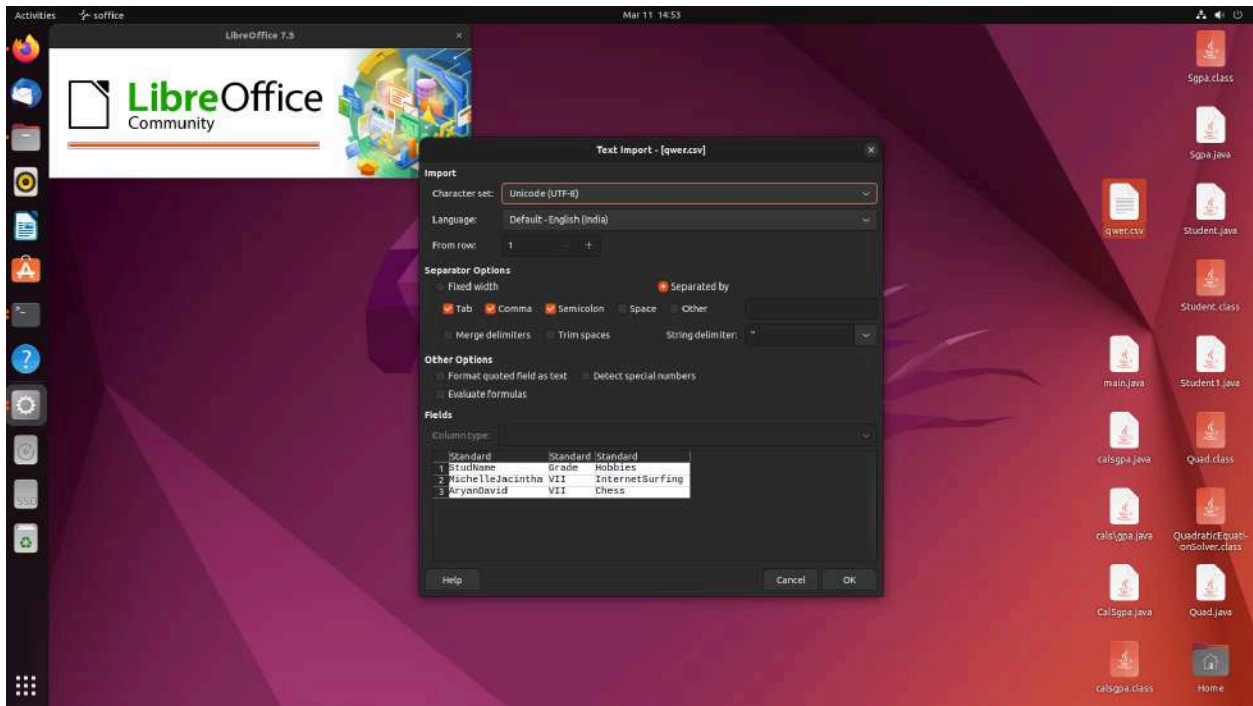
```
Activities Terminal Mar 11 14:31 mongosh mongodb+srv://<credentials>@cluster0.tud3r.mongodb.net/

{ "_id": 3, "Grade": "VII", "StudName": "AryanDavid", "Hobbies": "Skating" }
}
Atlas atlas-vvj19l-shard-0 [primary] myDB> db.Student1.update(
... { "_id": 3, "StudName": "AryanDavid", "Grade": "VII" },
... { $set: {Hobbies: "Chess"} }
... );
{
  acknowledged: true,
  insertedId: null,
  matchedCount: 1,
  modifiedCount: 1,
  upsertedCount: 0
}
Atlas atlas-vvj19l-shard-0 [primary] myDB> ds.Student1.find();
ReferenceError: ds is not defined
Atlas atlas-vvj19l-shard-0 [primary] myDB> db.Student1.find();
[
  {
    "_id": 1,
    "StudName": "MichelleJacinth",
    "Grade": "VII",
    "Hobbies": "InternetSurfing"
  },
  {
    "_id": 3, "Grade": "VII", "StudName": "AryanDavid", "Hobbies": "Chess" }
]
Atlas atlas-vvj19l-shard-0 [primary] myDB> db.Student1.find({StudName: "Aryan David"});
Atlas atlas-vvj19l-shard-0 [primary] myDB> db.Student1.find({StudName: "Aryan David"});
Atlas atlas-vvj19l-shard-0 [primary] myDB> db.Student1.find({StudName: "AryanDavid"});
[ { "_id": 3, "Grade": "VII", "StudName": "AryanDavid", "Hobbies": "Chess" } ]
Atlas atlas-vvj19l-shard-0 [primary] myDB> db.Student1.find({}, {StudName: 1, GrGrade: 1, _id: 0});
[
  { "StudName": "MichelleJacinth", "Grade": "VII" },
  { "Grade": "VII", "StudName": "AryanDavid" }
]
Atlas atlas-vvj19l-shard-0 [primary] myDB> db.Student1.find({}, {StudName: 1, Grade: 1, _id: 0});
[
  { "StudName": "MichelleJacinth", "Grade": "VII" },
  { "Grade": "VII", "StudName": "AryanDavid" }
]
Atlas atlas-vvj19l-shard-0 [primary] myDB> db.Student1.find({Grade: {Seq: "VII"}}).pretty();
[
  {
    "_id": 1,
    "StudName": "MichelleJacinth",
    "Grade": "VII",
    "Hobbies": "InternetSurfing"
  },
  {
    "_id": 3, "Grade": "VII", "StudName": "AryanDavid", "Hobbies": "Chess" }
]
Atlas atlas-vvj19l-shard-0 [primary] myDB> db.Student1.find({Hobbies: { $in: ["Chess", "Skating"] }}).pretty();
[ { "_id": 3, "Grade": "VII", "StudName": "AryanDavid", "Hobbies": "Chess" } ]
Atlas atlas-vvj19l-shard-0 [primary] myDB> db.Student1.find({StudName: /A/}).pretty();
[
  {
    "_id": 1,
    "StudName": "MichelleJacinth",
    "Grade": "VII",
    "Hobbies": "InternetSurfing"
  },
  {
    "_id": 3, "Grade": "VII", "StudName": "AryanDavid", "Hobbies": "Chess" }
]
Atlas atlas-vvj19l-shard-0 [primary] myDB> db.Student1.find({Hobbies: { $in: ["Chess", "Skating"] }}).pretty();
[ { "_id": 3, "Grade": "VII", "StudName": "AryanDavid", "Hobbies": "Chess" } ]
Atlas atlas-vvj19l-shard-0 [primary] myDB> db.Student1.find({StudName: /A/}).pretty();
[
  {
    "_id": 1,
    "StudName": "MichelleJacinth",
    "Grade": "VII",
    "Hobbies": "InternetSurfing"
  },
  {
    "_id": 3, "Grade": "VII", "StudName": "AryanDavid", "Hobbies": "Chess" }
]
Atlas atlas-vvj19l-shard-0 [primary] myDB> db.Student1.count();
DeprecationWarning: Collection.count() is deprecated. Use countDocuments or estimatedDocumentCount.
Atlas atlas-vvj19l-shard-0 [primary] myDB> db.Student1.find().sort({StudName: -1}).pretty();
[
  {
    "_id": 1,
    "StudName": "MichelleJacinth",
    "Grade": "VII",
    "Hobbies": "InternetSurfing"
  },
  {
    "_id": 3, "Grade": "VII", "StudName": "AryanDavid", "Hobbies": "Chess" }
]
Atlas atlas-vvj19l-shard-0 [primary] myDB> █
```









```
Activities Terminal Mar 11 15:22 mongosh mongodb+srv://<credentials>@cluster0.tud3r.mongodb.net/

{ "_id": 3, Grade: 'VII', StudName: 'AryanDavid', Hobbies: 'Skating' }

Atlas atlas-vvj19l-shard-0 [primary] myDB> db.Students.update(
... { _id: 3, StudName: 'AryanDavid', Grade: 'VII',
... { $set: { Hobbies: 'Chess' } }
... });
{
  acknowledged: true,
  insertedId: null,
  matchedCount: 1,
  modifiedCount: 1,
  upsertedCount: 0
}
Atlas atlas-vvj19l-shard-0 [primary] myDB> db.Students.find();
Warning: ds is not defined
Atlas atlas-vvj19l-shard-0 [primary] myDB> db.Students.find();
{
  "_id": 1,
  "StudName": "MichelleJacinth",
  "Grade": "VII",
  "Hobbies": "InternetSurfing"
},
{ "_id": 3, Grade: 'VII', StudName: 'AryanDavid', Hobbies: 'Chess' }
}
Atlas atlas-vvj19l-shard-0 [primary] myDB> db.Students.find({StudName: 'Aryan David'});
Atlas atlas-vvj19l-shard-0 [primary] myDB> db.Students.find({StudName: 'Aryan David'});
Atlas atlas-vvj19l-shard-0 [primary] myDB> db.Students.find({StudName: 'AryanDavid'});
{ { "_id": 3, Grade: 'VII', StudName: 'AryanDavid', Hobbies: 'Chess' } }
Atlas atlas-vvj19l-shard-0 [primary] myDB> db.Students.find({StudName: 1, Grade: 1, _id: 0});
{
  { StudName: 'MichelleJacinth', Grade: 'VII' },
  { Grade: 'VII', StudName: 'AryanDavid' }
}
Atlas atlas-vvj19l-shard-0 [primary] myDB> db.Students.find({StudName: 1, Grade: 1, _id: 0});
{
  { StudName: 'MichelleJacinth', Grade: 'VII' },
  { Grade: 'VII', StudName: 'AryanDavid' }
}
Atlas atlas-vvj19l-shard-0 [primary] myDB> db.Students.find({Grade: { $eq: 'VII' }}).pretty();
{
  "_id": 1,
  "StudName": "MichelleJacinth",
  "Grade": "VII",
  "Hobbies": "InternetSurfing"
},
{ "_id": 3, Grade: 'VII', StudName: 'AryanDavid', Hobbies: 'Chess' }
}
Atlas atlas-vvj19l-shard-0 [primary] myDB> db.Students.find({Hobbies: { $in: ['Chess', 'Skating'] }}).pretty();
{ { "_id": 3, Grade: 'VII', StudName: 'AryanDavid', Hobbies: 'Chess' } }
Atlas atlas-vvj19l-shard-0 [primary] myDB> db.Students.find({StudName: /A/}).pretty();
{
  {
    "_id": 1,
    "StudName": "MichelleJacinth",
    "Grade": "VII",
    "Hobbies": "InternetSurfing"
  },
  {
    "_id": 3,
    "Grade": "VII",
    "StudName": "AryanDavid",
    "Hobbies": "Chess"
  }
}
}

Atlas atlas-vvj19l-shard-0 [primary] myDB> db.Students.find({StudName: /A/}).pretty();
{
  {
    "_id": 1,
    "StudName": "MichelleJacinth",
    "Grade": "VII",
    "Hobbies": "InternetSurfing"
  },
  {
    "_id": 3,
    "Grade": "VII",
    "StudName": "AryanDavid",
    "Hobbies": "Chess"
  }
}
}

Atlas atlas-vvj19l-shard-0 [primary] myDB> db.Students.count();
DeprecationWarning: Collection.count() is deprecated. Use countDocuments or estimatedDocumentCount.
Atlas atlas-vvj19l-shard-0 [primary] myDB> db.Students.find().sort({StudName: -1}).pretty();
{
  {
    "_id": 1,
    "StudName": "MichelleJacinth",
    "Grade": "VII",
    "Hobbies": "InternetSurfing"
  },
  {
    "_id": 3,
    "Grade": "VII",
    "StudName": "AryanDavid",
    "Hobbies": "Chess"
  }
}
}
Atlas atlas-vvj19l-shard-0 [primary] myDB>
(To exit, press Ctrl+C again or Ctrl+D or type .exit)
Atlas atlas-vvj19l-shard-0 [primary] myDB> db.Students.find().pretty();
{
  {
    "_id": 1,
    "StudName": "MichelleJacinth",
    "Grade": "VII",
    "Hobbies": "InternetSurfing"
  },
  {
    "_id": 3,
    "Grade": "VII",
    "StudName": "AryanDavid",
    "Hobbies": "Chess"
  }
}
}
Atlas atlas-vvj19l-shard-0 [primary] myDB> db.Students.find()
{
  {
    "_id": ObjectId('67d0009c2043200f8c4ffc'),
    "StudName": "AryanDavid",
    "Grade": "VII",
    "Hobbies": "Chess"
  },
  {
    "_id": ObjectId('67d0009c2043200f8c4ffc'),
    "StudName": "MichelleJacinth",
    "Grade": "VII",
    "Hobbies": "InternetSurfing"
  }
}
```

```
Activities Terminal Mar 11 15:21
mongosh mongodb+srv://<credentials>@cluster0.tud3r.mongodb.net/

Unauthenticated session - HP-Elite-Tower-890-G9-Desktop-PC: $ mongosh 'mongodb+srv://<credentials>@cluster0.tud3r.mongodb.net/' --apiVersion 1 --username mariacs22 --password 8Ja9u6zrEjz8lot
Current Mongosh Log ID: 67c7fab5879ab5f31c0be7d
Connecting to:
  mongodb+srv://<credentials>@cluster0.tud3r.mongodb.net/?appName=mongosh+2.2.1

Using MongoDB:
  6.0.5 (API Version 1)
Using Mongosh:
  2.2.1
mongosh 2.4.2 is available for download: https://www.mongodb.com/try/download/shell
For mongosh info see: https://docs.mongodb.com/mongodb-shell/

Atlas atlas-vvj9l-shard-0 [primary] test>
Atlas atlas-vvj9l-shard-0 [primary] test> use myDB
switched to db myDB
Atlas atlas-vvj9l-shard-0 [primary] myDB> create.collection("Student1");
ReferenceError: create is not defined
Atlas atlas-vvj9l-shard-0 [primary] myDB> db.createCollection("Student1");
Uncaught:
  SyntaxError: Unexpected token, expected ',' (1:29)
    > | db.createCollection("Student1");
    > |
    > |
Atlas atlas-vvj9l-shard-0 [primary] myDB> db.createCollection("Student1");
{ ok: 1 }
Atlas atlas-vvj9l-shard-0 [primary] myDB> db.Student1.insert({
  ...
  _id: 1,
  ...
  StudName: 'MichelleJacintha',
  ...
  Grade: 'VII',
  ...
  Hobbies: 'InternetSurfing'
  ...
});
DeprecationWarning: Collection.insert() is deprecated. Use insertOne, insertMany, or bulkWrite.
{ acknowledged: true, insertedIds: { '0': 1 } }
Atlas atlas-vvj9l-shard-0 [primary] myDB> db.Student1.update(
  ...
  { _id: 1, StudName: 'AryanDavid', Grade: 'VII' },
  ...
  { $set: { Hobbies: 'Skating' } },
  ...
  { upsert: true }
  ...
);
DeprecationWarning: Collection.update() is deprecated. Use updateOne, updateMany, or bulkWrite.
{ acknowledged: true,
  insertedId: 1,
  matchedCount: 0,
  modifiedCount: 0,
  upsertedCount: 1 }
Atlas atlas-vvj9l-shard-0 [primary] myDB> db.Student1.find();
[
  {
    _id: 1,
    StudName: 'MichelleJacintha',
    Grade: 'VII',
    Hobbies: 'InternetSurfing'
  },
  {
    _id: 2, Grade: 'VII', StudName: 'AryanDavid', Hobbies: 'Skating'
  }
]

Atlas atlas-vvj9l-shard-0 [primary] myDB> db.Student1.update(
  ...
  { _id: 1, StudName: 'AryanDavid', Grade: 'VII' },
  ...
  { $set: { Hobbies: 'Chess' } }
  ...
);
{ acknowledged: true,
  insertedId: null,
  matchedCount: 1,
  modifiedCount: 1,
  upsertedCount: 0 }
Atlas atlas-vvj9l-shard-0 [primary] myDB> ds.Student1.find();
ReferenceError: ds is not defined
Atlas atlas-vvj9l-shard-0 [primary] myDB> db.Student1.find();
[
  {
    _id: 1,
    StudName: 'MichelleJacintha',
    Grade: 'VII',
    Hobbies: 'InternetSurfing'
  },
  {
    _id: 2, Grade: 'VII', StudName: 'AryanDavid', Hobbies: 'Chess'
  }
]
Atlas atlas-vvj9l-shard-0 [primary] myDB> db.Student.find({StudName: 'Aryan David'});
Atlas atlas-vvj9l-shard-0 [primary] myDB> db.Student.find({StudName: 'Aryan David'});
Atlas atlas-vvj9l-shard-0 [primary] myDB> db.Student1.find({StudName: 'AryanDavid'});
[ { _id: 1, Grade: 'VII', StudName: 'AryanDavid', Hobbies: 'Chess' } ]
Atlas atlas-vvj9l-shard-0 [primary] myDB> db.Student1.find({}, {StudName: 1, Grade: 1, _id: 0});
[ { StudName: 'MichelleJacintha', Grade: 'VII' },
  { Grade: 'VII', StudName: 'AryanDavid' } ]
Atlas atlas-vvj9l-shard-0 [primary] myDB> db.Student1.find({}, {StudName: 1, Grade: 1, _id: 0});
[ { StudName: 'MichelleJacintha', Grade: 'VII' },
  { Grade: 'VII', StudName: 'AryanDavid' } ]
Atlas atlas-vvj9l-shard-0 [primary] myDB> db.Student1.find({Grade: { $eq: 'VII' }}).pretty();
[
  {
    _id: 1,
    StudName: 'MichelleJacintha',
    Grade: 'VII',
    Hobbies: 'InternetSurfing'
  },
  {
    _id: 2, Grade: 'VII', StudName: 'AryanDavid', Hobbies: 'Chess'
  }
]
Atlas atlas-vvj9l-shard-0 [primary] myDB> db.Student1.find({Hobbies: { $in: [ 'Chess', 'Skating' ] }}).pretty();
[ { _id: 1, Grade: 'VII', StudName: 'AryanDavid', Hobbies: 'Chess' } ]
Atlas atlas-vvj9l-shard-0 [primary] myDB> db.Student1.find({StudName: /A/}).pretty();
[
  {
    _id: 1,
    StudName: 'MichelleJacintha',
    Grade: 'VII',
    Hobbies: 'InternetSurfing'
  },
  {
    _id: 2, Grade: 'VII', StudName: 'AryanDavid', Hobbies: 'Chess'
  }
]
```

```
Activities Terminal Mar 11 15:22 mongosh mongodb+srv://<credentials>@cluster0.tud3r.mongodb.net/

Atlas atlas-vvj19l-shard-0 [primary] myDB> db.Student1.find({StudName: 'Michelle'}).pretty();
{
  "_id": 1,
  "StudName": "MichelleJacinth",
  "Grade": "VII",
  "Hobbies": "InternetSurfing"
}

Atlas atlas-vvj19l-shard-0 [primary] myDB> db.Student1.find({StudName: 'e'}).pretty();
{
  "_id": 1,
  "StudName": "MichelleJacinth",
  "Grade": "VII",
  "Hobbies": "InternetSurfing"
}

Atlas atlas-vvj19l-shard-0 [primary] myDB> db.Student1.count();
DeprecationWarning: Collection.count() is deprecated. Use countDocuments or estimatedDocumentCount.

Atlas atlas-vvj19l-shard-0 [primary] myDB> db.Student1.find().sort({StudName: -1}).pretty();
{
  "_id": 1,
  "StudName": "MichelleJacinth",
  "Grade": "VII",
  "Hobbies": "InternetSurfing"
},
{ "_id": 3, "Grade": "VII", "StudName": "AryanDavid", "Hobbies": "Chess" }
}

Atlas atlas-vvj19l-shard-0 [primary] myDB>
(TypeError: db.Student1.save is not a function)
Atlas atlas-vvj19l-shard-0 [primary] myDB> db.Student1.find().pretty();
{
  "_id": 1,
  "StudName": "MichelleJacinth",
  "Grade": "VII",
  "Hobbies": "InternetSurfing"
},
{ "_id": 3, "Grade": "VII", "StudName": "AryanDavid", "Hobbies": "Chess" }
}

Atlas atlas-vvj19l-shard-0 [primary] myDB> db.Student3.find()
{
  "_id": ObjectId('67d0809c2043200f8c4fccc'),
  "StudName": "AryanDavid",
  "Grade": "VII",
  "Hobbies": "Chess"
},
{
  "_id": ObjectId('67d0809c2043200f8c4fcd'),
  "StudName": "MichelleJacinth",
  "Grade": "VII",
  "Hobbies": "InternetSurfing"
}
}

Hobbies: "InternetSurfing"
}

Atlas atlas-vvj19l-shard-0 [primary] myDB> db.Student1.save({StudName: "Vansh", Grade: "VI"})
TypeError: db.Student1.save is not a function
Atlas atlas-vvj19l-shard-0 [primary] myDB> use myDB;
already on db myDB
Atlas atlas-vvj19l-shard-0 [primary] myDB> db.Student1.insertOne({ StudName: "Vansh", Grade: "VI" });
{
  acknowledged: true,
  insertedId: ObjectId('67d0831a790ab5f31c0b0e7e')
}

Atlas atlas-vvj19l-shard-0 [primary] myDB> db.Student1.find()
{
  "_id": 1,
  "StudName": "MichelleJacinth",
  "Grade": "VII",
  "Hobbies": "InternetSurfing"
},
{ "_id": 3, "Grade": "VII", "StudName": "AryanDavid", "Hobbies": "Chess" },
{
  "_id": ObjectId('67d0831a790ab5f31c0b0e7e'),
  "StudName": "Vansh",
  "Grade": "VI"
}
}

Atlas atlas-vvj19l-shard-0 [primary] myDB> db.Student1.insertOne({
... _id: 4,
... StudName: "JohnDoe",
... Grade: "VIII",
... Hobbies: "Reading",
... Location: "abc"
... });
{ acknowledged: true, insertedId: 4 }

Atlas atlas-vvj19l-shard-0 [primary] myDB> db.Student1.find()
{
  "_id": 1,
  "StudName": "MichelleJacinth",
  "Grade": "VII",
  "Hobbies": "InternetSurfing"
},
{ "_id": 3, "Grade": "VII", "StudName": "AryanDavid", "Hobbies": "Chess" },
{
  "_id": ObjectId('67d0831a790ab5f31c0b0e7e'),
  "StudName": "Vansh",
  "Grade": "VI"
},
{
  "_id": 4,
  "StudName": "JohnDoe",
  "Grade": "VIII",
  "Hobbies": "Reading",
  "Location": "abc"
}
}
```

```
Activities Terminal Mar 11 15:23 mongosh mongodb+srv://<credentials>@cluster0.tud3r.mongodb.net/

Location: 'abc'
}
Atlas atlas-vvj19l-shard-0 [primary] myDB> db.Student1.updateOne(
... { _id: 4 }, // Find the document with _id: 4
... { $set: { Location: 'Network' } } // Set the Location field to "Network"
... );
{
  acknowledged: true,
  insertedId: null,
  matchedCount: 1,
  modifiedCount: 1,
  upsertedCount: 0
}
Atlas atlas-vvj19l-shard-0 [primary] myDB> db.Student1.find()
[
  {
    _id: 1,
    StudName: 'MichelleJacintho',
    Grade: 'VII',
    Hobbies: 'InternetSurfing'
  },
  {
    _id: 3, StudName: 'AryanDavid', Grade: 'VII', Hobbies: 'Chess'
  }
]
Atlas atlas-vvj19l-shard-0 [primary] myDB> db.Student1.find()
[
  {
    _id: 1,
    StudName: 'MichelleJacintho',
    Grade: 'VII',
    Hobbies: 'InternetSurfing'
  },
  {
    _id: 3, Grade: 'VII', StudName: 'AryanDavid', Hobbies: 'Chess'
  },
  {
    _id: ObjectId('67d8b31ab790ab5f31c0e7e'),
    StudName: 'Vansh',
    Grade: 'VI'
  },
  {
    _id: 4,
    StudName: 'JohnDoe',
    Grade: 'VIII',
    Hobbies: 'Reading',
    Location: 'Network'
  }
]
Atlas atlas-vvj19l-shard-0 [primary] myDB> db.Student1.find([{$or: 1}, {StudName: 1, Grade: 1, _id: 0}]);
[ { StudName: 'MichelleJacintho', Grade: 'VII' } ]
Atlas atlas-vvj19l-shard-0 [primary] myDB> db.Student1.find({Grade: {$in: 'VII'}});
[
  {
    _id: ObjectId('67d8b31ab790ab5f31c0e7e'),
    StudName: 'Vansh',
    Grade: 'VI'
  },
  {
    _id: 3,
    StudName: 'AryanDavid',
    Grade: 'VII',
    Hobbies: 'Chess'
  }
]
Atlas atlas-vvj19l-shard-0 [primary] myDB> db.Student1.find({StudName: 'Vansh'});
[
  {
    _id: 4,
    StudName: 'JohnDoe',
    Grade: 'VIII',
    Hobbies: 'Reading',
    Location: 'Network'
  }
]
Atlas atlas-vvj19l-shard-0 [primary] myDB> db.Student1.updateOne(
... { _id: 3 }, // Find the document with _id: 3
... { $set: { Location: 'Network' } } // Set the Location field to "Network"
... );
{
  acknowledged: true,
  insertedId: null,
  matchedCount: 1,
  modifiedCount: 1,
  upsertedCount: 0
}
Atlas atlas-vvj19l-shard-0 [primary] myDB> db.Student1.find()
[
  {
    _id: 1,
    StudName: 'MichelleJacintho',
    Grade: 'VII',
    Hobbies: 'InternetSurfing'
  },
  {
    _id: 3,
    StudName: 'AryanDavid',
    Grade: 'VII',
    Hobbies: 'Chess',
    Location: 'Network'
  },
  {
    _id: ObjectId('67d8b31ab790ab5f31c0e7e'),
    StudName: 'Vansh',
    Grade: 'VI'
  },
  {
    _id: 4,
    StudName: 'JohnDoe',
    Grade: 'VIII',
    Hobbies: 'Reading',
    Location: 'Network'
  }
]
Atlas atlas-vvj19l-shard-0 [primary] myDB> db.Student1.updateOne({ _id: 3 }, { $set: { Location: null } });
{
  acknowledged: true,
  insertedId: null,
  matchedCount: 1,
  modifiedCount: 1,
  upsertedCount: 0
}
```



```
Activities Terminal Mar 11 15:24 mongosh mongodb+srv://<credentials>@cluster0.tud3r.mongodb.net/

{
  "upsertedCount": 0
}
Atlas atlas-vvj19l-shard-0 [primary] myDB> db.Student1.find()
[
  {
    "_id": 1,
    "StudName": "MichelleJacinth",
    "Grade": "VII",
    "Hobbies": "InternetSurfing"
  },
  {
    "_id": 3,
    "Grade": "VII",
    "StudName": "AryanDavid",
    "Hobbies": "Chess",
    "Location": null
  },
  {
    "_id": ObjectId("67d0031ab790ab5f31c00e7e"),
    "StudName": "Vansh",
    "Grade": "VI"
  },
  {
    "_id": 4,
    "StudName": "JohnDoe",
    "Grade": "VIII",
    "Hobbies": "Reading",
    "Location": "Network"
  }
]
Atlas atlas-vvj19l-shard-0 [primary] myDB> db.Student1.count();
4
Atlas atlas-vvj19l-shard-0 [primary] myDB> db.Student1.countDocuments({Grade: "VII"});
2
Atlas atlas-vvj19l-shard-0 [primary] myDB> db.Student1.find({Grade: "VII"}).limit(2);
[
  {
    "_id": 1,
    "StudName": "MichelleJacinth",
    "Grade": "VII",
    "Hobbies": "InternetSurfing"
  },
  {
    "_id": 3,
    "Grade": "VII",
    "StudName": "AryanDavid",
    "Hobbies": "Chess",
    "Location": null
  }
]
Atlas atlas-vvj19l-shard-0 [primary] myDB> db.Student1.find().sort({StudName: 1});
[
  {
    "_id": 3,
    "Grade": "VII",
    "StudName": "AryanDavid",
    "Hobbies": "Chess",
    "Location": null
  },
  {
    "_id": 1,
    "StudName": "MichelleJacinth",
    "Grade": "VII",
    "Hobbies": "InternetSurfing"
  },
  {
    "_id": 4,
    "StudName": "JohnDoe",
    "Grade": "VIII",
    "Hobbies": "Reading",
    "Location": "Network"
  },
  {
    "_id": ObjectId("67d0031ab790ab5f31c00e7e"),
    "StudName": "Vansh",
    "Grade": "VI"
  }
]
Atlas atlas-vvj19l-shard-0 [primary] myDB> db.Student1.find().skip(2);
[
  {
    "_id": ObjectId("67d0031ab790ab5f31c00e7e"),
    "StudName": "Vansh",
    "Grade": "VI"
  },
  {
    "_id": 4,
    "StudName": "JohnDoe",
    "Grade": "VIII",
    "Hobbies": "Reading",
    "Location": "Network"
  }
]
Atlas atlas-vvj19l-shard-0 [primary] myDB> db.food.insert({ _id: 1, fruits: ['grapes', 'mango', 'apple' ] });
{ acknowledged: true, insertedIds: { '0': 1 } }
Atlas atlas-vvj19l-shard-0 [primary] myDB> db.food.insert({ _id: 2, fruits: ['grapes', 'mango', 'cherry' ] });
{ acknowledged: true, insertedIds: { '0': 2 } }
Atlas atlas-vvj19l-shard-0 [primary] myDB> db.food.insert({ _id: 3, fruits: ['banana', 'mango' ] });
{ acknowledged: true, insertedIds: { '0': 3 } }
Atlas atlas-vvj19l-shard-0 [primary] myDB>
Atlas atlas-vvj19l-shard-0 [primary] myDB> db.food.find()
[
  { "_id": 1, "fruits": [ "grapes", "mango", "apple" ] },
  { "_id": 2, "fruits": [ "grapes", "mango", "cherry" ] },
  { "_id": 3, "fruits": [ "banana", "mango" ] }
]
```



```
Activities Terminal Mar 11 15:24 mongosh mongodb+srv://<credentials>@cluster0.tud3r.mongodb.net/

{ _id: 1, fruits: [ 'grapes', 'mango', 'cherry' ] },
{ _id: 2, fruits: [ 'banana', 'mango' ] }
Atlas atlas-vvj19l-shard-0 [primary] myDB> db.food.find( { 'fruits': [ 'grapes', 'mango', 'apple' ] } );
[ { _id: 1, fruits: [ 'grapes', 'mango', 'apple' ] } ]
Atlas atlas-vvj19l-shard-0 [primary] myDB> db.food.find( { 'fruits': 'grapes' } );
[ { _id: 1, fruits: [ 'grapes', 'mango', 'apple' ] } ]
Atlas atlas-vvj19l-shard-0 [primary] myDB> db.food.find( { 'fruits': { $size: 3 } } );
[ { _id: 1, fruits: [ 'grapes', 'mango', 'apple' ] } ]
Atlas atlas-vvj19l-shard-0 [primary] myDB> db.food.find( { _id: 1, 'fruits': { $size: 3 } } );
[ { _id: 1, fruits: [ 'grapes', 'mango', 'apple' ] } ]
Atlas atlas-vvj19l-shard-0 [primary] myDB> db.food.find( { $or: [ 'mango', 'grapes' ] } );
[ { _id: 1, fruits: [ 'grapes', 'mango', 'apple' ] },
  { _id: 2, fruits: [ 'grapes', 'mango', 'cherry' ] } ]
Atlas atlas-vvj19l-shard-0 [primary] myDB> db.food.find( { 'fruits.0': 'grapes' } );
[ { _id: 1, fruits: [ 'grapes', 'mango', 'apple' ] },
  { _id: 2, fruits: [ 'grapes', 'mango', 'cherry' ] } ]
Atlas atlas-vvj19l-shard-0 [primary] myDB> db.food.find()
[ { _id: 1, fruits: [ 'grapes', 'mango', 'apple' ] },
  { _id: 2, fruits: [ 'grapes', 'mango', 'cherry' ] },
  { _id: 3, fruits: [ 'banana', 'mango' ] } ]
Atlas atlas-vvj19l-shard-0 [primary] myDB> db.food.update( { _id: 3 }, { $set: { 'fruits.1': 'apple' } } );
{ acknowledged: true,
  insertedId: null,
  matchedCount: 1,
  modifiedCount: 1,
  upsertedCount: 0 }
Atlas atlas-vvj19l-shard-0 [primary] myDB> db.food.find()
[ { _id: 1, fruits: [ 'grapes', 'mango', 'apple' ] },
  { _id: 2, fruits: [ 'grapes', 'mango', 'cherry' ] },
  { _id: 3, fruits: [ 'banana', 'apple' ] } ]
Atlas atlas-vvj19l-shard-0 [primary] myDB> db.food.update( { _id: 2 }, { $push: { prices: [grapes: 80, mango: 200, cherries: 100] } } );
{ acknowledged: true,
  insertedId: null,
  matchedCount: 1,
  modifiedCount: 1,
  upsertedCount: 0 }
Atlas atlas-vvj19l-shard-0 [primary] myDB> db.food.find()
[ { _id: 1, fruits: [ 'grapes', 'mango', 'apple' ] },
  { _id: 2,
    fruits: [ 'grapes', 'mango', 'cherry' ],
    price: [ { grapes: 80, mango: 200, cherry: 100 } ] },
  { _id: 3, fruits: [ 'banana', 'apple' ] } ]
Atlas atlas-vvj19l-shard-0 [primary] myDB> db.food.update( { _id: 3 }, { $addToSet: { fruits: 'orange' } } );
{ acknowledged: true,
  insertedId: null,
  matchedCount: 1,
  modifiedCount: 1,
  upsertedCount: 0 }
Atlas atlas-vvj19l-shard-0 [primary] myDB> db.food.find()
[ { _id: 1, fruits: [ 'grapes', 'mango', 'apple' ] },
  { _id: 2,
    fruits: [ 'grapes', 'mango', 'cherry' ],
    price: [ { grapes: 80, mango: 200, cherry: 100 } ] },
  { _id: 3, fruits: [ 'banana', 'apple', 'orange' ] } ]
Atlas atlas-vvj19l-shard-0 [primary] myDB> db.food.update( { _id: 2 }, { $pullAll: { fruits: [ 'grapes', 'mango' ] } } );
{ acknowledged: true,
  insertedId: null,
  matchedCount: 1,
  modifiedCount: 1,
  upsertedCount: 0 }
Atlas atlas-vvj19l-shard-0 [primary] myDB> db.food.find()
[ { _id: 1, fruits: [ 'grapes', 'mango', 'apple' ] },
  { _id: 2,
    fruits: [ 'cherry' ],
    price: [ { grapes: 80, mango: 200, cherry: 100 } ] },
  { _id: 3, fruits: [ 'banana', 'apple', 'orange' ] } ]
```

```
Activities Terminal Mar 11 15:25
mongosh mongodb+srv://<credentials>@cluster0.tud3r.mongodb.net/

...
  price: [ { grapes: 40, nango: 200, cherry: 100 } ]
},
{ _id: 3, fruits: [ 'banana', 'apple' ] }
]
Atlas atlas-vvj19l-shard-0 [primary] myDB> db.food.update({_id: 3}, {$pull: {fruits: 'banana'}});
{
  acknowledged: true,
  insertedIds: null,
  matchedCount: 1,
  modifiedCount: 1,
  upsertedCount: 0
}
Atlas atlas-vvj19l-shard-0 [primary] myDB> db.food.find()
[
  { _id: 1, fruits: [ 'grapes', 'nango', 'orange' ] },
  {
    _id: 2,
    fruits: [ 'cherry' ],
    price: [ { grapes: 40, nango: 200, cherry: 100 } ]
  },
  { _id: 3, fruits: [ 'apple' ] }
]
Atlas atlas-vvj19l-shard-0 [primary] myDB> db.createCollection("Customers");
{ ok: 1 }
Atlas atlas-vvj19l-shard-0 [primary] myDB>
Atlas atlas-vvj19l-shard-0 [primary] myDB> db.Customers.insertMany([
... { custID: 1, AcctBal: 1000, AcctType: "S" },
... { custID: 2, AcctBal: 1500, AcctType: "S" },
... { custID: 3, AcctBal: 1200, AcctType: "C" },
... { custID: 4, AcctBal: 800, AcctType: "S" },
... { custID: 5, AcctBal: 500, AcctType: "S" }
... ]);
{
  acknowledged: true,
  insertedIds: {
    "0": ObjectId("67d0b709b790ab5f31c0b67f"),
    "1": ObjectId("67d0b709b790ab5f31c0b680"),
    "2": ObjectId("67d0b709b790ab5f31c0b681"),
    "3": ObjectId("67d0b709b790ab5f31c0b682"),
    "4": ObjectId("67d0b709b790ab5f31c0b683")
  }
}
Atlas atlas-vvj19l-shard-0 [primary] myDB> db.Customer.find()
Atlas atlas-vvj19l-shard-0 [primary] myDB> db.Customers.find()
[
  {
    _id: ObjectId("67d0b709b790ab5f31c0b67f"),
    CustID: 1,
    AcctBal: 1000,
    AcctType: "S"
  },
  {
    _id: ObjectId("67d0b709b790ab5f31c0b680"),
    CustID: 2,
    AcctBal: 1500,
    AcctType: "S"
  },
  {
    _id: ObjectId("67d0b709b790ab5f31c0b681"),
    CustID: 3,
    AcctBal: 1200,
    AcctType: "C"
  },
  {
    _id: ObjectId("67d0b709b790ab5f31c0b682"),
    CustID: 4,
    AcctBal: 800,
    AcctType: "S"
  },
  {
    _id: ObjectId("67d0b709b790ab5f31c0b683"),
    CustID: 5,
    AcctBal: 500,
    AcctType: "S"
  }
]
Atlas atlas-vvj19l-shard-0 [primary] myDB> db.Customers.aggregate([
... {
...   $group: {
...     _id: "CustID",
...     TotAcctBal: { $sum: "$AcctBal" }
...   }
... })
... ]);
[
  { _id: 1, TotAcctBal: 1000 },
  { _id: 2, TotAcctBal: 800 },
  { _id: 3, TotAcctBal: 1200 }
]
Atlas atlas-vvj19l-shard-0 [primary] myDB> db.Customers.aggregate([
... {
...   $group: {
...     _id: "CustID",
...     TotAcctBal: { $sum: "$AcctBal" }
...   }
... })
... ]);
[
  { _id: 1, TotAcctBal: 1000 },
  { _id: 2, TotAcctBal: 800 },
  { _id: 3, TotAcctBal: 1200 }
]
```

```
Activities Terminal Mar 11 15:26 bmsccse@bmsccse-HP-Elite-Tower-800-G9-Desktop-PC: ~  
bmsccse@bmsccse-HP-Elite-Tower-800-G9-Desktop-PC: $ mongoexport  
2025-03-11T14:39:05.282+0530 must specify a collection  
2025-03-11T14:39:05.282+0530 try 'mongoexport --help' for more information  
bmsccse@bmsccse-HP-Elite-Tower-800-G9-Desktop-PC: $ "[[00-mongosh "mongodb+srv://cluster0.tud3r.mongodb.net/" --apiVersion 1 --username marlacs22 --password 83a9u62irEj2Blot  
mongosh: command not found  
bmsccse@bmsccse-HP-Elite-Tower-800-G9-Desktop-PC: $ mongoexport --url "mongodb+srv://cluster0.tud3r.mongodb.net/" --apiVersion 1 --username marlacs22 --password 83a9u62irEj2Blot  
bmsccse@bmsccse-HP-Elite-Tower-800-G9-Desktop-PC: $ mongoexport --url "mongodb+srv://cluster0.tud3r.mongodb.net/" --apiVersion 1 --username marlacs22 --password 83a9u62irEj2Blot  
Current Mongosh Log ID: 67cfdffcb3c86154c8be7d  
Connecting to: mongodb+srv://[redacted]@cluster0.tud3r.mongodb.net/?appName=mongosh+2.2.1  
Using MongoDB: 6.0.5 (API Version 1)  
Using Mongosh: 2.2.1  
Mongosh 2.4.2 is available for download: https://www.mongodb.com/try/download/shell  
For mongosh info see: https://docs.mongodb.com/mongodb-shell/  
Atlas atlas-vvj19l-shard-0 [primary] test> use myDB;  
switched to db myDB  
Atlas atlas-vvj19l-shard-0 [primary] myDB> mongoexport --url "mongodb+srv://marlacs22:83a9u62irEj2Blot@cluster0.tud3r.mongodb.net/myDB" --collection Student1 --type csv --out "D:/output.txt" --fields "Stud  
Name,Grade,Hobbies"  
Uncaught: SyntaxError: Missing semicolon. (1:14)  
> 1 | mongoexport --url "mongodb+srv://marlacs22:83a9u62irEj2Blot@cluster0.tud3r.mongodb.net/myDB" --collection Student1 --type csv --out "D:/output.txt" --fields "StudName,Grade,Hobbies"  
2 |  
Atlas atlas-vvj19l-shard-0 [primary] myDB> mongoexport --url "mongodb+srv://marlacs22:83a9u62irEj2Blot@cluster0.tud3r.mongodb.net/myDB" --collection Student1 --type csv --out "D:/output.txt" --fields "Stud  
Name,Grade,Hobbies"  
Uncaught: SyntaxError: Missing semicolon. (1:14)  
> 1 | mongoexport --url "mongodb+srv://marlacs22:83a9u62irEj2Blot@cluster0.tud3r.mongodb.net/myDB" --collection Student1 --type csv --out "D:/output.txt" --fields "StudName,Grade,Hobbies"  
2 |  
Atlas atlas-vvj19l-shard-0 [primary] myDB>  
(To exit, press Ctrl+C again or Ctrl+D or type .exit)  
Atlas atlas-vvj19l-shard-0 [primary] myDB> mongoexport --url "mongodb+srv://marlacs22:83a9u62irEj2Blot@cluster0.tud3r.mongodb.net/myDB" --collection Student1 --type csv --out "D:/output.txt" --fields "Stud  
Name,Grade,Hobbies"  
Uncaught: SyntaxError: Missing semicolon. (1:14)  
> 1 | mongoexport --url "mongodb+srv://marlacs22:83a9u62irEj2Blot@cluster0.tud3r.mongodb.net/myDB" --collection Student1 --type csv --out "D:/output.txt" --fields "StudName,Grade,Hobbies"  
2 |  
Atlas atlas-vvj19l-shard-0 [primary] myDB> exit  
bmsccse@bmsccse-HP-Elite-Tower-800-G9-Desktop-PC: $ mongoexport --url "mongodb+srv://marlacs22:83a9u62irEj2Blot@cluster0.tud3r.mongodb.net/myDB" --collection Student1 --type csv --out "D:/output.txt" --fi  
elds "StudName,Grade,Hobbies"  
2025-03-11T14:42:20.815+0530 connected to: mongodb+srv://[redacted]@cluster0.tud3r.mongodb.net/myDB  
2025-03-11T14:42:21.808+0530 exported 2 records  
bmsccse@bmsccse-HP-Elite-Tower-800-G9-Desktop-PC: $  
bmsccse@bmsccse-HP-Elite-Tower-800-G9-Desktop-PC: $ mongoexport --url "mongodb+srv://marlacs22:83a9u62irEj2Blot@cluster0.tud3r.mongodb.net/myDB" --collection Student1 --type csv --out "D:/out.txt" --fi  
elds "StudName,Grade,Hobbies"  
2025-03-11T14:42:21.808+0530 connected to: mongodb+srv://[redacted]@cluster0.tud3r.mongodb.net/myDB  
2025-03-11T14:42:21.808+0530 exported 2 records  
bmsccse@bmsccse-HP-Elite-Tower-800-G9-Desktop-PC: $  
bmsccse@bmsccse-HP-Elite-Tower-800-G9-Desktop-PC: $ mongoexport --url "mongodb+srv://marlacs22:83a9u62irEj2Blot@cluster0.tud3r.mongodb.net/myDB" --collection Student1 --type csv --out "D:/out.txt" --fi  
elds "StudName,Grade,Hobbies"  
2025-03-11T14:43:35.174+0530 connected to: mongodb+srv://[redacted]@cluster0.tud3r.mongodb.net/myDB  
2025-03-11T14:43:35.437+0530 exported 2 records  
bmsccse@bmsccse-HP-Elite-Tower-800-G9-Desktop-PC: $  
bmsccse@bmsccse-HP-Elite-Tower-800-G9-Desktop-PC: $ mongoexport --url "mongodb+srv://marlacs22:83a9u62irEj2Blot@cluster0.tud3r.mongodb.net/myDB" --collection Student1 --type csv --out "D:/qwer.txt" --fi  
elds "StudName,Grade,Hobbies"  
2025-03-11T14:44:44.222+0530 connected to: mongodb+srv://[redacted]@cluster0.tud3r.mongodb.net/myDB  
2025-03-11T14:44:44.529+0530 exported 2 records  
bmsccse@bmsccse-HP-Elite-Tower-800-G9-Desktop-PC: $ whoami  
bmsccse  
bmsccse@bmsccse-HP-Elite-Tower-800-G9-Desktop-PC: $ ^C  
bmsccse@bmsccse-HP-Elite-Tower-800-G9-Desktop-PC: $ mongoexport --url "mongodb+srv://marlacs22:83a9u62irEj2Blot@cluster0.tud3r.mongodb.net/myDB" --collection Student1 --type csv --out "/home/bmsccse/De  
sktop/qwer.csv" --fields "StudName,Grade,Hobbies"  
2025-03-11T14:47:52.993+0530 connected to: mongodb+srv://[redacted]@cluster0.tud3r.mongodb.net/myDB  
bmsccse@bmsccse-HP-Elite-Tower-800-G9-Desktop-PC: $ mongoimport --url "mongodb+srv://marlacs22:83a9u62irEj2Blot@cluster0.tud3r.mongodb.net/myDB" --collection Student3 --type csv --file "/home/bmsccse/D  
esktop/qwer.csv" --headerline  
2025-03-11T14:51:32.165+0530 connected to: mongodb+srv://[redacted]@cluster0.tud3r.mongodb.net/myDB  
2025-03-11T14:51:32.316+0530 2 document(s) imported successfully, 0 document(s) failed to import.  
bmsccse@bmsccse-HP-Elite-Tower-800-G9-Desktop-PC: $  
bmsccse@bmsccse-HP-Elite-Tower-800-G9-Desktop-PC: $ mongoimport --url "mongodb+srv://marlacs22:83a9u62irEj2Blot@cluster0.tud3r.mongodb.net/myDB" --collection Student1 --type csv --headerline --file ~/Do  
wnloads/airline.csv  
2025-03-11T14:55:09.112+0530 Failed: open /home/bmsccse/Downloads/airline.csv: no such file or directory  
2025-03-11T14:55:09.112+0530 0 document(s) imported successfully, 0 document(s) failed to import.  
bmsccse@bmsccse-HP-Elite-Tower-800-G9-Desktop-PC: $ mongoimport --url "mongodb+srv://marlacs22:83a9u62irEj2Blot@cluster0.tud3r.mongodb.net/myDB" --collection Student1 --type csv --headerline --file ~/De  
sktop/airline.csv  
2025-03-11T14:56:05.413+0530 Failed: open /home/bmsccse/Desktop/airline.csv: no such file or directory  
2025-03-11T14:56:05.413+0530 0 document(s) imported successfully, 0 document(s) failed to import.  
bmsccse@bmsccse-HP-Elite-Tower-800-G9-Desktop-PC: $ mongoimport --local localhost:27020 --collection Student1 --type csv --out ~/Downloads/airline.csv --fields "Year","Quarter"
```



Activities Terminal Mar 11 15:26

bmsecse@bmsecse-HP-Elite-Tower-800-G9-Desktop-PC: ~

```

dName,Grade,Hobbies"
Uncaught:
SyntaxError: Missing semicolon. (1:14)
> 1 | mongoexport --url "mongodb+srv://narlacs22:8j9u6zrEjz8lot@cluster0.tud3r.mongodb.net/my08" --collection Student1 --type csv --out "D:/output.txt" --fields "StudName,Grade,Hobbies"
  |
  |
Atlas atlas-vvj9l-shard-0 [primary] myDB> exit
bmsecse@bmsecse-HP-Elite-Tower-800-G9-Desktop-PC: $ mongoexport --url "mongodb+srv://narlacs22:8j9u6zrEjz8lot@cluster0.tud3r.mongodb.net/my08" --collection Student1 --type csv --out "D:/output.txt" --fields "StudName,Grade,Hobbies"
2025-03-11T14:42:20.815+0530 connected to: mongodb+srv://[REDACTED]@cluster0.tud3r.mongodb.net/my08
2025-03-11T14:42:21.800+0530 exported 2 records
bmsecse@bmsecse-HP-Elite-Tower-800-G9-Desktop-PC: $
bmsecse@bmsecse-HP-Elite-Tower-800-G9-Desktop-PC: $ mongoexport --url "mongodb+srv://narlacs22:8j9u6zrEjz8lot@cluster0.tud3r.mongodb.net/my08" --collection Student1 --type csv --out "D:/output.txt" --file
ids "StudName,Grade,Hobbies"
2025-03-11T14:43:35.174+0530 connected to: mongodb+srv://[REDACTED]@cluster0.tud3r.mongodb.net/my08
2025-03-11T14:43:35.437+0530 exported 2 records
bmsecse@bmsecse-HP-Elite-Tower-800-G9-Desktop-PC: $
bmsecse@bmsecse-HP-Elite-Tower-800-G9-Desktop-PC: $ mongoexport --url "mongodb+srv://narlacs22:8j9u6zrEjz8lot@cluster0.tud3r.mongodb.net/my08" --collection Student1 --type csv --out "D:/qwer.txt" --fi
elds "StudName,Grade,Hobbies"
2025-03-11T14:44:44.529+0530 connected to: mongodb+srv://[REDACTED]@cluster0.tud3r.mongodb.net/my08
2025-03-11T14:44:44.529+0530 exported 2 records
bmsecse@bmsecse-HP-Elite-Tower-800-G9-Desktop-PC: $ whoami
bmsecse
bmsecse@bmsecse-HP-Elite-Tower-800-G9-Desktop-PC: $ nc
bmsecse@bmsecse-HP-Elite-Tower-800-G9-Desktop-PC: $ mongoexport --url "mongodb+srv://narlacs22:8j9u6zrEjz8lot@cluster0.tud3r.mongodb.net/my08" --collection Student1 --type csv --out "/home/bmsecse/be
sktop/qwer.csv" --fields "StudName,Grade,Hobbies"
2025-03-11T14:47:53.277+0530 connected to: mongodb+srv://[REDACTED]@cluster0.tud3r.mongodb.net/my08
2025-03-11T14:47:53.277+0530 exported 2 records
bmsecse@bmsecse-HP-Elite-Tower-800-G9-Desktop-PC: $ mongoimport --url "mongodb+srv://narlacs22:8j9u6zrEjz8lot@cluster0.tud3r.mongodb.net/my08" --collection Student3 --type csv --file "/home/bmsecse/D
esktop/qwer.csv" --headerline
2025-03-11T14:51:32.165+0530 connected to: mongodb+srv://[REDACTED]@cluster0.tud3r.mongodb.net/my08
2025-03-11T14:51:32.316+0530 2 document(s) imported successfully. 0 document(s) failed to import.
bmsecse@bmsecse-HP-Elite-Tower-800-G9-Desktop-PC: $
bmsecse@bmsecse-HP-Elite-Tower-800-G9-Desktop-PC: $ mongoimport --url "mongodb+srv://narlacs22:8j9u6zrEjz8lot@cluster0.tud3r.mongodb.net/my08" --collection Student1 --type csv --headerline --file ~/Do
wnloads/airline.csv
2025-03-11T14:55:09.112+0530 Failed: open /home/bmsecse/Downloads/airline.csv: no such file or directory
2025-03-11T14:55:09.112+0530 0 document(s) imported successfully. 0 document(s) failed to import.
bmsecse@bmsecse-HP-Elite-Tower-800-G9-Desktop-PC: $ mongoimport --url "mongodb+srv://narlacs22:8j9u6zrEjz8lot@cluster0.tud3r.mongodb.net/my08" --collection Student1 --type csv --headerline --file ~/De
sktop/airline.csv
2025-03-11T14:56:05.413+0530 Failed: open /home/bmsecse/Desktop/airline.csv: no such file or directory
2025-03-11T14:56:05.413+0530 0 document(s) imported successfully. 0 document(s) failed to import.
bmsecse@bmsecse-HP-Elite-Tower-800-G9-Desktop-PC: $ mongoexport --host localhost --db my08 --collection Student1 --type=csv --out ~/Downloads/airline.csv --fields "Year","Quarter"
2025-03-11T14:58:48.070+0530 connected to: mongodb://localhost/
2025-03-11T14:58:48.071+0530 exported 0 records
bmsecse@bmsecse-HP-Elite-Tower-800-G9-Desktop-PC: $

```

Activities LibreOffice 7.3 Mar 11 15:27

LibreOffice Community

Text Import - [qwer.csv]

Import

Character set: Unicode (UTF-8)

Language: Default - English (India)

From row: 1

Separator Options

Fixed width ☐ Separated by ☒

☒ Tab ☒ Comma ☒ Semicolon ☐ Space ☐ Other

☐ Merge delimiters ☐ Trim spaces String delimiter: "

Other Options

☐ Format quoted field as text ☐ Detect special numbers

☐ Evaluate formulas

Fields

Column type:

	Standard	Standard	Standard
1	StudName	Grade	Hobbies
2	NichelleJacintha	VII	InternetSurfing
3	AryanDavid	VII	Chess

Help Cancel OK

sgpa.class  
sgpa.java  
qwer.csv  
Student.java  
Student.class  
main.java  
Student.java  
calsgpa.java  
Quad.class  
calsgpa.java  
QuadraticEquationSolver.class  
CalSgpa.java  
Quad.java  
calsgpa.class  
Home

cluster0.mongodb.net > test > New\_Student Open MongoDB shell

Documents Aggregations Schema Indexes Validation

Type a query: { field: 'value' } or [Generate Query](#) Explain Reset Find Options

[ADD DATA](#) [EXPORT DATA](#) [UPDATE](#) [DELETE](#) 25 1 - 6 of 6 < > ⌵ ⌶

<pre>{ "_id": "67c6c029acf8a6e28fa226",   "RollNo": 2,   "Age": 21,   "Cont": 5576,   "email": "anubhav.de9@gmail.com" }</pre>
<pre>{ "_id": "67c6c03aacf8a6e28fa225",   "RollNo": 2,   "Age": 22,   "Cont": 9976,   "email": "anushka.de9@gmail.com" }</pre>
<pre>{ "_id": "67c6c04aacf8a6e28fa224",   "RollNo": 1,   "Age": 21,   "Cont": 9876,   "email": "antara.de9@gmail.com" }</pre>
<pre>{ "_id": "67c6c05aacf8a6e28fa223",   "RollNo": 11,   "Age": 22,   "Name": "FER",   "Cont": 2276,   "email": "frea.de9@gmail.com" }</pre>
<pre>{ "_id": "67c6c07aacf8a6e28fa227",   "RollNo": 4,   "Age": 20,   "Cont": 4476,   "email": "pani.de9@gmail.com" }</pre>
<pre>{ "_id": "67c6c08aacf8a6e28fa228",   "RollNo": 10,   "Age": 23,   "Cont": 2276,   "email": "Abhinav@gmail.com" }</pre>
<pre>RollNo: 3 Age: 21 Cont: 5576 email: "anubhav.de9@gmail.com"</pre>
<pre>{ "_id": "67c6c01aacf8a6e28fa215",   "RollNo": 2,   "Age": 22,   "Cont": 9976,   "email": "anushka.de9@gmail.com" }</pre>
<pre>{ "_id": "67c6c014aacf8a6e28fa214",   "RollNo": 1,   "Age": 21,   "Cont": 9876,   "email": "antara.de9@gmail.com" }</pre>
<pre>{ "_id": "67c6c05aacf8a6e28fa219",   "RollNo": 11,   "Age": 22,   "Name": "FER",   "Cont": 2276,   "email": "frea.de9@gmail.com" }</pre>
<pre>{ "_id": "67c6c02f1aacf8a6e28fa217",   "RollNo": 4,   "Age": 20,   "Cont": 4476,   "email": "pani.de9@gmail.com" }</pre>
<pre>{ "_id": "67c6c03aacf8a6e28fa218",   "RollNo": 10,   "Age": 23,   "Cont": 2276,   "email": "Abhinav@gmail.com" }</pre>

Windows taskbar: Search, 15:20:44, 04-03-2022

Compass

cluster0.tud3.mongodb.net > test > Student

Documents | Aggregations | Schema | Indexes | Validation

Type a query: { field: 'value' } or [Generate Query](#)

EXPLAIN | RESET | FIND | Options

25 | 1-6 of 6

ADD DATA | EXPORT DATA | UPDATE | DELETE

_id: ObjectId('67c6b141bc78a628fa21a4')
RollNo: 1
Age: 21
Cont: 9876
email: "anirudh.deshgani@gmail.com"

_id: ObjectId('67c6b141bc78a628fa21a5')
RollNo: 2
Age: 22
Cont: 9876
email: "anirudh.deshgani@gmail.com"

_id: ObjectId('67c6b141bc78a628fa21a6')
RollNo: 3
Age: 21
Cont: 9876
email: "anirudh.deshgani@gmail.com"

_id: ObjectId('67c6b141bc78a628fa21a7')
RollNo: 4
Age: 20
Cont: 4476
email: "pari.deshgani@gmail.com"

_id: ObjectId('67c6b141bc78a628fa21a8')
RollNo: 10
Age: 23
Cont: 2276
email: "Abhinav@gmail.com"

_id: ObjectId('67c6b141bc78a628fa21a9')
RollNo: 11
Age: 22
Name: "Rishi"
Cont: 2276
email: "res.deshgani@gmail.com"

32°C  
Partly sunny

15:21:58  
04-09-2025

MongoDB Compass - cluster0.tud3.mongodb.net/test/Student

Connections | Settings | View | Collection | Help

Compass

My Queries

Search connections

Connections

- BDALab
- cluster0.tud3.mongodb.net
  - admin
  - config
  - local
  - myDB
  - sample\_mflix
  - test
    - New\_Student
    - Student

cluster0.tud3.mongodb.net > test > Student

Documents | Aggregations | Schema | Indexes | Validation

Type a query: { field: 'value' } or [Generate Query](#)

EXPLAIN | RESET | FIND | Options

25 | 1-6 of 6

ADD DATA | EXPORT DATA | UPDATE | DELETE

RollNo: 1
Age: 21
Cont: 9876
email: "anirudh.deshgani@gmail.com"

_id: ObjectId('67c6b141bc78a628fa21a5')
RollNo: 2
Age: 22
Cont: 9876
email: "anirudh.deshgani@gmail.com"

_id: ObjectId('67c6b141bc78a628fa21a6')
RollNo: 3
Age: 21
Cont: 9876
email: "anirudh.deshgani@gmail.com"

_id: ObjectId('67c6b141bc78a628fa21a7')
RollNo: 4
Age: 20
Cont: 4476
email: "pari.deshgani@gmail.com"

_id: ObjectId('67c6b141bc78a628fa21a8')
RollNo: 10
Age: 23
Cont: 2276
email: "Abhinav@gmail.com"

_id: ObjectId('67c6b141bc78a628fa21a9')
RollNo: 11
Age: 22
Name: "Rishi"
Cont: 2276
email: "res.deshgani@gmail.com"

32°C  
Partly sunny

15:22:11  
04-09-2025





```
Command Prompt
Microsoft Windows [Version 10.0.22631.4890]
(c) Microsoft Corporation. All rights reserved.

C:\Users\STUDENT>mongoexport --uri="mongodb+srv://mariacs22:8Ja9u6zIrEjzBlot@cluster0.tud3r.mongodb.net/test" --collection=Student --out=abc.csv
2025-03-04T15:09:05.137+0530    connected to: mongodb+srv://[**REDACTED**]@cluster0.tud3r.mongodb.net/test
2025-03-04T15:09:05.359+0530    exported 6 records

C:\Users\STUDENT>mongoexport --uri="mongodb+srv://mariacs22:8Ja9u6zIrEjzBlot@cluster0.tud3r.mongodb.net/test" --collection=Student --out=efg.json
2025-03-04T15:14:50.067+0530    connected to: mongodb+srv://[**REDACTED**]@cluster0.tud3r.mongodb.net/test
2025-03-04T15:14:50.288+0530    exported 6 records

C:\Users\STUDENT>mongoimport --uri="mongodb+srv://mariacs22:8Ja9u6zIrEjzBlot@cluster0.tud3r.mongodb.net/test" --collection=New_Student --type=json --file="C:\Users\STUDENT\efg.json"
2025-03-04T15:15:35.501+0530    connected to: mongodb+srv://[**REDACTED**]@cluster0.tud3r.mongodb.net/test
2025-03-04T15:15:35.590+0530    6 document(s) imported successfully. 0 document(s) failed to import.

C:\Users\STUDENT>
```

## EXPERIMENT-2

Neo4j

Code with Screenshots:

The top screenshot shows the Neo4j Cypher Shell interface. The query entered is `neo4j$ CREATE (a:STUDENT {name:"Ram", id: 12})`. The execution status bar indicates: "Added 1 label, created 1 node, set 2 properties, completed after 27 ms." The left sidebar shows the "Table" and "Code" views.

The bottom screenshot shows the Neo4j Graph Explorer interface. The query entered is `neo4j$ MATCH (n) RETURN n`. The graph view displays a single node labeled "Ram". The right sidebar shows the "Overview" panel with the following information:

- Node labels:** \* (1) STUDENT (1)
- Displaying 1 nodes, 0 relationships.

neo4j\$ CREATE (a:STUDENT {name:"Abhi" , id: 14})

Table

Code

Added 1 label, created 1 node, set 2 properties, completed after 6 ms.

neo4j\$ CREATE (a:STUDENT {name:"Raj" , id: 13})

Table

Code


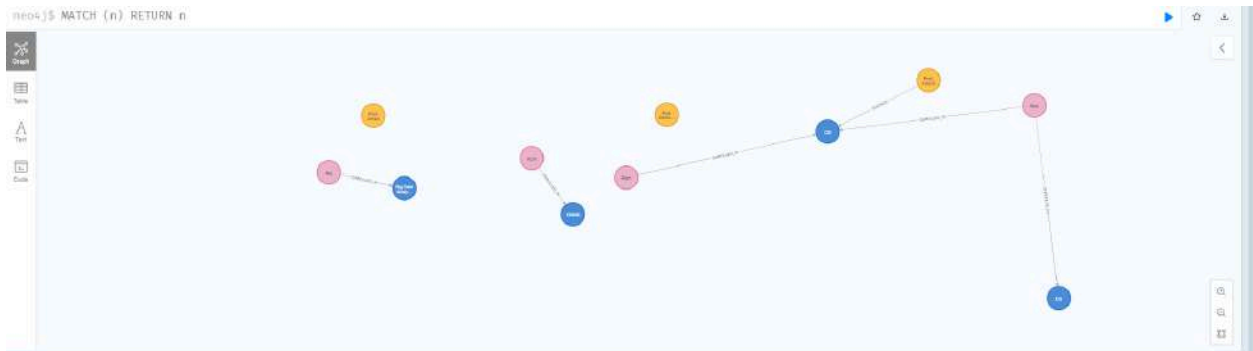
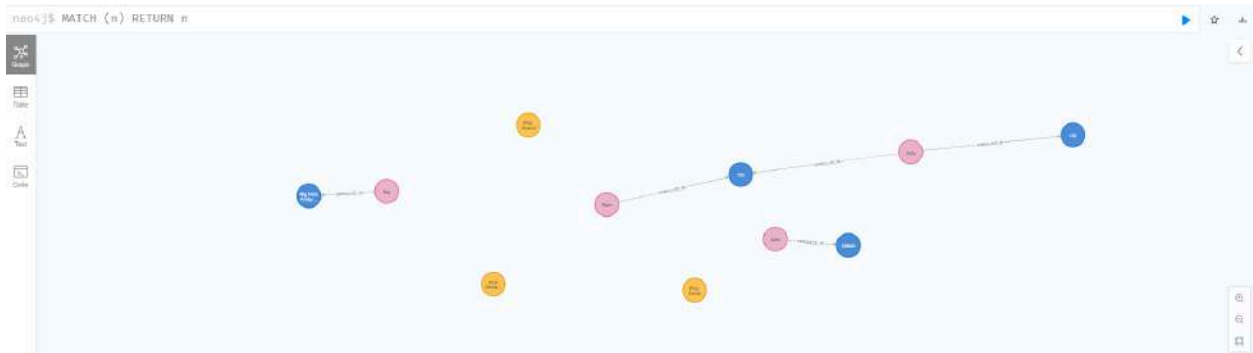
Added 1 label, created 1 node, set 2 properties, completed after 6 ms.

neo4j\$ CREATE (a:STUDENT {name:"Ajay" , id: 15})

Table

Code

Added 1 label, created 1 node, set 2 properties, completed after 6 ms.







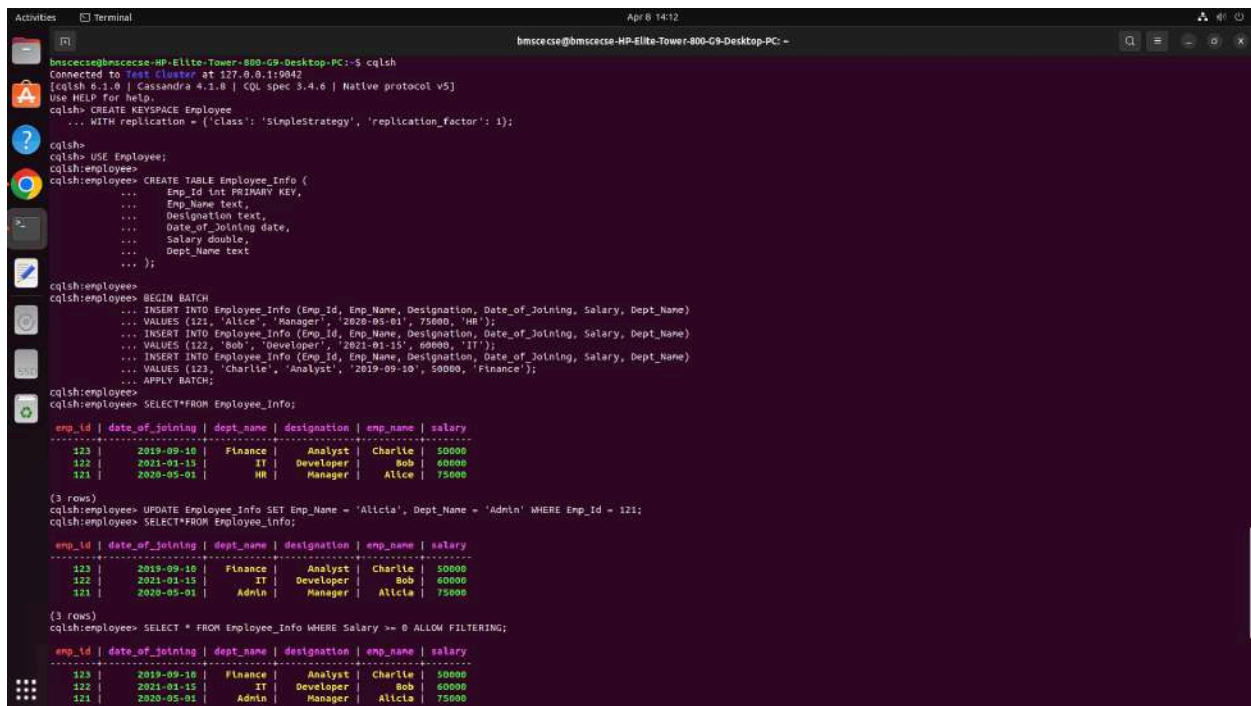


## EXPERIMENT-3

Perform the following DB operations using Cassandra.

- Create a keyspace by name Employee
- Create a column family by name Employee-Info with attributes Emp\_Id Primary Key, Emp\_Name, Designation, Date\_of\_Joining, Salary, Dept\_Name
- Insert the values into the table in batch
- Update Employee name and Department of Emp-Id 121
- Sort the details of Employee records based on salary
- Alter the schema of the table Employee\_Info to add a column Projects which stores a set of Projects done by the corresponding Employee.
- Update the altered table to add project names.
- Create a TTL of 15 seconds to display the values of Employees.

CODE WITH SCREENSHOTS:



```
Activities Terminal Apr 8 14:12
bmscscse@bmscscse-HP-Elite-Tower-800-G9-Desktop-PC: ~
bmscscse@bmscscse-HP-Elite-Tower-800-G9-Desktop-PC:~$ cqlsh
Connected to Test cluster at 127.0.0.1:9042
[cqlsh 6.1.0 | Cassandra 4.1.0 | CQL spec 3.4.0 | Native protocol v5]
Use HELP for help.
cqlsh> CREATE KEYSPACE Employee
... WITH replication = {'class': 'SimpleStrategy', 'replication_factor': 1};

cqlsh> USE Employee;
cqlsh:Employee>
cqlsh:Employee> CREATE TABLE Employee_Info (
...     Emp_Id int PRIMARY KEY,
...     Emp_Name text,
...     Designation text,
...     Date_of_Joining date,
...     Salary double,
...     Dept_Name text
... );

cqlsh:Employee>
cqlsh:Employee> BEGIN BATCH
... INSERT INTO Employee_Info (Emp_Id, Emp_Name, Designation, Date_of_Joining, Salary, Dept_Name)
... VALUES (121, 'Alice', 'Manager', '2020-05-01', 75000, 'HR');
... INSERT INTO Employee_Info (Emp_Id, Emp_Name, Designation, Date_of_Joining, Salary, Dept_Name)
... VALUES (122, 'Bob', 'Developer', '2021-01-15', 60000, 'IT');
... INSERT INTO Employee_Info (Emp_Id, Emp_Name, Designation, Date_of_Joining, Salary, Dept_Name)
... VALUES (123, 'Charlie', 'Analyst', '2019-09-10', 50000, 'Finance');
... APPLY BATCH;

cqlsh:Employee> SELECT * FROM Employee_Info;

emp_id | date_of_joining | dept_name | designation | emp_name | salary
-----|-----|-----|-----|-----|-----
123 | 2019-09-10 | Finance | Analyst | Charlie | 50000
122 | 2021-01-15 | IT | Developer | Bob | 60000
121 | 2020-05-01 | HR | Manager | Alice | 75000

(3 rows)
cqlsh:Employee> UPDATE Employee_Info SET Emp_Name = 'Alicia', Dept_Name = 'Admin' WHERE Emp_Id = 121;
cqlsh:Employee> SELECT * FROM Employee_Info;

emp_id | date_of_joining | dept_name | designation | emp_name | salary
-----|-----|-----|-----|-----|-----
122 | 2019-09-10 | Finance | Analyst | Charlie | 50000
122 | 2021-01-15 | IT | Developer | Bob | 60000
121 | 2020-05-01 | Admin | Manager | Alicia | 75000

(3 rows)
cqlsh:Employee> SELECT * FROM Employee_Info WHERE Salary >= 0 ALLOW FILTERING;

emp_id | date_of_joining | dept_name | designation | emp_name | salary
-----|-----|-----|-----|-----|-----
123 | 2019-09-10 | Finance | Analyst | Charlie | 50000
122 | 2021-01-15 | IT | Developer | Bob | 60000
121 | 2020-05-01 | Admin | Manager | Alicia | 75000
```

```
Activities Terminal Apr 8 14:12 bmsccse@bmsccse-HP-Elite-Tower-800-G9-Desktop-PC: -

121 | 2020-05-01 | HR | Manager | Alice | 75000
(3 rows)
cqlsh:employee> UPDATE Employee_Info SET Emp_Name = 'Alicia', Dept_Name = 'Admin' WHERE Emp_Id = 121;
cqlsh:employee> SELECT*FROM Employee_Info;

emp_id | date_of_joining | dept_name | designation | emp_name | salary
-----|-----|-----|-----|-----|-----
123 | 2019-09-10 | Finance | Analyst | Charlie | 50000
122 | 2021-01-15 | IT | Developer | Bob | 60000
121 | 2020-05-01 | Admin | Manager | Alicia | 75000
(3 rows)
cqlsh:employee> SELECT * FROM Employee_Info WHERE Salary >= 0 ALLOW FILTERING;

emp_id | date_of_joining | dept_name | designation | emp_name | salary
-----|-----|-----|-----|-----|-----
123 | 2019-09-10 | Finance | Analyst | Charlie | 50000
122 | 2021-01-15 | IT | Developer | Bob | 60000
121 | 2020-05-01 | Admin | Manager | Alicia | 75000
(3 rows)
cqlsh:employee> ALTER TABLE Employee_Info ADD Projects set<text>;
cqlsh:employee> SELECT*FROM Employee_Info;

emp_id | date_of_joining | dept_name | designation | emp_name | projects | salary
-----|-----|-----|-----|-----|-----|-----
123 | 2019-09-10 | Finance | Analyst | Charlie | null | 50000
122 | 2021-01-15 | IT | Developer | Bob | null | 60000
121 | 2020-05-01 | Admin | Manager | Alicia | null | 75000
(3 rows)
cqlsh:employee> UPDATE Employee_Info SET Projects = ('Payroll System', 'Attendance Tracker') WHERE Emp_Id = 121;
cqlsh:employee> SELECT*FROM Employee_Info;

emp_id | date_of_joining | dept_name | designation | emp_name | projects | salary
-----|-----|-----|-----|-----|-----|-----
123 | 2019-09-10 | Finance | Analyst | Charlie | null | 50000
122 | 2021-01-15 | IT | Developer | Bob | null | 60000
121 | 2020-05-01 | Admin | Manager | Alicia | ('Attendance Tracker', 'Payroll System') | 75000
(3 rows)
cqlsh:employee> INSERT INTO Employee_Info (Emp_Id, Emp_Name, Designation, Date_of_Joining, Salary, Dept_Name)
... VALUES (124, 'David', 'Tester', '2022-03-05', 45000, 'QA') USING TTL 15;
cqlsh:employee> select*from employee_info;

emp_id | date_of_joining | dept_name | designation | emp_name | projects | salary
-----|-----|-----|-----|-----|-----|-----
123 | 2019-09-10 | Finance | Analyst | Charlie | null | 50000
122 | 2021-01-15 | IT | Developer | Bob | null | 60000
121 | 2020-05-01 | Admin | Manager | Alicia | ('Attendance Tracker', 'Payroll System') | 75000
(3 rows)
cqlsh:employee>

bmsccse@bmsccse-HP-Elite-Tower-800-G9-Desktop-PC: -
cqlsh
Connected to Test Cluster at 127.0.0.1:9042
[cqlsh 4.1.0 | Cassandra 4.1.0 | Cql spec 3.4.6 | Native protocol v5]
Use HELP for help.
cqlsh> CREATE KEYSPACE Library
... WITH REPLICATION = {'class': 'SimpleStrategy', 'replication_factor': 1};
cqlsh> USE Library;
cqlsh:library>
cqlsh:library> CREATE TABLE Library_Info (
... Stud_id int PRIMARY KEY,
... Counter_value counter,
... Stud_Name text,
... Book_Name text,
... Book_id int,
... Date_of_issue date
... );
InvalidRequest: Error from server: code=2200 [Invalid query] message="Cannot mix counter and non counter columns in the same table"
cqlsh:library> CREATE TABLE Library_Info (
... Stud_id int PRIMARY KEY,
... Stud_Name text,
... Book_Name text,
... Book_id int,
... Date_of_issue date
... );
cqlsh:library>
cqlsh:library> CREATE TABLE Library_Counter (
... Stud_id int PRIMARY KEY,
... Counter_value counter
... );
cqlsh:library>
cqlsh:library> SELECT*FROM Library_Info;

stud_id | book_id | book_name | date_of_issue | stud_name
-----|-----|-----|-----|-----
(0 rows)
cqlsh:library> CREATE TABLE Library_Info1(
... Stud_id int,
... Book_Name text,
... Counter_Value counter,
... PRIMARY KEY (Stud_id, Book_Name)
... );
cqlsh:library> DROP TABLE LIBRARY_INFO;

cqlsh:library>
cqlsh:library>
cqlsh:library> DROP TABLE LIBRARY_INFO1;
cqlsh:library> CREATE TABLE Library_Info (
... Stud_id int,
... Book_Name text,
... Counter_Value counter,
... PRIMARY KEY (Stud_id, Book_Name)
... );
cqlsh:library>
```

```
Activities Terminal Apr 6 14:25 bmsccse@bmsccse-HP-Elite-Tower-800-G9-Desktop-PC: -

... Counter_Value counter,
... PRIMARY KEY (Stud_Id, Book_Name)
... );

cqlsh:library> BEGIN BATCH
... UPDATE Library_Info SET Counter_Value = Counter_Value + 1 WHERE Stud_Id = 112 AND Book_Name = 'BDA';
... UPDATE Library_Info SET Counter_Value = Counter_Value + 1 WHERE Stud_Id = 113 AND Book_Name = 'DSA';
... UPDATE Library_Info SET Counter_Value = Counter_Value + 1 WHERE Stud_Id = 112 AND Book_Name = 'BDA';
... APPLY BATCH;
InvalidRequest: Error from server: code=2280 [Invalid query] message="Cannot include a counter statement in a logged batch"
cqlsh:library> BEGIN UNLOGGED BATCH
... UPDATE Library_Info SET Counter_Value = Counter_Value + 1 WHERE Stud_Id = 112 AND Book_Name = 'BDA';
... UPDATE Library_Info SET Counter_Value = Counter_Value + 1 WHERE Stud_Id = 113 AND Book_Name = 'DSA';
... UPDATE Library_Info SET Counter_Value = Counter_Value + 1 WHERE Stud_Id = 112 AND Book_Name = 'BDA';
... APPLY BATCH;
cqlsh:library> SELECT * FROM Library_Info;

stud_id | book_name | counter_value
-----|-----|-----
113 | DSA | 1
112 | BDA | 2
(2 rows)
cqlsh:library> SELECT * FROM Library_Info WHERE Stud_Id = 112 AND Book_Name = 'BDA';

stud_id | book_name | counter_value
-----|-----|-----
112 | BDA | 2
(1 rows)
cqlsh:library> COPY Library_Info TO '/home/your_username/library_info.csv' WITH HEADER = TRUE;
Using 16 child processes
Can't open '/home/your_username/library_info.csv' for writing: [Errno 2] No such file or directory: '/home/your_username/library_info.csv'
cqlsh:library> COPY Library_Info (stud_id, book_name, counter_value) TO '/home/bmsccse/library_info.csv';
Using 16 child processes

Starting copy of library.library_info with columns [stud_id, book_name, counter_value].
Processed: 2 rows; Rate: 22 rows/s; Avg. rate: 22 rows/s
2 rows exported to 1 files in 0.105 seconds.
cqlsh:library> COPY Library_Info (stud_id, book_name, counter_value) FROM '/home/bmsccse/library_info.csv' WITH HEADER = true;
Using 16 child processes

Starting copy of library.library_info with columns [stud_id, book_name, counter_value].
Processed: 1 rows; Rate: 2 rows/s; Avg. rate: 3 rows/s
1 rows imported from 1 files in 0.158 seconds (0 skipped).
cqlsh:library> select * from library.library_info;

stud_id | book_name | counter_value
-----|-----|-----
113 | DSA | 1
112 | BDA | 4
(2 rows)
cqlsh:library>
```

Perform the following DB operations using Cassandra.

- Create a keyspace by name Library
- Create a column family by name Library-Info with attributes  
Stud\_Id Primary Key,  
Counter\_value of type Counter,  
Stud\_Name, Book-Name, Book-Id,  
Date\_of\_issue
- Insert the values into the table in batch
- Display the details of the table created and increase the value of the counter
- Write a query to show that a student with id 112 has taken a book “BDA” 2 times.
- Export the created column to a csv file
- Import a given csv dataset from local file system into Cassandra column family

```

bmcscse@bmcscse-HP-Elite-Tower-800-G9-Desktop-PC:~$ [200-docker run --name cassandra -d -p 9042:9042 cassandra:latest
docker: command not found
bmcscse@bmcscse-HP-Elite-Tower-800-G9-Desktop-PC:~$ -docker run --name cassandra -d -p 9042:9042 cassandra:latest
Command '-docker' not found, did you mean:
  command 'docker' from snap docker (27.5.1)
  command 'docker' from deb docker.io (20.1.3-0ubuntu1-22.04.1)
  command 'docker' from deb podman-docker (3.4.4+ds1-1ubuntu1-22.04.3)
  command 'kdocker' from deb kdocker (5.3.1)
See 'snap info <snapname>' for additional versions.
bmcscse@bmcscse-HP-Elite-Tower-800-G9-Desktop-PC:~$ sudo docker logs -f cassandra
Command 'docker' not found, but can be installed with:
sudo snap install docker # version 27.5.1, or
sudo apt install docker.io # version 20.1.3-0ubuntu1-22.04.1
sudo apt install podman-docker # version 3.4.4+ds1-1ubuntu1-22.04.3
See 'snap info docker' for additional versions.
bmcscse@bmcscse-HP-Elite-Tower-800-G9-Desktop-PC:~$ sudo apt update
sudo apt install cassandra -y
[sudo] password for bmcscse:
bmcscse: try again.
[sudo] password for bmcscse:
bmcscse: try again.
[sudo] password for bmcscse:
bmcscse: try again.
[sudo] password for bmcscse:
Reading package lists... Done
Building dependency tree... Done
Reading state information... Done
The following packages were automatically installed and are no longer required:
  linux-headers-6.2.0-26-generic linux-hwe-6.2-headers-6.2.0-26
  linux-image-6.2.0-26-generic linux-modules-6.2.0-26-generic
  linux-modules-extra-6.2.0-26-generic
Use 'sudo apt autoremove' to remove them.
Suggested packages:
  cassandra-tools
The following packages will be upgraded:
  cassandra
1 upgraded, 0 newly installed, 0 to remove and 316 not upgraded.
Need to get 48.7 MB of archives.
After this operation, 68.4 MB of additional disk space will be used.
Get:1 https://apache.jfrog.io/artifactory/cassandra-deb 4ix/main amd64 cassandra all 4.1.8 [48.7 MB]
Fetched 48.7 MB in 25s (1,949 kB/s)
(Reading database ... 250222 files and directories currently installed.)
Preparing to unpack .../cassandra-4.1.8_0_all.deb ...
Unpacking cassandra (4.1.8) over (4.1.5) ...
Setting up cassandra (4.1.8) ...
Installing new version of config file /etc/cassandra/cassandra.yaml ...
Installing new version of config file /etc/cassandra/commitlog_archiving.properties ...
Installing new version of config file /etc/cassandra/credentials.sample ...
Installing new version of config file /etc/cassandra/jvm-server.options ...

```



```
vm.max_map_count = 1048575
net.ipv4.tcp_keepalive_time = 300
update-rc.d warning: start and stop actions are no longer supported; falling ba
ck to defaults
bmscsec@bmscsec-HP-Elite-Tower-800-G9-Desktop-PC:~$ cassandra -f
OpenJDK 64-Bit Server VM warning: Option UseConcMarkSweepGC was deprecated in version 9.0 and will likely be removed in a future release.
Could not rename log file '/var/log/cassandra/gc.log' to '/var/log/cassandra/gc.log.1' (Permission denied).
[0.000s][error][logging] Error opening log file '/var/log/cassandra/gc.log': Permission denied
[0.000s][error][logging] Initialization of output 'file:/var/log/cassandra/gc.log' using options 'filecount=10,fileSize=10485760' failed.
Invalid 'Xlog option 'Xlog:gc=info,heap*=trace,age*=debug,safepoint=info,promotion*=trace:file=/var/log/cassandra/gc.log:time,uptime,pid,tid,level:filecount=10,fileSize=10485760', see error log for detail
ls.
Error: Could not create the Java Virtual Machine.
Error: A fatal exception has occurred. Program will exit.
bmscsec@bmscsec-HP-Elite-Tower-800-G9-Desktop-PC:~$ sudo mkdir -p /var/log/cassandra
sudo chown -R $(whoami):$(whoami) /var/log/cassandra
bmscsec@bmscsec-HP-Elite-Tower-800-G9-Desktop-PC:~$ sudo chmod -R 777 /var/log/cassandra
bmscsec@bmscsec-HP-Elite-Tower-800-G9-Desktop-PC:~$ sudo systemctl restart cassandra

bmscsec@bmscsec-HP-Elite-Tower-800-G9-Desktop-PC:~$ ^[[200-sudo nano /etc/cassandra/cassandra-env.sh
sudo: command not found
bmscsec@bmscsec-HP-Elite-Tower-800-G9-Desktop-PC:~$ -sudo nano /etc/cassandra/cassandra-env.sh
Command '-sudo' not found, did you mean:
  command 'sudo' from deb sudo (1:9.9-1ubuntu2.4)
  command 'sudo' from deb sudo-ldap (1:9.9-1ubuntu2.4)
Try: sudo apt install <deb name>
bmscsec@bmscsec-HP-Elite-Tower-800-G9-Desktop-PC:~$ sudo nano /etc/cassandra/cassandra-env.sh
bmscsec@bmscsec-HP-Elite-Tower-800-G9-Desktop-PC:~$
bmscsec@bmscsec-HP-Elite-Tower-800-G9-Desktop-PC:~$ sudo systemctl restart cassandra

bmscsec@bmscsec-HP-Elite-Tower-800-G9-Desktop-PC:~$
bmscsec@bmscsec-HP-Elite-Tower-800-G9-Desktop-PC:~$ nodetool status
Datacenter: Datacenter1
=====
Status=Up/Down
// State=Normal/Leaving/Joining/Moving
-- Address Load Tokens Owns Host ID Rack
UN 127.0.0.1 300.28 KiB 16 7 rack1

Note: Non-system keyspaces don't have the same replication settings, effective ownership information is meaningless
bmscsec@bmscsec-HP-Elite-Tower-800-G9-Desktop-PC:~$ cqlsh
Connected to test cluster at 127.0.0.1:9042
[cqlsh 0.1.0 | Cassandra 4.1.0 | CQL spec 3.4.0 | Native protocol v5]
Use HELP for help.
cqlsh> CREATE KEYSPACE Students WITH REPLICATION =
... {'class': 'SimpleStrategy', 'replication_factor':1};

cqlsh>
cqlsh>
cqlsh>
```

```
cqlsh> USE Students;
cqlsh> BEGIN BATCH
... INSERT INTO Students_Info (Roll_No, StudName, DateOfJoining, Last_exan_Percent)
... VALUES (1, 'Asha', '2012-03-12', 79.9);
...
... INSERT INTO Students_Info (Roll_No, StudName, DateOfJoining, Last_exan_Percent)
... VALUES (2, 'Krian', '2012-03-12', 89.9);
...
... INSERT INTO Students_Info (Roll_No, StudName, DateOfJoining, Last_exan_Percent)
... VALUES (3, 'Tarun', '2012-03-12', 78.9);
...
... INSERT INTO Students_Info (Roll_No, StudName, DateOfJoining, Last_exan_Percent)
... VALUES (4, 'Sanrith', '2012-03-12', 90.9);
...
... INSERT INTO Students_Info (Roll_No, StudName, DateOfJoining, Last_exan_Percent)
... VALUES (5, 'Sinhha', '2012-03-12', 67.9);
...
... INSERT INTO Students_Info (Roll_No, StudName, DateOfJoining, Last_exan_Percent)
... VALUES (6, 'Rohan', '2012-03-12', 50.9);
... APPLY BATCH;
InvalidRequest: Error from server: code=2200 [Invalid query] message="table students_info does not exist"
cqlsh>
cqlsh> DESCRIBE KEYSPACES;
company students system_distributed system_views
emp system system_schema system_virtual_schema
emp_data system_auth system_traces

cqlsh> DESCRIBE TABLES;
cqlsh> DESCRIBE TABLES;
cqlsh> DESCRIBE TABLES;
cqlsh> CREATE TABLE Students_Info (
... Roll_No int PRIMARY KEY,
... StudName text,
... DateOfJoining timestamp,
... Last_exan_Percent double
... );
cqlsh> BEGIN BATCH
... INSERT INTO Students_Info (Roll_No, StudName, DateOfJoining, Last_exan_Percent)
... VALUES (1, 'Asha', '2012-03-12', 79.9);
...
... INSERT INTO Students_Info (Roll_No, StudName, DateOfJoining, Last_exan_Percent)
... VALUES (2, 'Krian', '2012-03-12', 89.9);
...
[1]+ Stopped cqlsh
bmscsec@bmscsec-HP-Elite-Tower-800-G9-Desktop-PC:~$ BEGIN BATCH
INSERT INTO Students_Info (Roll_No, StudName, DateOfJoining, Last_exan_Percent)
VALUES (1, 'Asha', '2012-03-12', 79.9);

INSERT INTO Students_Info (Roll_No, StudName, DateOfJoining, Last_exan_Percent)
VALUES (2, 'Krian', '2012-03-12', 89.9);

INSERT INTO Students_Info (Roll_No, StudName, DateOfJoining, Last_exan_Percent)
VALUES (3, 'Tarun', '2012-03-12', 78.9);

INSERT INTO Students_Info (Roll_No, StudName, DateOfJoining, Last_exan_Percent)
```



```

bmsecse@bmsecse-HP-Elite-Tower-800-G9-Desktop-PC: ~
cqlsh:students> BEGIN BATCH
... INSERT INTO Students_Info (Roll_No, StudName, DateOfJoining, Last_exan_Percent)
... VALUES (1, 'Asha', '2012-03-12', 79.9);
...
... INSERT INTO Students_Info (Roll_No, StudName, DateOfJoining, Last_exan_Percent)
... VALUES (2, 'Krian', '2012-03-12', 89.9);
...
... INSERT INTO Students_Info (Roll_No, StudName, DateOfJoining, Last_exan_Percent)
... VALUES (3, 'Tarun', '2012-03-12', 78.9);
...
... INSERT INTO Students_Info (Roll_No, StudName, DateOfJoining, Last_exan_Percent)
... VALUES (4, 'Sanrth', '2012-03-12', 90.9);
...
... INSERT INTO Students_Info (Roll_No, StudName, DateOfJoining, Last_exan_Percent)
... VALUES (5, 'Sinha', '2012-03-12', 67.9);
...
... INSERT INTO Students_Info (Roll_No, StudName, DateOfJoining, Last_exan_Percent)
... VALUES (6, 'Rohan', '2012-03-12', 56.9);
...
... APPLY BATCH;
cqlsh:students> SELECT*FROM students_info;

roll_no | dateofjoining | last_exan_percent | studname
-----|-----|-----|-----
5 | 2012-03-11 18:30:00.000000+0000 | 67.9 | Sinha
1 | 2012-03-11 18:30:00.000000+0000 | 79.9 | Asha
2 | 2012-03-11 18:30:00.000000+0000 | 89.9 | Krian
4 | 2012-03-11 18:30:00.000000+0000 | 90.9 | Sanrth
6 | 2012-03-11 18:30:00.000000+0000 | 56.9 | Rohan
3 | 2012-03-11 18:30:00.000000+0000 | 78.9 | Tarun

(6 rows)
cqlsh:students> SELECT*FROM student_info WHERE roll_no IN(1,2,3);
InvalidRequest: Error from server: code=2200 [Invalid query] message="table student_info does not exist"
cqlsh:students> SELECT*FROM student_info WHERE roll_no IN(1,2,3);
InvalidRequest: Error from server: code=2200 [Invalid query] message="table student_info does not exist"
cqlsh:students> SELECT*FROM students_info WHERE roll_no IN(1,2,3);

roll_no | dateofjoining | last_exan_percent | studname
-----|-----|-----|-----
1 | 2012-03-11 18:30:00.000000+0000 | 79.9 | Asha
2 | 2012-03-11 18:30:00.000000+0000 | 89.9 | Krian
3 | 2012-03-11 18:30:00.000000+0000 | 78.9 | Tarun

(3 rows)
cqlsh:students> SELECT*FROM students_info WHERE studname='Asha';
SyntaxException: line 1:47 no viable alternative at input '1' (SELECT*FROM students_info WHERE studname=('Ash'a'))
cqlsh:students> SELECT * FROM students_info WHERE studname = 'Asha';
InvalidRequest: Error from server: code=2200 [Invalid query] message="Cannot execute this query as it might involve data filtering and thus may have unpredictable performance. If you want to execute this query despite the performance unpredictability, use ALLOW FILTERING"
cqlsh:students> SELECT * FROM students_info WHERE studname = 'Asha' ALLOW FILTERING;

roll_no | dateofjoining | last_exan_percent | studname
-----|-----|-----|-----
1 | 2012-03-11 18:30:00.000000+0000 | 79.9 | Asha

(1 rows)
cqlsh:students> select roll_no as "USN" from students_info;

USN
---
5
1
2
4
6
3

(6 rows)
cqlsh:students> UPDATE roll_no as "USN" from students_info;
SyntaxException: line 1:15 mismatched input 'as' expecting K_SET (UPDATE roll_no [as] 'US...')
cqlsh:students> ALTER TABLE students_info RENAME roll_no TO USN;
cqlsh:students> select*from students_info;

usn | dateofjoining | last_exan_percent | studname
---|-----|-----|-----
5 | 2012-03-11 18:30:00.000000+0000 | 67.9 | Sinha
1 | 2012-03-11 18:30:00.000000+0000 | 79.9 | Asha
2 | 2012-03-11 18:30:00.000000+0000 | 89.9 | Krian
4 | 2012-03-11 18:30:00.000000+0000 | 90.9 | Sanrth
6 | 2012-03-11 18:30:00.000000+0000 | 56.9 | Rohan
3 | 2012-03-11 18:30:00.000000+0000 | 78.9 | Tarun

(6 rows)
cqlsh:students> UPDATE students_info SET studname='David Sheen' WHERE usn=2;
cqlsh:students> select*from students_info;

usn | dateofjoining | last_exan_percent | studname
---|-----|-----|-----
4 | 2012-03-11 18:30:00.000000+0000 | 67.9 | Sinha
1 | 2012-03-11 18:30:00.000000+0000 | 79.9 | Asha
2 | 2012-03-11 18:30:00.000000+0000 | 89.9 | David Sheen
4 | 2012-03-11 18:30:00.000000+0000 | 90.9 | Sanrth
6 | 2012-03-11 18:30:00.000000+0000 | 56.9 | Rohan
3 | 2012-03-11 18:30:00.000000+0000 | 78.9 | Tarun

(6 rows)
cqlsh:students> UPDATE students_info SET rollno=6 WHERE rollno=3;
InvalidRequest: Error from server: code=2200 [Invalid query] message="Undefined column name rollno in table students.students_info"
cqlsh:students> cqlsh:students> UPDATE students_info SET usn=6 WHERE usn=3;
[2]- Stopped cqlsh
bmsecse@bmsecse-HP-Elite-Tower-800-G9-Desktop-PC: ~$ cqlsh
Connected to Test Cluster at 127.0.0.1:9042
[cqlsh 6.1.0 | Cassandra 4.1.0 | CQL spec 3.4.0 | Native protocol v5]
Use HELP for help.
cqlsh> use students
...
...
...
cqlsh>

```

```

[3]+ Stopped                  cqlsh
bmsecse@bmsecse-HP-Elite-Tower-800-G9-Desktop-PC:~$ cqlsh
Connected to Test Cluster at 127.0.0.1:9042
[cqlsh 6.1.0 | Cassandra 4.1.0 | CQL spec 3.4.0 | Native protocol v4]
Use HELP for help.
cqlsh> USE students;
cqlsh:students> SELECT * FROM students_info;

   usn | dateofjoining | last_exam_percent | studname
-----|-----|-----|-----
  5 | 2012-03-11 18:30:00.000000+0000 | 67.9 | Sinha
  1 | 2012-03-11 18:30:00.000000+0000 | 79.9 | Asha
  2 | 2012-03-11 18:30:00.000000+0000 | 89.9 | David Sheen
  4 | 2012-03-11 18:30:00.000000+0000 | 90.9 | Sanrth
  6 | 2012-03-11 18:30:00.000000+0000 | 56.9 | Rohan
  3 | 2012-03-11 18:30:00.000000+0000 | 78.9 | Tarun

(6 rows)
cqlsh:students> UPDATE students_info SET usn=6 WHERE usn=3;
InvalidRequest: Error from server: code=2200 [Invalid query] message="PRIMARY KEY part usn found in SET part"
cqlsh:students> UPDATE students_info SET usn = 10 WHERE roll_no = 3;
InvalidRequest: Error from server: code=2200 [Invalid query] message="PRIMARY KEY part usn found in SET part"
cqlsh:students> UPDATE students_info SET usn = 6 WHERE usn= 3;
InvalidRequest: Error from server: code=2200 [Invalid query] message="PRIMARY KEY part usn found in SET part"
cqlsh:students> DELETE FROM students_info WHERE usn = 3;
cqlsh:students> INSERT INTO students_info (usn, dateofjoining, last_exam_percent, studname)
... VALUES (6, '2012-03-11 18:30:00', 78.9, 'Tarun');
cqlsh:students> select*FROM STUDENTS_INFO;

   usn | dateofjoining | last_exam_percent | studname
-----|-----|-----|-----
  5 | 2012-03-11 18:30:00.000000+0000 | 67.9 | Sinha
  1 | 2012-03-11 18:30:00.000000+0000 | 79.9 | Asha
  2 | 2012-03-11 18:30:00.000000+0000 | 89.9 | David Sheen
  4 | 2012-03-11 18:30:00.000000+0000 | 90.9 | Sanrth
  6 | 2012-03-11 13:00:00.000000+0000 | 78.9 | Tarun

(5 rows)
cqlsh:students> DELETE last_exam_percent from students_info where usn=2;
cqlsh:students> select*from students_info;

   usn | dateofjoining | last_exam_percent | studname
-----|-----|-----|-----
  5 | 2012-03-11 18:30:00.000000+0000 | 67.9 | Sinha
  1 | 2012-03-11 18:30:00.000000+0000 | 79.9 | Asha
  2 | 2012-03-11 18:30:00.000000+0000 | null | David Sheen
  4 | 2012-03-11 18:30:00.000000+0000 | 90.9 | Sanrth
  6 | 2012-03-11 13:00:00.000000+0000 | 78.9 | Tarun

(5 rows)
cqlsh:students> DELETE FROM students_info where usn=2;
cqlsh:students> SELECT*FROM STUDENTS_INFO;

   usn | dateofjoining | last_exam_percent | studname
-----|-----|-----|-----
  5 | 2012-03-11 18:30:00.000000+0000 | 67.9 | Sinha

(1 rows)
cqlsh:students> DELETE last_exam_percent from students_info where usn=2;
cqlsh:students> select*from students_info;

   usn | dateofjoining | last_exam_percent | studname
-----|-----|-----|-----
  5 | 2012-03-11 18:30:00.000000+0000 | 67.9 | Sinha
  1 | 2012-03-11 18:30:00.000000+0000 | 79.9 | Asha
  2 | 2012-03-11 18:30:00.000000+0000 | null | David Sheen
  4 | 2012-03-11 18:30:00.000000+0000 | 90.9 | Sanrth
  6 | 2012-03-11 13:00:00.000000+0000 | 78.9 | Tarun

(5 rows)
cqlsh:students> DELETE FROM students_info where usn=2;
cqlsh:students> SELECT*FROM STUDENTS_INFO;

   usn | dateofjoining | last_exam_percent | studname
-----|-----|-----|-----
  5 | 2012-03-11 18:30:00.000000+0000 | 67.9 | Sinha
  1 | 2012-03-11 18:30:00.000000+0000 | 79.9 | Asha
  4 | 2012-03-11 18:30:00.000000+0000 | 90.9 | Sanrth
  6 | 2012-03-11 13:00:00.000000+0000 | 78.9 | Tarun

(4 rows)
cqlsh:students> ALTER TABLE STUDENTS_INFO ADD hobbies set<text>;
...
...
...
...
Statements are terminated with a ';'. You can press CTRL-C to cancel an incomplete statement.
cqlsh:students> ALTER TABLE STUDENTS_INFO ADD hobbies set<text>;
...
cqlsh:students> ALTER TABLE STUDENTS_INFO ADD hobbies set<text>;
cqlsh:students>
cqlsh:students> select*from students_info;

   usn | dateofjoining | hobbies | last_exam_percent | studname
-----|-----|-----|-----|-----
  5 | 2012-03-11 18:30:00.000000+0000 | null | 67.9 | Sinha
  1 | 2012-03-11 18:30:00.000000+0000 | null | 79.9 | Asha
  4 | 2012-03-11 18:30:00.000000+0000 | null | 90.9 | Sanrth
  6 | 2012-03-11 13:00:00.000000+0000 | null | 78.9 | Tarun

(4 rows)
cqlsh:students> ALTER TABLE students_info ADD language list<text>;
cqlsh:students> select*from students_info;

   usn | dateofjoining | hobbies | language | last_exam_percent | studname
-----|-----|-----|-----|-----|-----
  5 | 2012-03-11 18:30:00.000000+0000 | null | null | 67.9 | Sinha
  1 | 2012-03-11 18:30:00.000000+0000 | null | null | 79.9 | Asha
  4 | 2012-03-11 18:30:00.000000+0000 | null | null | 90.9 | Sanrth
  6 | 2012-03-11 13:00:00.000000+0000 | null | null | 78.9 | Tarun

(4 rows)

```

```

cqlsh:students>
cqlsh:students> select*from students_info;

usrn | dateofjoining | hobbies | last_exan_percent | studname
-----|-----|-----|-----|-----
5 | 2012-03-11 18:30:00.000000+0000 | null | 67.9 | Sinha
1 | 2012-03-11 18:30:00.000000+0000 | null | 79.9 | Asha
4 | 2012-03-11 18:30:00.000000+0000 | null | 90.9 | Sanrth
6 | 2012-03-11 13:00:00.000000+0000 | null | 78.9 | Tarun
(4 rows)

cqlsh:students> ALTER TABLE students_info ADD language list<text>;
cqlsh:students> select*from students_info;

usrn | dateofjoining | hobbies | language | last_exan_percent | studname
-----|-----|-----|-----|-----|-----
5 | 2012-03-11 18:30:00.000000+0000 | null | null | 67.9 | Sinha
1 | 2012-03-11 18:30:00.000000+0000 | null | null | 79.9 | Asha
4 | 2012-03-11 18:30:00.000000+0000 | null | null | 90.9 | Sanrth
6 | 2012-03-11 13:00:00.000000+0000 | null | null | 78.9 | Tarun
(4 rows)

cqlsh:students> UPDATE students_info SET hobbies=hobbies+('Chess,Table Tennis') where usrn=1;
cqlsh:students> SELECT*FROM STUDENTS_INFO;

usrn | dateofjoining | hobbies | language | last_exan_percent | studname
-----|-----|-----|-----|-----|-----
5 | 2012-03-11 18:30:00.000000+0000 | null | null | 67.9 | Sinha
1 | 2012-03-11 18:30:00.000000+0000 | ('Chess,Table Tennis') | null | 79.9 | Asha
4 | 2012-03-11 18:30:00.000000+0000 | null | null | 90.9 | Sanrth
6 | 2012-03-11 13:00:00.000000+0000 | null | null | 78.9 | Tarun
(4 rows)

cqlsh:students> SELECT*FROM student_info where usrn=1;
InvalidRequest: Error from server: code=2200 [Invalid query] message="table student_info does not exist"
cqlsh:students> select*from students_info where usrn=1;

usrn | dateofjoining | hobbies | language | last_exan_percent | studname
-----|-----|-----|-----|-----|-----
1 | 2012-03-11 18:30:00.000000+0000 | ('Chess,Table Tennis') | null | 79.9 | Asha
(1 rows)

cqlsh:students> update students_info set language=language+('Hindi,English') where usrn=1;
cqlsh:students> select*from students_info;

usrn | dateofjoining | hobbies | language | last_exan_percent | studname
-----|-----|-----|-----|-----|-----
5 | 2012-03-11 18:30:00.000000+0000 | null | null | 67.9 | Sinha
1 | 2012-03-11 18:30:00.000000+0000 | ('Chess,Table Tennis') | ('Hindi,English') | 79.9 | Asha
4 | 2012-03-11 18:30:00.000000+0000 | null | null | 90.9 | Sanrth
6 | 2012-03-11 13:00:00.000000+0000 | null | null | 78.9 | Tarun
(4 rows)

cqlsh:students>

```

```

Activities Terminal Apr 1 14:40
bmscscse@bmscscse-HP-Elite-Tower-800-G9-Desktop-PC: ~

(4 rows)
cqlsh:students> CREATE TABLE library_book (
... counter_value counter,
... book_name varchar,
... stud_name varchar,
... PRIMARY KEY (book_name, stud_name)
... );
cqlsh:students> SELECT*FROM LIBRARY_BOOK;

book_name | stud_name | counter_value
-----|-----|-----
(0 rows)

cqlsh:students> BEGIN BATCH
... -- Inserting 4 records into library_book table
... INSERT INTO library_book (counter_value, book_name, stud_name)
... VALUES (0, 'Big Data Analytics', 'Jeet');
...
... INSERT INTO library_book (counter_value, book_name, stud_name)
... VALUES (0, 'Machine Learning', 'Anit');
...
... INSERT INTO library_book (counter_value, book_name, stud_name)
... VALUES (0, 'Data Science Essentials', 'Ravi');
...
... INSERT INTO library_book (counter_value, book_name, stud_name)
... VALUES (0, 'Cloud Computing', 'Priya');
InvalidRequest: Error from server: code=2200 [Invalid query] message="INSERT statements are not allowed on counter tables, use UPDATE instead"
cqlsh:students> BEGIN BATCH
... -- Inserting 4 records into library_book table with counter initialized to 0
... INSERT INTO library_book (counter_value, book_name, stud_name)
... VALUES (0, 'Big Data Analytics', 'Jeet');
...
... INSERT INTO library_book (counter_value, book_name, stud_name)
... VALUES (0, 'Machine Learning', 'Anit');
...
... INSERT INTO library_book (counter_value, book_name, stud_name)
... VALUES (0, 'Data Science Essentials', 'Ravi');
...
... INSERT INTO library_book (counter_value, book_name, stud_name)
... VALUES (0, 'Cloud Computing', 'Priya');
InvalidRequest: Error from server: code=2200 [Invalid query] message="INSERT statements are not allowed on counter tables, use UPDATE instead"
cqlsh:students> -- Load data into the counter column
cqlsh:students> UPDATE library_book
... SET counter_value = counter_value + 1
... WHERE book_name = 'Big data Analytics'
... AND stud_name = 'jeet';
cqlsh:students> SELECT*FROM library_book;

book_name | stud_name | counter_value
-----|-----|-----
Big data Analytics | jeet | 1
(1 rows)

```

```
Activities Terminal Apr 1 14:41 bmscscse@bmscscse-HP-Elite-Tower-800-G9-Desktop-PC: ~  
... INSERT INTO library_book (counter_value, book_name, stud_name)  
... VALUES (0, 'Data Science Essentials', 'Ravi');  
...  
... INSERT INTO library_book (counter_value, book_name, stud_name)  
... VALUES (0, 'Cloud Computing', 'Priya');  
... APPLY BATCH;  
InvalidRequest: Error from server: code=5200 [Invalid query] message="INSERT statements are not allowed on counter tables, use UPDATE instead"  
cqlsh:students> -- Load data into the counter column  
cqlsh:students> UPDATE library_book  
... SET counter_value = counter_value + 1  
... WHERE book_name = 'Big data Analytics'  
... AND stud_name = 'jeet';  
cqlsh:students> SELECT * FROM library_book;  


| book_name          | stud_name | counter_value |
|--------------------|-----------|---------------|
| Big data Analytics | jeet      | 1             |

  
(1 rows)  
cqlsh:students> -- TIME TO LIVE (TTL) Example  
cqlsh:students> CREATE TABLE userlogin (  
... user_id int PRIMARY KEY,  
... password text  
... );  
cqlsh:students>  
cqlsh:students>  
cqlsh:students> SELECT * FROM userlogin;  


| user_id | password |
|---------|----------|
|---------|----------|

  
(0 rows)  
cqlsh:students> INSERT INTO userlogin(user_id, password)  
... VALUES (1, 'luffy')  
... USING TTL 30;  
cqlsh:students> SELECT * FROM userlogin;  


| user_id | password |
|---------|----------|
| 1       | luffy    |

  
(1 rows)  
cqlsh:students> SELECT TTL(password)  
... FROM userlogin  
... WHERE user_id = 1;  


| ttr(password) |
|---------------|
|---------------|

  
(0 rows)  
cqlsh:students>
```



```
Activities Terminal Apr 1 14:49 bmsccse@bmsccse-HP-Elite-Tower-800-G9-Desktop-PC: ~

CREATE TABLE students.elearninglists (
  id uuid PRIMARY KEY,
  course_id uuid,
  course_order int,
  courseowner text,
  title text
) WITH additional_write_policy = '99p'
AND bloom_filter_fp_chance = 0.01
AND caching = {'keys': 'ALL', 'rows_per_partition': 'NONE'}
AND cdc = false
AND comment = ''
AND compaction = {'class': 'org.apache.cassandra.db.compaction.SizeTieredCompactionStrategy', 'max_threshold': '32', 'min_threshold': '4'}
AND compression = {'chunk_length_in_kb': '16', 'class': 'org.apache.cassandra.io.compress.LZ4Compressor'}
AND dclocal_read_repair = 'BLOCKING'
AND default_time_to_live = 0
AND extensions = {}
AND gc_grace_seconds = 864000
AND max_index_interval = 2048
AND memtable_flush_period_in_ms = 0
AND min_index_interval = 128
AND read_repair = 'BLOCKING'
AND speculative_retry = '99p';
cqlsh:students> COPY elearninglists(id, course_order, course_id, courseowner, title)
... TO 'd:/elearninglists.csv' WITH HEADER = TRUE;
Using 16 child processes
Can't open 'd:/elearninglists.csv' for writing: [Errno 2] No such file or directory: 'd:/elearninglists.csv'
cqlsh:students> COPY elearninglists(id, course_order, course_id, courseowner, title)
... TO '/home/bmsccse/Desktop/elearninglists.csv' WITH HEADER = TRUE;
Using 16 child processes
Starting copy of students.elearninglists with columns [id, course_order, course_id, courseowner, title].
Processed: 0 rows; Rate: 0 rows/s; Avg. rate: 0 rows/s
0 rows exported to 1 files in 0.133 seconds.
cqlsh:students> BEGIN BATCH
... INSERT INTO elearninglists (id, course_order, course_id, courseowner, title)
... VALUES (uuid(), 1, uuid(), 'John Doe', 'Introduction to Python');
... INSERT INTO elearninglists (id, course_order, course_id, courseowner, title)
... VALUES (uuid(), 2, uuid(), 'Jane Smith', 'Machine Learning Basics');
... INSERT INTO elearninglists (id, course_order, course_id, courseowner, title)
... VALUES (uuid(), 3, uuid(), 'Alice Johnson', 'Data Structures and Algorithms');
... APPLY BATCH;
cqlsh:students> select*from elearninglists;
+-----+-----+-----+-----+-----+
| id | course_id | course_order | courseowner | title |
+-----+-----+-----+-----+-----+
| f6cb9530-33c6-4619-b26a-7f1f14d75290 | 10cde27-2d81-4481-ae73-d156a2b01a0f | 3 | Alice Johnson | Data Structures and Algorithms |
| f9892ac5-ae62-4156-a496-2f62e0f77962 | 3b58b8f5-55bb-4726-9460-786a98b38883 | 2 | Jane Smith | Machine Learning Basics |
| 59ca570d-0253-440b-b137-003f51ff31f1 | 5d5c2b42-97ef-483d-97ba-33e201ef7b36 | 1 | John Doe | Introduction to Python |
+-----+-----+-----+-----+-----+
(3 rows)
cqlsh:students> COPY elearninglists(id, course_order, course_id, courseowner, title)
... TO '/home/bmsccse/Desktop/elearninglists.csv' WITH HEADER = TRUE;
Using 16 child processes
Can't open 'd:/elearninglists.csv' for writing: [Errno 2] No such file or directory: 'd:/elearninglists.csv'
cqlsh:students> COPY elearninglists(id, course_order, course_id, courseowner, title)
... TO '/home/bmsccse/Desktop/elearninglists.csv' WITH HEADER = TRUE;
Using 16 child processes
Starting copy of students.elearninglists with columns [id, course_order, course_id, courseowner, title].
Processed: 0 rows; Rate: 0 rows/s; Avg. rate: 0 rows/s
0 rows exported to 1 files in 0.133 seconds.
cqlsh:students> BEGIN BATCH
... INSERT INTO elearninglists (id, course_order, course_id, courseowner, title)
... VALUES (uuid(), 1, uuid(), 'John Doe', 'Introduction to Python');
... INSERT INTO elearninglists (id, course_order, course_id, courseowner, title)
... VALUES (uuid(), 2, uuid(), 'Jane Smith', 'Machine Learning Basics');
... INSERT INTO elearninglists (id, course_order, course_id, courseowner, title)
... VALUES (uuid(), 3, uuid(), 'Alice Johnson', 'Data Structures and Algorithms');
... APPLY BATCH;
cqlsh:students> select*from elearninglists;
+-----+-----+-----+-----+-----+
| id | course_id | course_order | courseowner | title |
+-----+-----+-----+-----+-----+
| f6cb9530-33c6-4619-b26a-7f1f14d75290 | 10cde27-2d81-4481-ae73-d156a2b01a0f | 3 | Alice Johnson | Data Structures and Algorithms |
| f9892ac5-ae62-4156-a496-2f62e0f77962 | 3b58b8f5-55bb-4726-9460-786a98b38883 | 2 | Jane Smith | Machine Learning Basics |
| 59ca570d-0253-440b-b137-003f51ff31f1 | 5d5c2b42-97ef-483d-97ba-33e201ef7b36 | 1 | John Doe | Introduction to Python |
+-----+-----+-----+-----+-----+
(3 rows)
cqlsh:students> COPY elearninglists(id, course_order, course_id, courseowner, title)
... TO '/home/bmsccse/Desktop/elearninglists.csv' WITH HEADER = TRUE;
Using 16 child processes
Starting copy of students.elearninglists with columns [id, course_order, course_id, courseowner, title].
Processed: 3 rows; Rate: 4 rows/s; Avg. rate: 7 rows/s
3 rows exported to 1 files in 0.452 seconds (0 skipped).
cqlsh:students> cat persons.csv | cqlsh -e "COPY persons(id, fname, lname) FROM STDIN WITH HEADER = TRUE;"
Invalid syntax at char 17
cat persons.csv | cqlsh -e "COPY persons(id, fname, lname) FROM STDIN WITH HEADER = TRUE;"
cqlsh:students> COPY elearninglists(id, course_order, course_id, courseowner, title)
... TO STDOUT WITH HEADER = TRUE;
cqlsh:students>
+-----+-----+-----+-----+-----+
| id | course_order | course_id | courseowner | title |
+-----+-----+-----+-----+-----+
| f6cb9530-33c6-4619-b26a-7f1f14d75290,3,10cde27-2d81-4481-ae73-d156a2b01a0f,Alice Johnson,Data Structures and Algorithms |
+-----+-----+-----+-----+-----+
cqlsh:students>
```

## EXPERIMENT-5

Execution of HDFS Commands for interaction with Hadoop Environment. (Minimum 10 commands to be executed)

```
hadoop@bmsciece-HP-Ellite-Tower-800-G3-Desktop-PC: ~/hadoop/sbin
version          print the version

Daemon Commands:
kms              run KMS, the Key Management Server
registrydms     run the registry DNS server

SUBCOMMAND may print help when invoked w/o parameters or with -h.
hadoop@bmsciece-HP-Ellite-Tower-800-G3-Desktop-PC: ~/hadoop/sbin
hadoop@bmsciece-HP-Ellite-Tower-800-G3-Desktop-PC: $ which hadoop
/home/hadoop/hadoop/bin/hadoop
hadoop@bmsciece-HP-Ellite-Tower-800-G3-Desktop-PC: $ cd /home/hadoop/hadoop/sbin
hadoop@bmsciece-HP-Ellite-Tower-800-G3-Desktop-PC: ~/hadoop/sbin$ ./start-dfs.sh
Starting namenodes on [localhost]
localhost: namenode is running as process 9900. Stop it first and ensure /tmp/hadoop-hadoop-namenode.pid file is empty before retry.
Starting datanodes
localhost: datanode is running as process 10056. Stop it first and ensure /tmp/hadoop-hadoop-datanode.pid file is empty before retry.
Starting secondary namenodes [bmsciece-HP-Ellite-Tower-800-G3-Desktop-PC]
bmsciece-HP-Ellite-Tower-800-G3-Desktop-PC: secondarynamenode is running as process 10334. Stop it first and ensure /tmp/hadoop-hadoop-secondarynamenode.pid file is empty before retry.
hadoop@bmsciece-HP-Ellite-Tower-800-G3-Desktop-PC: ~/hadoop/sbin$ ./start-yarn.sh
Starting resource manager
resource manager is running as process 10644. Stop it first and ensure /tmp/hadoop-hadoop-resource manager.pid file is empty before retry.
Starting node managers
localhost: nodemanager is running as process 10807. Stop it first and ensure /tmp/hadoop-hadoop-nodemanager.pid file is empty before retry.
hadoop@bmsciece-HP-Ellite-Tower-800-G3-Desktop-PC: ~/hadoop/sbin$ jps
10644 ResourceManager
10807 NodeManager
10056 DataNode
9900 NameNode
10334 SecondaryNameNode
16111 Jps
hadoop@bmsciece-HP-Ellite-Tower-800-G3-Desktop-PC: ~/hadoop/sbin$ hdfs dfs -mkdir /abc
mkdir: '/abc': File exists
hadoop@bmsciece-HP-Ellite-Tower-800-G3-Desktop-PC: ~/hadoop/sbin$ hdfs dfs -rm -r /abc
Deleted /abc
hadoop@bmsciece-HP-Ellite-Tower-800-G3-Desktop-PC: ~/hadoop/sbin$ hdfs dfs -mkdir /abc
hadoop@bmsciece-HP-Ellite-Tower-800-G3-Desktop-PC: ~/hadoop/sbin$ hdfs dfs -ls /
Found 4 items
drwxr-xr-x 1 hadoop supergroup 0 2025-04-15 14:32 /abc
drwxr-xr-x 1 hadoop supergroup 0 2024-05-13 15:12 /homedataflair
drwxr-xr-x 1 hadoop supergroup 0 2024-05-13 14:30 /olddata
drwxr-xr-x 1 hadoop supergroup 0 2025-04-15 14:19 /rgs
hadoop@bmsciece-HP-Ellite-Tower-800-G3-Desktop-PC: ~/hadoop/sbin$ hdfs dfs -put /home/hadoop/Desktop/file1.txt /abc/file1.txt
hadoop@bmsciece-HP-Ellite-Tower-800-G3-Desktop-PC: ~/hadoop/sbin$ hdfs dfs -cat /abc/file1.txt
ht how are you
how is your job
how is your family
how is your brother
how is your sister
hadoop@bmsciece-HP-Ellite-Tower-800-G3-Desktop-PC: ~/hadoop/sbin$ hdfs dfs -copyfromLocal /home/hadoop/Desktop/file1.txt /abc/file1.txt
copyfromLocal: '/abc/file1.txt': File exists
hadoop@bmsciece-HP-Ellite-Tower-800-G3-Desktop-PC: ~/hadoop/sbin$ hdfs dfs -copyfromLocal /home/hadoop/Desktop/file1.txt /abc/file2.txt
hadoop@bmsciece-HP-Ellite-Tower-800-G3-Desktop-PC: ~/hadoop/sbin$ hdfs dfs -ls /abc
Found 2 items
-rw-r--r-- 1 hadoop supergroup 90 2025-04-15 14:34 /abc/file1.txt
-rw-r--r-- 1 hadoop supergroup 90 2025-04-15 14:35 /abc/file2.txt
hadoop@bmsciece-HP-Ellite-Tower-800-G3-Desktop-PC: ~/hadoop/sbin$ hdfs dfs -cat /abc/file2.txt
ht how are you
how is your job
how is your family
how is your brother
how is your sister
hadoop@bmsciece-HP-Ellite-Tower-800-G3-Desktop-PC: ~/hadoop/sbin$ hdfs dfs -get /abc/WC.txt /home/hduser/Downloads/WC.txt
get: '/home/hduser/Downloads/WC.txt': No such file or directory: 'file:///home/hduser/Downloads/WC.txt'
hadoop@bmsciece-HP-Ellite-Tower-800-G3-Desktop-PC: ~/hadoop/sbin$ ls /home/hduser/Downloads/
ls: cannot access '/home/hduser/Downloads/': No such file or directory
hadoop@bmsciece-HP-Ellite-Tower-800-G3-Desktop-PC: ~/hadoop/sbin$ mkdir -p /home/hduser/Downloads/
mkdir: cannot create directory '/home/hduser/': Permission denied
hadoop@bmsciece-HP-Ellite-Tower-800-G3-Desktop-PC: ~/hadoop/sbin$ sudo mkdir -p /home/hduser/Downloads/
[sudo] password for hadoop:
hadoop@bmsciece-HP-Ellite-Tower-800-G3-Desktop-PC: ~/hadoop/sbin$ ls /home/hduser/
Downloads
hadoop@bmsciece-HP-Ellite-Tower-800-G3-Desktop-PC: ~/hadoop/sbin$ hdfs dfs -get /abc/WC.txt /home/hduser/Downloads/WC.txt
get: '/abc/WC.txt': No such file or directory
hadoop@bmsciece-HP-Ellite-Tower-800-G3-Desktop-PC: ~/hadoop/sbin$ hdfs dfs -get /abc/WC.txt /home/hduser/Downloads/WC.txt
get: '/abc/WC.txt': No such file or directory
hadoop@bmsciece-HP-Ellite-Tower-800-G3-Desktop-PC: ~/hadoop/sbin$ hdfs dfs -ls /abc/
Found 2 items
-rw-r--r-- 1 hadoop supergroup 90 2025-04-15 14:34 /abc/file1.txt
-rw-r--r-- 1 hadoop supergroup 90 2025-04-15 14:35 /abc/file2.txt
hadoop@bmsciece-HP-Ellite-Tower-800-G3-Desktop-PC: ~/hadoop/sbin$ hdfs dfs -get /abc/file1.txt /home/hduser/Downloads/file1.txt
get: '/home/hduser/Downloads/file1.txt': No such file or directory
hadoop@bmsciece-HP-Ellite-Tower-800-G3-Desktop-PC: ~/hadoop/sbin$ ls -ld /home/hduser/Downloads/
drwxr-xr-x 2 root root 4096 Apr 15 14:37 /home/hduser/Downloads/
hadoop@bmsciece-HP-Ellite-Tower-800-G3-Desktop-PC: ~/hadoop/sbin$ sudo chown -R hadoop:hadoop /home/hduser/Downloads/
hadoop@bmsciece-HP-Ellite-Tower-800-G3-Desktop-PC: ~/hadoop/sbin$ ls -ld /home/hduser/Downloads/
drwxr-xr-x 2 hadoop hadoop 4096 Apr 15 14:37 /home/hduser/Downloads/
hadoop@bmsciece-HP-Ellite-Tower-800-G3-Desktop-PC: ~/hadoop/sbin$ hdfs dfs -get /abc/file1.txt /home/hduser/Downloads/file1.txt
hadoop@bmsciece-HP-Ellite-Tower-800-G3-Desktop-PC: ~/hadoop/sbin$ hdfs dfs -getmerge /abc/WC.txt /abc/WC2.txt /home/hduser/Desktop/Merge.txt
getmerge: '/abc/WC2.txt': No such file or directory
getmerge: '/abc/WC.txt': No such file or directory
hadoop@bmsciece-HP-Ellite-Tower-800-G3-Desktop-PC: ~/hadoop/sbin$ hdfs dfs -getmerge /abc/file1.txt /abc/file2.txt /home/hduser/Desktop/Merge.txt
getmerge: Mkdir failed to create file:/home/hduser/Desktop (exit=1): file:/home/hduser/Desktop (exit=1): file:/home/hduser/Desktop
hadoop@bmsciece-HP-Ellite-Tower-800-G3-Desktop-PC: ~/hadoop/sbin$ sudo mkdir -p /home/hduser/Desktop
hadoop@bmsciece-HP-Ellite-Tower-800-G3-Desktop-PC: ~/hadoop/sbin$ sudo chown -R hadoop:hadoop /home/hduser/Desktop
hadoop@bmsciece-HP-Ellite-Tower-800-G3-Desktop-PC: ~/hadoop/sbin$ hdfs dfs -getmerge /abc/file1.txt /abc/file2.txt /home/hduser/Desktop/Merge.txt
hadoop@bmsciece-HP-Ellite-Tower-800-G3-Desktop-PC: ~/hadoop/sbin$ cat /home/hduser/Desktop/Merge.txt
ht how are you
how is your job
how is your family
how is your brother
how is your sister
```



```
hadoop@bmsccse-HP-Elite-Tower-800-G9-Desktop-PC: ~/hadoop/sbin
hadoop@bmsccse-HP-Elite-Tower-800-G9-Desktop-PC:~/hadoop/sbin$ hdfs dfs -cat /abc/file2.txt
ht how are you
how is your job
how is your family
how is your brother
how is your sister

hadoop@bmsccse-HP-Elite-Tower-800-G9-Desktop-PC:~/hadoop/sbin$ hdfs dfs -get /abc/WC.txt /home/hduser/Downloads/WC.txt
get: /home/hduser/Downloads/WC.txt: No such file or directory: "file:///home/hduser/Downloads/WC.txt"
hadoop@bmsccse-HP-Elite-Tower-800-G9-Desktop-PC:~/hadoop/sbin$ ls /home/hduser/Downloads/
ls: cannot access '/home/hduser/Downloads/': No such file or directory
hadoop@bmsccse-HP-Elite-Tower-800-G9-Desktop-PC:~/hadoop/sbin$ mkdir -p /home/hduser/Downloads/
mkdir: cannot create directory '/home/hduser/': Permission denied
hadoop@bmsccse-HP-Elite-Tower-800-G9-Desktop-PC:~/hadoop/sbin$ sudo mkdir -p /home/hduser/Downloads/
[sudo] password for hadoop:
hadoop@bmsccse-HP-Elite-Tower-800-G9-Desktop-PC:~/hadoop/sbin$ ls /home/hduser/
Downloads
hadoop@bmsccse-HP-Elite-Tower-800-G9-Desktop-PC:~/hadoop/sbin$ hdfs dfs -get /abc/WC.txt /home/hduser/Downloads/WC.txt
get: /abc/WC.txt: No such file or directory
hadoop@bmsccse-HP-Elite-Tower-800-G9-Desktop-PC:~/hadoop/sbin$ hdfs dfs -get /abc/WC.txt /home/hduser/Downloads/WC.txt
get: /abc/WC.txt: No such file or directory
hadoop@bmsccse-HP-Elite-Tower-800-G9-Desktop-PC:~/hadoop/sbin$ hdfs dfs -ls /abc/
Found 2 items
-rw-r--r-- 1 hadoop supergroup          90 2025-04-15 14:34 /abc/file1.txt
-rw-r--r-- 1 hadoop supergroup          90 2025-04-15 14:35 /abc/file2.txt
hadoop@bmsccse-HP-Elite-Tower-800-G9-Desktop-PC:~/hadoop/sbin$ hdfs dfs -get /abc/file1.txt /home/hduser/Downloads/file1.txt
get: /home/hduser/Downloads/file1.txt: COPIING (Permission denied)
hadoop@bmsccse-HP-Elite-Tower-800-G9-Desktop-PC:~/hadoop/sbin$ ls -ld /home/hduser/Downloads/
drwx-r-xr-x 2 root root 4096 Apr 15 14:37 /home/hduser/Downloads/
hadoop@bmsccse-HP-Elite-Tower-800-G9-Desktop-PC:~/hadoop/sbin$ sudo chown -R hadoop:hadoop /home/hduser/Downloads/
hadoop@bmsccse-HP-Elite-Tower-800-G9-Desktop-PC:~/hadoop/sbin$ ls -ld /home/hduser/Downloads/
drwx-r-xr-x 2 hadoop hadoop 4096 Apr 15 14:37 /home/hduser/Downloads/
hadoop@bmsccse-HP-Elite-Tower-800-G9-Desktop-PC:~/hadoop/sbin$ hdfs dfs -get /abc/file1.txt /home/hduser/Downloads/file1.txt
hadoop@bmsccse-HP-Elite-Tower-800-G9-Desktop-PC:~/hadoop/sbin$ hdfs dfs -getmerge /abc/WC.txt /abc/WC2.txt /home/hduser/Desktop/Merge.txt
getmerge: /abc/WC.txt: No such file or directory
getmerge: /abc/WC2.txt: No such file or directory
hadoop@bmsccse-HP-Elite-Tower-800-G9-Desktop-PC:~/hadoop/sbin$ hdfs dfs -getmerge /abc/file1.txt /abc/file2.txt /home/hduser/Desktop/Merge.txt
getmerge: Mkdirs failed to create file:/home/hduser/Desktop (exists=false, cwd=file:/home/hadoop/hadoop/sbin)
hadoop@bmsccse-HP-Elite-Tower-800-G9-Desktop-PC:~/hadoop/sbin$ sudo mkdir -p /home/hduser/Desktop
hadoop@bmsccse-HP-Elite-Tower-800-G9-Desktop-PC:~/hadoop/sbin$ sudo chown -R hadoop:hadoop /home/hduser/Desktop
hadoop@bmsccse-HP-Elite-Tower-800-G9-Desktop-PC:~/hadoop/sbin$ hdfs dfs -getmerge /abc/file1.txt /abc/file2.txt /home/hduser/Desktop/Merge.txt
hadoop@bmsccse-HP-Elite-Tower-800-G9-Desktop-PC:~/hadoop/sbin$ cat /home/hduser/Desktop/Merge.txt
ht how are you
how is your job
how is your family
how is your brother
how is your sister

ht how are you
how is your job
how is your family
how is your brother
how is your sister

hadoop@bmsccse-HP-Elite-Tower-800-G9-Desktop-PC:~/hadoop/sbin$ hdfs dfs -copyToLocal /abc/WC.txt /home/hduser/Desktop
copyToLocal: /abc/WC.txt: No such file or directory
hadoop@bmsccse-HP-Elite-Tower-800-G9-Desktop-PC:~/hadoop/sbin$ hdfs dfs -copyToLocal /abc/file1.txt /home/hduser/Desktop
hadoop@bmsccse-HP-Elite-Tower-800-G9-Desktop-PC:~/hadoop/sbin$ cat /home/hduser/Desktop/file1.txt
ht how are you
how is your job
how is your family
how is your brother
how is your sister

hadoop@bmsccse-HP-Elite-Tower-800-G9-Desktop-PC:~/hadoop/sbin$ hdfs dfs -cat /abc/file1.txt
ht how are you
how is your job
how is your family
how is your brother
how is your sister

hadoop@bmsccse-HP-Elite-Tower-800-G9-Desktop-PC:~/hadoop/sbin$ hdfs dfs -cat /abc/file2.txt
ht how are you
how is your job
how is your family
how is your brother
how is your sister

hadoop@bmsccse-HP-Elite-Tower-800-G9-Desktop-PC:~/hadoop/sbin$ hdfs dfs -mv /abc/FFF
hadoop@bmsccse-HP-Elite-Tower-800-G9-Desktop-PC:~/hadoop/sbin$ hdfs dfs -ls /FFF
Found 2 items
-rw-r--r-- 1 hadoop supergroup          90 2025-04-15 14:34 /FFF/file1.txt
-rw-r--r-- 1 hadoop supergroup          90 2025-04-15 14:35 /FFF/file2.txt
hadoop@bmsccse-HP-Elite-Tower-800-G9-Desktop-PC:~/hadoop/sbin$ 9. Cp Hadoop HDFS
cp Command Usage cp Hadoop HDFS
cp Command Example
hadoop fs -cp /CSE/ /LLL
hadoop fs -ls /LLL
The cp command copies a file from one directory to another directory within the
HDFS.
9.: command not found
cp: target 'HDFS' is not a directory
cp: cannot stat 'Command': No such file or directory
cp: /CSE/: No such file or directory
ls: /LLL: No such file or directory
Command 'The' not found, did you mean:
  Command 'he' from deb node-he (1.2.0-3)
  Command 'the' from deb the (3.3-rc1-bu11e1)
Try: sudo apt install <deb name>
HDFS.: command not found
hadoop@bmsccse-HP-Elite-Tower-800-G9-Desktop-PC:~/hadoop/sbin$ hadoop fs -cp /abc/file1.txt /newDir/
cp: /newDir/: No such file or directory: 'hdfs://localhost:9000/newDir/'
hadoop@bmsccse-HP-Elite-Tower-800-G9-Desktop-PC:~/hadoop/sbin$ hdfs dfs -mkdir /destinationDir
hadoop@bmsccse-HP-Elite-Tower-800-G9-Desktop-PC:~/hadoop/sbin$ hadoop fs -cp /abc/file1.txt /destinationDir/
```

```
hadoop@bmsccese-HP-Ellite-Tower-800-G9-Desktop-PC: ~/hadoop/sbin
how is your brother
how is your sister

hadoop@bmsccese-HP-Ellite-Tower-800-G9-Desktop-PC: ~/hadoop/sbin$ hdfs dfs -mv /abc /FFF
hadoop@bmsccese-HP-Ellite-Tower-800-G9-Desktop-PC: ~/hadoop/sbin$ hdfs dfs -ls /FFF
Found 2 items
-rw-r--r-- 1 hadoop supergroup          0 2025-04-15 14:34 /FFF/file1.txt
-rw-r--r-- 1 hadoop supergroup          0 2025-04-15 14:35 /FFF/file2.txt
hadoop@bmsccese-HP-Ellite-Tower-800-G9-Desktop-PC: ~/hadoop/sbin$ 9, Cp Hadoop HDFS

cp Command Usage cp Hadoop HDFS

cp Command Example

hadoop fs -cp /CSE/ /LLL
hadoop fs -ls /LLL

The cp command copies a file from one directory to another directory within the
HDFS.
9.: command not found
cp: target 'HDFS' is not a directory
cp: cannot stat 'command': No such file or directory
cp: '/CSE/': No such file or directory
ls: '/LLL': No such file or directory
Command 'the' not found, did you mean:
  Command 'he' from deb node-he (1.2.0-3)
  Command 'the' from deb the (3.3-rc1-3build1)
Try: sudo apt install <deb name>
HDFS: command not found
hadoop@bmsccese-HP-Ellite-Tower-800-G9-Desktop-PC: ~/hadoop/sbin$ hadoop fs -cp /abc/file1.txt /newDir/
cp: '/newDir/': No such file or directory: 'hdfs://localhost:9000/newDir'
hadoop@bmsccese-HP-Ellite-Tower-800-G9-Desktop-PC: ~/hadoop/sbin$ hdfs dfs -mkdir /destinationDir
hadoop@bmsccese-HP-Ellite-Tower-800-G9-Desktop-PC: ~/hadoop/sbin$ hadoop fs -cp /abc/file1.txt /destinationDir/
cp: '/abc/file1.txt': No such file or directory
hadoop@bmsccese-HP-Ellite-Tower-800-G9-Desktop-PC: ~/hadoop/sbin$ hdfs dfs -ls /abc/
ls: '/abc/': No such file or directory
hadoop@bmsccese-HP-Ellite-Tower-800-G9-Desktop-PC: ~/hadoop/sbin$ hdfs dfs -ls /
Found 5 items
drwxr-xr-x 1 hadoop supergroup          0 2025-04-15 14:35 /FFF
drwxr-xr-x 1 hadoop supergroup          0 2025-04-15 14:47 /destinationDir
drwxr-xr-x 1 hadoop supergroup          0 2024-05-13 15:12 /newDataFlair
drwxr-xr-x 1 hadoop supergroup          0 2024-05-13 14:36 /olddata
drwxr-xr-x 1 hadoop supergroup          0 2025-04-15 14:19 /rgs
hadoop@bmsccese-HP-Ellite-Tower-800-G9-Desktop-PC: ~/hadoop/sbin$ hdfs dfs -ls /FFF
Found 2 items
-rw-r--r-- 1 hadoop supergroup          0 2025-04-15 14:34 /FFF/file1.txt
-rw-r--r-- 1 hadoop supergroup          0 2025-04-15 14:35 /FFF/file2.txt
hadoop@bmsccese-HP-Ellite-Tower-800-G9-Desktop-PC: ~/hadoop/sbin$ hdfs dfs -cp /FFF/file1.txt /destinationDir/
hadoop@bmsccese-HP-Ellite-Tower-800-G9-Desktop-PC: ~/hadoop/sbin$ hdfs dfs -cp /FFF/file2.txt /destinationDir/
hadoop@bmsccese-HP-Ellite-Tower-800-G9-Desktop-PC: ~/hadoop/sbin$ hdfs dfs -ls /destinationDir/
Found 2 items
-rw-r--r-- 1 hadoop supergroup          0 2025-04-15 14:49 /destinationDir/file1.txt
-rw-r--r-- 1 hadoop supergroup          0 2025-04-15 14:49 /destinationDir/file2.txt
hadoop@bmsccese-HP-Ellite-Tower-800-G9-Desktop-PC: ~/hadoop/sbin$ 5-5-
```

```
hadoop@bmsccese-HP-Ellite-Tower-800-G9-Desktop-PC: ~/hadoop/sbin
SUBCOMMAND may print help when invoked w/o parameters or with -h.
hadoop@bmsccese-HP-Ellite-Tower-800-G9-Desktop-PC: $ cd /home/hadoop/hadoop/sbin
hadoop@bmsccese-HP-Ellite-Tower-800-G9-Desktop-PC: ~/hadoop/sbin$ ./start-dfs.sh
Starting namenodes on [localhost]
localhost: namenode is running as process 9900. Stop it first and ensure /tmp/hadoop-hadoop-namenode.pid file is empty before retry.
Starting datanodes
localhost: datanode is running as process 10056. Stop it first and ensure /tmp/hadoop-hadoop-datanode.pid file is empty before retry.
Starting secondary namenodes [bmsccese-HP-Ellite-Tower-800-G9-Desktop-PC]
bmsccese-HP-Ellite-Tower-800-G9-Desktop-PC: secondarynamenode is running as process 10334. Stop it first and ensure /tmp/hadoop-hadoop-secondarynamenode.pid file is empty before retry.
hadoop@bmsccese-HP-Ellite-Tower-800-G9-Desktop-PC: ~/hadoop/sbin$ ./start-yarn.sh
Starting resourcemanager
resourcemanager is running as process 10644. Stop it first and ensure /tmp/hadoop-hadoop-resourcemanager.pid file is empty before retry.
Starting nodemanagers
localhost: nodemanager is running as process 10807. Stop it first and ensure /tmp/hadoop-hadoop-nodemanager.pid file is empty before retry.
hadoop@bmsccese-HP-Ellite-Tower-800-G9-Desktop-PC: ~/hadoop/sbin$ jps
12226 Jps
10644 ResourceManager
10807 NodeManager
10056 DataNode
9900 NameNode
10334 SecondaryNameNode
hadoop@bmsccese-HP-Ellite-Tower-800-G9-Desktop-PC: ~/hadoop/sbin$ ./start-dfs.sh
Starting namenodes on [localhost]
localhost: namenode is running as process 9900. Stop it first and ensure /tmp/hadoop-hadoop-namenode.pid file is empty before retry.
Starting datanodes
localhost: datanode is running as process 10056. Stop it first and ensure /tmp/hadoop-hadoop-datanode.pid file is empty before retry.
Starting secondary namenodes [bmsccese-HP-Ellite-Tower-800-G9-Desktop-PC]
bmsccese-HP-Ellite-Tower-800-G9-Desktop-PC: secondarynamenode is running as process 10334. Stop it first and ensure /tmp/hadoop-hadoop-secondarynamenode.pid file is empty before retry.
hadoop@bmsccese-HP-Ellite-Tower-800-G9-Desktop-PC: ~/hadoop/sbin$ ./start-yarn.sh
Starting resourcemanager
resourcemanager is running as process 10644. Stop it first and ensure /tmp/hadoop-hadoop-resourcemanager.pid file is empty before retry.
Starting nodemanagers
localhost: nodemanager is running as process 10807. Stop it first and ensure /tmp/hadoop-hadoop-nodemanager.pid file is empty before retry.
hadoop@bmsccese-HP-Ellite-Tower-800-G9-Desktop-PC: ~/hadoop/sbin$ jps
13012 Jps
10644 ResourceManager
10807 NodeManager
10056 DataNode
9900 NameNode
10334 SecondaryNameNode
hadoop@bmsccese-HP-Ellite-Tower-800-G9-Desktop-PC: ~/hadoop/sbin$ hadoop fs -ls /
Found 2 items
drwxr-xr-x 1 hadoop supergroup          0 2024-05-13 15:12 /newDataFlair
drwxr-xr-x 1 hadoop supergroup          0 2024-05-13 14:36 /olddata
hadoop@bmsccese-HP-Ellite-Tower-800-G9-Desktop-PC: ~/hadoop/sbin$ hadoop fs -copyFromLocal /home/rekhags/Desktop/file1.txt /rgs/test.txt
copyFromLocal: /rgs/test.txt: No such file or directory: 'hdfs://localhost:9000/rgs/test.txt'
hadoop@bmsccese-HP-Ellite-Tower-800-G9-Desktop-PC: ~/hadoop/sbin$ hadoop fs -mkdir /rgs
hadoop@bmsccese-HP-Ellite-Tower-800-G9-Desktop-PC: ~/hadoop/sbin$ hadoop fs -copyFromLocal /home/rekhags/Desktop/file1.txt /rgs/test.txt
copyFromLocal: /home/rekhags/Desktop/file1.txt: No such file or directory
hadoop@bmsccese-HP-Ellite-Tower-800-G9-Desktop-PC: ~/hadoop/sbin$ ls /home/hadoop/Desktop/
file1.txt  first.txt  test.txt
hadoop@bmsccese-HP-Ellite-Tower-800-G9-Desktop-PC: ~/hadoop/sbin$ nano /home/hadoop/Desktop/file1.txt
hadoop@bmsccese-HP-Ellite-Tower-800-G9-Desktop-PC: ~/hadoop/sbin$ hadoop fs -copyFromLocal /home/hadoop/Desktop/file1.txt /rgs/test.txt
hadoop@bmsccese-HP-Ellite-Tower-800-G9-Desktop-PC: ~/hadoop/sbin$ hadoop fs -mkdir /rgs
```

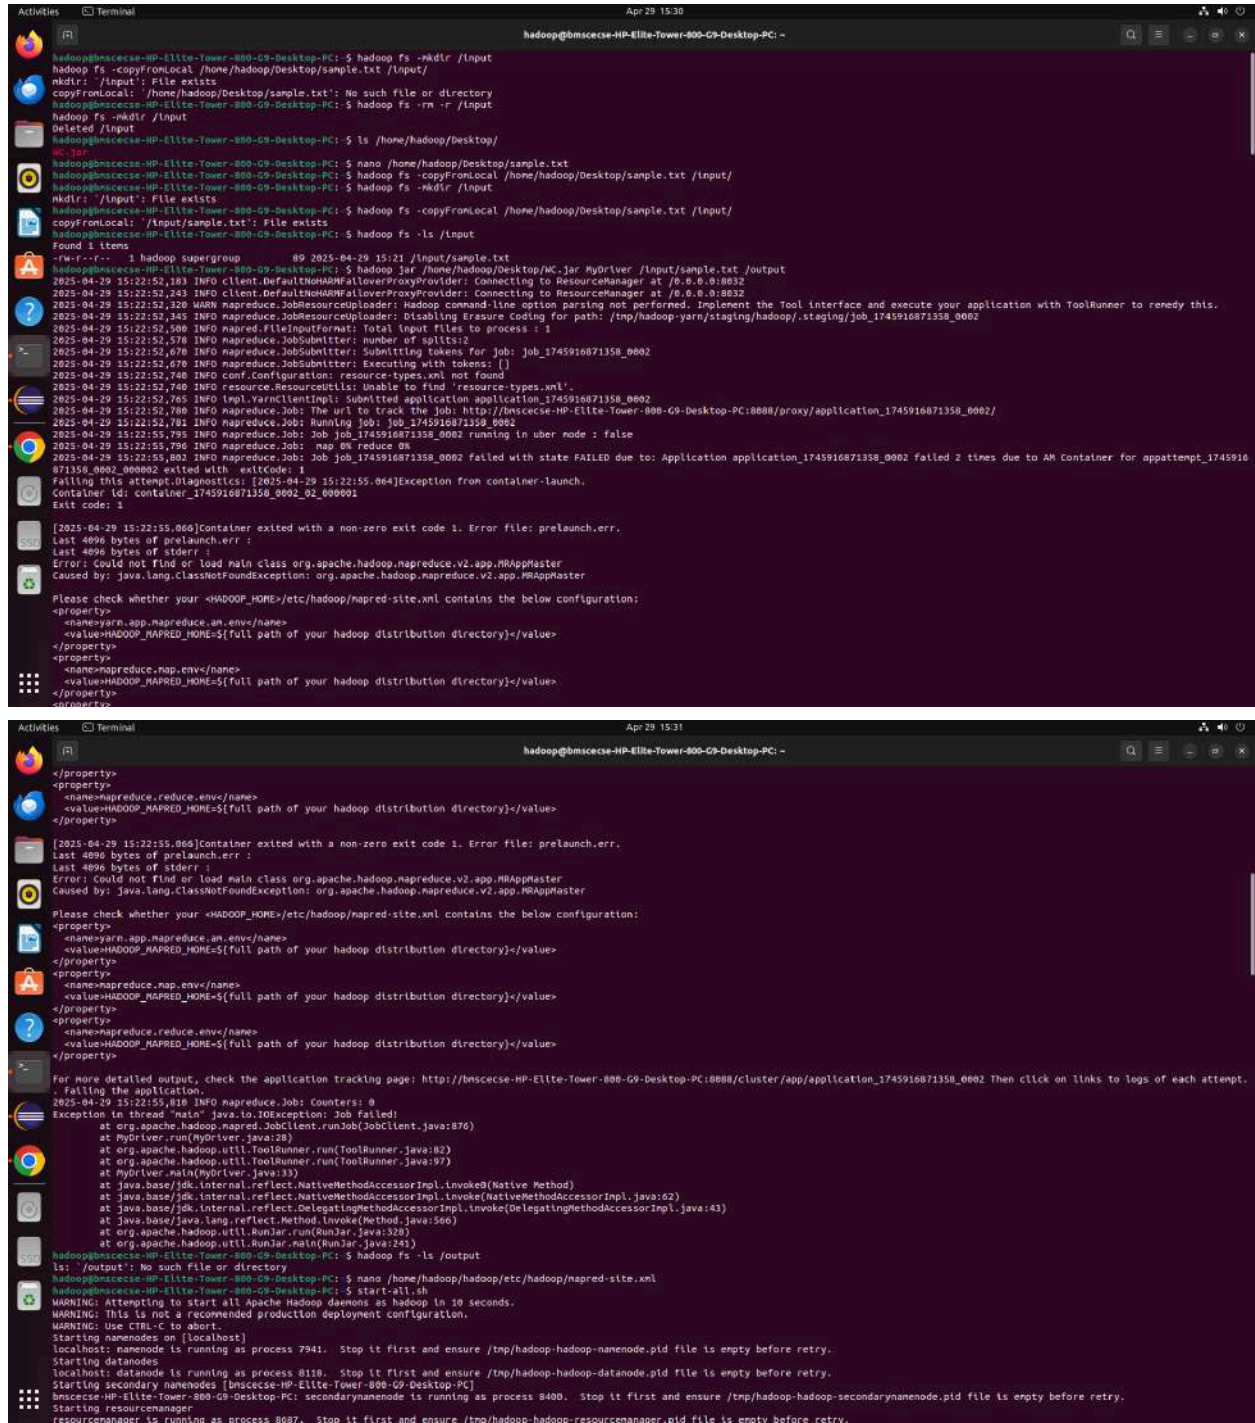
```
hadoop@bmsccse-HP-Elite-Tower-800-G9-Desktop-PC: ~/hadoop_wordcount
hadoop@bmsccse-HP-Elite-Tower-800-G9-Desktop-PC: $ ^[[200-sudo apt update
sudo: command not found
hadoop@bmsccse-HP-Elite-Tower-800-G9-Desktop-PC: $ -sudo apt update
Command '-sudo' not found, did you mean:
  command 'sudo' from deb sudo (1.9.9-1ubuntu2.4)
  command 'sudo' from deb sudo-ldap (1.9.9-1ubuntu2.4)
Try: sudo apt install -deb names
hadoop@bmsccse-HP-Elite-Tower-800-G9-Desktop-PC: $ tar -xzf hadoop-3.x.x.tar.gz
tar (chlid): hadoop-3.x.x.tar.gz: Cannot open: No such file or directory
tar (chlid): Error is not recoverable: exiting now
tar: Child returned Status 2
tar: Error is not recoverable: exiting now
hadoop@bmsccse-HP-Elite-Tower-800-G9-Desktop-PC: $ sudo apt update
[sudo] password for hadoop:
Hit:1 https://packages.microsoft.com/repos/code stable InRelease
Hit:2 http://security.ubuntu.com/ubuntu jammy-security InRelease
Get:3 https://dl.google.com/linux/chrome/deb stable InRelease [1,825 B]
Hit:4 http://in.archive.ubuntu.com/ubuntu jammy InRelease
Get:5 http://in.archive.ubuntu.com/ubuntu jammy-updates InRelease [128 kB]
Hit:6 https://repo.mongodb.org/apt/ubuntu jammy/mongodb-org/6.0 InRelease
Get:7 https://dl.google.com/linux/chrome/deb stable/main amd64 Packages [1,214 B]
Hit:9 https://repo.mongodb.org/apt/ubuntu jammy/mongodb-org/7.0 InRelease
Hit:10 http://in.archive.ubuntu.com/ubuntu jammy-backports InRelease
Hit:8 https://apache.jfrog.io/artifactory/cassandra-deb 41x InRelease
Fetched 131 kB in 3s (42.9 kB/s)
Reading package lists... Done
Building dependency tree... Done
Reading state information... Done
344 packages can be upgraded. Run 'apt list --upgradable' to see them.
W: https://repo.mongodb.org/apt/ubuntu/dists/jammy/mongodb-org/6.0/InRelease: Key is stored in legacy trusted.gpg keyring (/etc/apt/trusted.gpg), see the DEPRECATION section in apt-key(8) for details.
W: https://debian.cassandra.apache.org/dists/41x/InRelease: Key is stored in legacy trusted.gpg keyring (/etc/apt/trusted.gpg), see the DEPRECATION section in apt-key(8) for details.
hadoop@bmsccse-HP-Elite-Tower-800-G9-Desktop-PC: $ sudo apt install openjdk-11-jdk
Reading package lists... Done
Building dependency tree... Done
Reading state information... Done
openjdk-11-jdk is already the newest version (11.0.26+4-1ubuntu1-22.04).
The following packages were automatically installed and are no longer required:
  linux-headers-6.2.0-26-generic linux-hwe-6.2-headers-6.2.0-26
  linux-image-6.2.0-26-generic linux-modules-6.2.0-26-generic
  linux-modules-extra-6.2.0-26-generic
Use 'sudo apt autoremove' to remove them.
0 upgraded, 0 newly installed, 0 to remove and 344 not upgraded.
hadoop@bmsccse-HP-Elite-Tower-800-G9-Desktop-PC: $ java -version
openjdk version "11.0.20" 2023-01-21
OpenJDK Runtime Environment (build 11.0.20+4-post-Ubuntu-1ubuntu122.04)
OpenJDK 64-Bit Server VM (build 11.0.20+4-post-Ubuntu-1ubuntu122.04, mixed mode, sharing)
hadoop@bmsccse-HP-Elite-Tower-800-G9-Desktop-PC: $ mkdir -p ~/hadoop_wordcount
hadoop@bmsccse-HP-Elite-Tower-800-G9-Desktop-PC: ~/hadoop_wordcount$ nano WordCount.java
hadoop@bmsccse-HP-Elite-Tower-800-G9-Desktop-PC: ~/hadoop_wordcount$ javac -classpath $(hadoop classpath) -d /home/hadoop/wordcount_classes WordCount.java
WordCount.java:18: error: cannot find symbol
    public static class TokenizerMapper extends MapReduceBase implements Mapper<Object, Text, Text, IntWritable> {
                                   ^
  symbol:   class MapReduceBase
  location: class WordCount
WordCount.java:22: error: cannot find symbol
    ...
    ^
  symbol:   class MapReduceBase
  location: class WordCount
```



## EXPERIMENT-6

### Implement Wordcount program on Hadoop framework

#### Code with Screenshots:



```
Activities Terminal Apr 29 15:30
hadoop@bmscsece-hp-elite-tower-800-g9-desktop-pc: ~
hadoop@bmscsece-hp-elite-tower-800-g9-desktop-pc: ~$ hadoop fs -mkdir /input
hadoop fs -copyFromLocal /home/hadoop/Desktop/sample.txt /input/
mkdirt: '/input': File exists
copyFromLocal: '/home/hadoop/Desktop/sample.txt': No such file or directory
hadoop@bmscsece-hp-elite-tower-800-g9-desktop-pc: ~$ hadoop fs -rm -r /input
hadoop fs -mkdir /input
Deleted /input
hadoop@bmscsece-hp-elite-tower-800-g9-desktop-pc: ~$ ls /home/hadoop/Desktop/
sample.txt
hadoop@bmscsece-hp-elite-tower-800-g9-desktop-pc: ~$ nano /home/hadoop/Desktop/sample.txt
hadoop@bmscsece-hp-elite-tower-800-g9-desktop-pc: ~$ hadoop fs -copyFromLocal /home/hadoop/Desktop/sample.txt /input/
hadoop@bmscsece-hp-elite-tower-800-g9-desktop-pc: ~$ hadoop fs -mkdir /input
mkdirt: '/input': File exists
hadoop@bmscsece-hp-elite-tower-800-g9-desktop-pc: ~$ hadoop fs -copyFromLocal /home/hadoop/Desktop/sample.txt /input/
copyFromLocal: '/input/sample.txt': File exists
hadoop@bmscsece-hp-elite-tower-800-g9-desktop-pc: ~$ hadoop fs -ls /input
Found 1 items
-rw-r--r-- 1 hadoop supergroup 89 2025-04-29 15:21 /input/sample.txt
hadoop@bmscsece-hp-elite-tower-800-g9-desktop-pc: ~$ hadoop jar /home/hadoop/Desktop/PC.jar MyDriver /input/sample.txt /output
2025-04-29 15:22:52.181 INFO client.DefaultHadoopFollowerProxyProvider: Connecting to ResourceManager at /0.0.0.0:8032
2025-04-29 15:22:52.243 INFO client.DefaultHadoopFollowerProxyProvider: Connecting to ResourceManager at /0.0.0.0:8032
2025-04-29 15:22:52.320 WARN mapreduce.JobResourceUploader: Hadoop command-line option parsing not performed. Implement the Tool interface and execute your application with ToolRunner to remedy this.
2025-04-29 15:22:52.345 INFO mapreduce.JobResourceUploader: Disabling Erasure Coding for path: /tmp/hadoop-yarn/staging/hadoop/.staging/job_1745916871358_0002
2025-04-29 15:22:52.580 INFO mapreduce.FileInputFormat: Total input files to process : 1
2025-04-29 15:22:52.578 INFO mapreduce.JobSubmitter: number of splits:2
2025-04-29 15:22:52.670 INFO mapreduce.JobSubmitter: Submitting tokens for job: job_1745916871358_0002
2025-04-29 15:22:52.670 INFO mapreduce.JobSubmitter: Executing with tokens: []
2025-04-29 15:22:52.740 INFO conf.Configuration: resource-types.xml not found
2025-04-29 15:22:52.740 INFO resource.ResourceUtils: Unable to find 'resource-types.xml'.
2025-04-29 15:22:52.765 INFO Impl.YarnClientImpl: Submitted application application_1745916871358_0002
2025-04-29 15:22:52.781 INFO mapreduce.Job: The url to track the job: http://bmscsece-hp-elite-tower-800-g9-desktop-pc:8088/proxy/application_1745916871358_0002/
2025-04-29 15:22:52.781 INFO mapreduce.Job: Running job: job_1745916871358_0002
2025-04-29 15:22:52.795 INFO mapreduce.Job: Job job_1745916871358_0002 running in uber mode : false
2025-04-29 15:22:52.796 INFO mapreduce.Job: map 0% reduce 0%
2025-04-29 15:22:55.882 INFO mapreduce.Job: Job job_1745916871358_0002 failed with state FAILED due to: Application application_1745916871358_0002 failed 2 times due to AM Container for appattempt_1745916871358_0002_000002 exited with : exitCode: 1
Falling this attempt.Diagnostics: [2025-04-29 15:22:55.064]Exception from container-launch.
Container id: container_1745916871358_0002_02_000001
Exit code: 1
[2025-04-29 15:22:55.066]Container exited with a non-zero exit code 1. Error file: prelaunch.err.
Last 4096 bytes of prelaunch.err :
Last 4096 bytes of stderr :
Error: Could not find or load main class org.apache.hadoop.mapreduce.v2.app.MRAppMaster
Caused by: java.lang.ClassNotFoundException: org.apache.hadoop.mapreduce.v2.app.MRAppMaster
Please check whether your <HADOOP_HOME>/etc/hadoop/mapred-site.xml contains the below configuration:
<property>
<name>yarn.app.mapreduce.am.env</name>
<value>HADOOP_MAPRED_HOME=${full path of your hadoop distribution directory}</value>
</property>
<property>
<name>mapreduce.map.env</name>
<value>HADOOP_MAPRED_HOME=${full path of your hadoop distribution directory}</value>
</property>
</properties>
For more detailed output, check the application tracking page: http://bmscsece-hp-elite-tower-800-g9-desktop-pc:8088/cluster/app/application_1745916871358_0002 Then click on links to logs of each attempt.
Falling the application.
2025-04-29 15:22:55.810 INFO mapreduce.Job: Counters: 0
Exception in thread "main" java.io.IOException: Job failed!
at org.apache.hadoop.mapred.JobClient.runJob(JobClient.java:876)
at MyDriver.run(MyDriver.java:28)
at org.apache.hadoop.util.ToolRunner.run(ToolRunner.java:82)
at org.apache.hadoop.util.ToolRunner.run(ToolRunner.java:97)
at MyDriver.main(MyDriver.java:33)
at java.base/jdk.internal.reflect.NativeMethodAccessorImpl.invoke0(Native Method)
at java.base/jdk.internal.reflect.NativeMethodAccessorImpl.invoke(NativeMethodAccessorImpl.java:62)
at java.base/jdk.internal.reflect.DelegatingMethodAccessorImpl.invoke(DelegatingMethodAccessorImpl.java:43)
at java.base/java.lang.reflect.Method.invoke(Method.java:566)
at org.apache.hadoop.util.RunJar.run(RunJar.java:328)
at org.apache.hadoop.util.RunJar.main(RunJar.java:245)
hadoop@bmscsece-hp-elite-tower-800-g9-desktop-pc: ~$ hadoop fs -ls /output
ls: '/output': No such file or directory
hadoop@bmscsece-hp-elite-tower-800-g9-desktop-pc: ~$ nano /home/hadoop/hadoop/etc/hadoop/mapred-site.xml
hadoop@bmscsece-hp-elite-tower-800-g9-desktop-pc: ~$ start-all.sh
WARNING: Attempting to start all Apache Hadoop daemons as hadoop in 10 seconds.
WARNING: This is not a recommended production deployment configuration.
WARNING: Use CTRL-C to abort.
Starting namenodes on [localhost]
localhost: namenode is running as process 7941. Stop it first and ensure /tmp/hadoop-hadoop-namenode.pid file is empty before retry.
Starting datanodes
localhost: datanode is running as process 8118. Stop it first and ensure /tmp/hadoop-hadoop-datanode.pid file is empty before retry.
Starting secondary namenodes [bmscsece-hp-elite-tower-800-g9-desktop-pc]
bmscsece-hp-elite-tower-800-g9-desktop-pc: secondarynamenode is running as process 8400. Stop it first and ensure /tmp/hadoop-hadoop-secondarynamenode.pid file is empty before retry.
Starting resourcemanager
resourcemanager is running as process 8687. Stop it first and ensure /tmp/hadoop-hadoop-resourcemanager.pid file is empty before retry.
```

```
Activities Terminal Apr 29 15:31 hadoop@bmscscse-HP-Elite-Tower-800-G9-Desktop-PC:~$

resourceManager is running as process 8687. Stop it first and ensure /tmp/hadoop-hadoop-resourceManager.pid file is empty before retry.
Starting nodemanagers
localhost: nodemanager is running as process 8852. Stop it first and ensure /tmp/hadoop-hadoop-nodemanager.pid file is empty before retry.
hadoop@bmscscse-HP-Elite-Tower-800-G9-Desktop-PC:~$ jps
8400 SecondaryNameNode
15338 Jps
8852 NodeManager
7941 NameNode
8118 DataNode
18092 org.eclipse.equinox.launcher_1.6.1000.v20250227-1734.jar
8687 ResourceManager
hadoop@bmscscse-HP-Elite-Tower-800-G9-Desktop-PC:~$ cd ~/Desktop # or navigate to any other directory where you want to create the file
hadoop@bmscscse-HP-Elite-Tower-800-G9-Desktop-PC:~/Desktop$ nano sample.txt
hadoop@bmscscse-HP-Elite-Tower-800-G9-Desktop-PC:~/Desktop$ cd -
hadoop@bmscscse-HP-Elite-Tower-800-G9-Desktop-PC:~$ jps
8400 SecondaryNameNode
8852 NodeManager
7941 NameNode
8118 DataNode
15338 Jps
18092 org.eclipse.equinox.launcher_1.6.1000.v20250227-1734.jar
8687 ResourceManager
hadoop@bmscscse-HP-Elite-Tower-800-G9-Desktop-PC:~$ hadoop fs -mkdir /rgs
mkdir: '/rgs': File exists
hadoop@bmscscse-HP-Elite-Tower-800-G9-Desktop-PC:~$ hadoop fs -copyFromLocal ~/Desktop/sample.txt /rgs/test.txt
copyFromLocal: '/rgs/test.txt': File exists
hadoop@bmscscse-HP-Elite-Tower-800-G9-Desktop-PC:~$ hadoop fs -ls /rgs
Found 1 items
-rw-r--r-- 1 hadoop supergroup 90 2025-04-15 14:19 /rgs/test.txt
hadoop@bmscscse-HP-Elite-Tower-800-G9-Desktop-PC:~$ hadoop fs -cat /rgs/test.txt
hl how are you
how is your job
how is your family
how is your brother
how is your sister

hadoop@bmscscse-HP-Elite-Tower-800-G9-Desktop-PC:~$ hadoop jar /home/hadoop/Desktop/HC.jar MyDriver /rgs /output
2025-04-29 15:29:14,097 INFO impl.MetricsConfig: Loaded properties from hadoop-metrics2.properties
2025-04-29 15:29:14,935 INFO impl.MetricsSystemImpl: Scheduled Metric snapshot period at 10 second(s).
2025-04-29 15:29:14,935 INFO impl.MetricsSystemImpl: JobTracker metrics system started
2025-04-29 15:29:14,941 WARN impl.MetricsSystemImpl: JobTracker metrics system already initialized!
2025-04-29 15:29:14,996 WARN mapreduce.JobResourceUploader: Hadoop command-line option parsing not performed. Implement the Tool interface and execute your application with ToolRunner to remedy this.
2025-04-29 15:29:15,059 INFO mapred.FileInputFormat: Total input files to process : 1
2025-04-29 15:29:15,064 INFO mapreduce.JobSubmitter: number of splits: 1
2025-04-29 15:29:15,126 INFO mapreduce.JobSubmitter: Submitting tokens for job: job_local537170479_0001
2025-04-29 15:29:15,126 INFO mapreduce.JobSubmitter: Executing with tokens: []
2025-04-29 15:29:15,182 INFO mapreduce.Job: The url to track the job: http://localhost:8080/
2025-04-29 15:29:15,183 INFO mapreduce.Job: Running job: job_local537170479_0001
2025-04-29 15:29:15,183 INFO mapred.LocalJobRunner: OutputCommitter set in config null
2025-04-29 15:29:15,184 INFO mapred.LocalJobRunner: OutputCommitter is org.apache.hadoop.mapred.FileOutputCommitter
2025-04-29 15:29:15,186 INFO output.FileOutputCommitter: File Output Committer Algorithm version is 2
2025-04-29 15:29:15,187 INFO output.FileOutputCommitter: FileOutputCommitter skip cleanup_temporary folders under output directory:false, ignore cleanup failures: false
2025-04-29 15:29:15,221 INFO mapred.LocalJobRunner: Waiting for map tasks
2025-04-29 15:29:15,222 INFO mapred.LocalJobRunner: Starting task: attempt_local537170479_0001_n_000000_0
2025-04-29 15:29:15,233 INFO output.FileOutputCommitter: File Output Committer Algorithm version is 2
2025-04-29 15:29:15,233 INFO output.FileOutputCommitter: FileOutputCommitter skip cleanup_temporary folders under output directory:false, ignore cleanup failures: false
2025-04-29 15:29:15,233 INFO output.FileOutputCommitter: FileOutputCommitter skip cleanup_temporary folders under output directory:false, ignore cleanup failures: false
2025-04-29 15:29:15,242 INFO mapred.MapTask: Processing split: hdfs://localhost:9000/rgs/test.txt:0+90
2025-04-29 15:29:15,252 INFO mapred.MapTask: numReduceTasks: 1
2025-04-29 15:29:15,281 INFO mapred.MapTask: (EQUATOR) 0 kv1 26214396(104857584)
2025-04-29 15:29:15,281 INFO mapred.MapTask: mapreduce.task.io.sort.mb: 100
2025-04-29 15:29:15,281 INFO mapred.MapTask: soft limit at 83886080
2025-04-29 15:29:15,281 INFO mapred.MapTask: bufstart = 0; bufend = 104857600
2025-04-29 15:29:15,281 INFO mapred.MapTask: kvstart = 26214396(104857584); kvend = 26214320(104857280); length = 77/6553600
2025-04-29 15:29:15,319 INFO mapred.MapTask: Finished spill 0
2025-04-29 15:29:15,325 INFO mapred.Task: Task:attempt_local537170479_0001_n_000000_0 is done. And is in the process of committing
2025-04-29 15:29:15,327 INFO mapred.LocalJobRunner: hdfs://localhost:9000/rgs/test.txt:0+90
2025-04-29 15:29:15,329 INFO mapred.Task: Task:attempt_local537170479_0001_n_000000_0 done.
2025-04-29 15:29:15,329 INFO mapred.Task: Final Counters for attempt_local537170479_0001_n_000000_0: Counters: 23
File System Counters
FILE: Number of bytes read=4150
FILE: Number of bytes written=42471
FILE: Number of read operations=0
FILE: Number of large read operations=0
FILE: Number of write operations=0
HDFS: Number of bytes read=90
HDFS: Number of bytes written=0
HDFS: Number of read operations=5
HDFS: Number of large read operations=0
HDFS: Number of write operations=1
HDFS: Number of bytes read erasure-coded=0
Map-Reduce Framework
Map input records=0
Map output records=20
Map output bytes=160
Map output materialized bytes=215
```



```
Activities Terminal Apr 29 15:32
hadoop@bmscscse-HP-Elite-Tower-800-G9-Desktop-PC: ~
2025-04-29 15:29:15.316 INFO mapred.MapTask: bufstart = 0; bufend = 169; bufveto = 104857600
2025-04-29 15:29:15.316 INFO mapred.MapTask: kvstart = 26214396(104857584); kvend = 26214320(104857280); length = 77/6553600
2025-04-29 15:29:15.319 INFO mapred.MapTask: Finished split 0
2025-04-29 15:29:15.325 INFO mapred.Task: Task:attempt_local537170479_0001_r_000000_0 is done. And is in the process of committing
2025-04-29 15:29:15.327 INFO mapred.LocalJobRunner: hdfs://localhost:9000/rgs/test.txt:0-90
2025-04-29 15:29:15.327 INFO mapred.Task: Task: attempt_local537170479_0001_r_000000_0 done
2025-04-29 15:29:15.329 INFO mapred.Task: Final Counters for attempt_local537170479_0001_r_000000_0: Counters: 23
File System Counters
  FILE: Number of bytes read=4156
  FILE: Number of bytes written=42471
  FILE: Number of read operations=0
  FILE: Number of large read operations=0
  FILE: Number of write operations=0
  HDFS: Number of bytes read=90
  HDFS: Number of bytes written=6
  HDFS: Number of read operations=5
  HDFS: Number of large read operations=0
  HDFS: Number of write operations=1
  HDFS: Number of bytes read erasure-coded=0
Map-Reduce Framework
  Map input records=6
  Map output records=20
  Map output bytes=169
  Map output materialized bytes=215
  Input split bytes=60
  Combine input records=0
  Spilled Records=20
  Failed Shuffles=0
  Merged Map outputs=0
  GC time elapsed (ms)=0
  Total committed heap usage (bytes)=633339984
File Input Format Counters
  Bytes Read=60
2025-04-29 15:29:15.329 INFO mapred.LocalJobRunner: Finishing task: attempt_local537170479_0001_r_000000_0
2025-04-29 15:29:15.330 INFO mapred.LocalJobRunner: map task executor complete.
2025-04-29 15:29:15.331 INFO mapred.LocalJobRunner: Waiting for reduce tasks
2025-04-29 15:29:15.331 INFO mapred.LocalJobRunner: Starting task: attempt_local537170479_0001_r_000000_0
2025-04-29 15:29:15.335 INFO output.FileOutputCommitter: FileOutputCommitter Algorithm Version is 2
2025-04-29 15:29:15.335 INFO output.FileOutputCommitter: FileOutputCommitter skip cleanup_temporary folders under output directory:false, ignore cleanup failures: false
2025-04-29 15:29:15.336 INFO mapred.ReduceTask: Using ResourceCalculatorProcessTree : [ ]
2025-04-29 15:29:15.336 INFO mapred.ReduceTask: Using ShuffleConsumerPlugin: org.apache.hadoop.mapreduce.task.reduce.Shuffle27a9263f
2025-04-29 15:29:15.337 WARN Impl.MetricsSystemImpl: JobTracker metrics system already initialized!
2025-04-29 15:29:15.344 INFO reduce.MergeManagerImpl: MergeManager: memoryLimit=5827985408, maxSingleShuffleLimit=1456990352, mergeThreshold=3846470400, ioSortFactor=10, nonTenMergeOutputsThreshold=10
2025-04-29 15:29:15.344 INFO reduce.LocalFetcher: attempt_local537170479_0001_r_000000_0 Thread started: EventFetcher for fetching Map Completion Events
2025-04-29 15:29:15.356 INFO reduce.LocalFetcher: about to shuffle output of map attempt_local537170479_0001_r_000000_0 decomp: 211 len: 215 to MEMORY
2025-04-29 15:29:15.357 INFO reduce.InMemoryMapOutput: Read 211 bytes from map-output for attempt_local537170479_0001_r_000000_0
2025-04-29 15:29:15.358 INFO reduce.MergeManagerImpl: closeInMemoryFile -> map-output of size: 211, inMemoryMapOutputs.size() -> 1, commitMemory -> 0, usedMemory -> 211
2025-04-29 15:29:15.359 INFO reduce.EventFetcher: EventFetcher is interrupted.. Returning
2025-04-29 15:29:15.359 INFO mapred.LocalJobRunner: 1 / 1 copied.
2025-04-29 15:29:15.359 INFO reduce.MergeManagerImpl: finalMerge called with 1 in-memory map-outputs and 0 on-disk map-outputs
2025-04-29 15:29:15.362 INFO mapred.Reducer: Merging 1 sorted segments
2025-04-29 15:29:15.362 INFO reduce.MergeManagerImpl: Merged 1 segments, 211 bytes to disk to satisfy reduce memory limit
2025-04-29 15:29:15.362 INFO reduce.MergeManagerImpl: Merged 1 segments, 211 bytes to disk to satisfy reduce memory limit
2025-04-29 15:29:15.362 INFO reduce.MergeManagerImpl: Merging 1 files, 215 bytes from disk
2025-04-29 15:29:15.363 INFO reduce.MergeManagerImpl: Merging 0 segments, 0 bytes from memory into reduce
2025-04-29 15:29:15.363 INFO reduce.MergeManagerImpl: Merging 1 sorted segments
2025-04-29 15:29:15.363 INFO mapred.Reducer: Down to the last merge-pass, with 1 segments left of total size: 265 bytes
2025-04-29 15:29:15.363 INFO reduce.MergeManagerImpl: Merging 1 files, 215 bytes from disk
2025-04-29 15:29:15.363 INFO reduce.MergeManagerImpl: Merging 0 segments, 0 bytes from memory into reduce
2025-04-29 15:29:15.363 INFO mapred.Reducer: 1 / 1 copied.
2025-04-29 15:29:15.415 INFO mapred.Task: Task:attempt_local537170479_0001_r_000000_0 is done. And is in the process of committing
2025-04-29 15:29:15.416 INFO mapred.LocalJobRunner: 1 / 1 copied.
2025-04-29 15:29:15.416 INFO mapred.Task: Task:attempt_local537170479_0001_r_000000_0 is allowed to commit now
2025-04-29 15:29:15.423 INFO output.FileOutputCommitter: Saved output of task 'attempt_local537170479_0001_r_000000_0' to hdfs://localhost:9000/output
2025-04-29 15:29:15.424 INFO mapred.LocalJobRunner: reduce > reduce
2025-04-29 15:29:15.424 INFO mapred.Task: Task: attempt_local537170479_0001_r_000000_0 done.
2025-04-29 15:29:15.424 INFO mapred.Task: Final Counters for attempt_local537170479_0001_r_000000_0: Counters: 30
File System Counters
  FILE: Number of bytes read=4618
  FILE: Number of bytes written=42686
  FILE: Number of read operations=0
  FILE: Number of large read operations=0
  FILE: Number of write operations=0
  HDFS: Number of bytes read=90
  HDFS: Number of bytes written=69
  HDFS: Number of read operations=10
  HDFS: Number of large read operations=0
  HDFS: Number of write operations=3
  HDFS: Number of bytes read erasure-coded=0
Map-Reduce Framework
  Combine input records=0
  Combine output records=6
  Reduce input groups=10
  Reduce shuffle bytes=215
  Reduce input records=20
  Reduce output records=10
  Spilled Records=20
  Shuffled Maps=1
  Failed Shuffles=0
  Merged Map outputs=1
```



```
Activities Terminal Apr 29 15:33
hadop@bmscscce-HP-Elite-Tower-800-G9-Desktop-PC: ~

File System Counters
  FILE: Number of bytes read=4618
  FILE: Number of bytes written=442686
  FILE: Number of read operations=0
  FILE: Number of large read operations=0
  FILE: Number of write operations=0
  HDFS: Number of bytes read=56
  HDFS: Number of bytes written=69
  HDFS: Number of read operations=10
  HDFS: Number of large read operations=0
  HDFS: Number of write operations=1
  HDFS: Number of bytes read erasure-coded=0
Map-Reduce Framework
  Combine input records=0
  Combine output records=0
  Reduce input groups=10
  Reduce shuffle bytes=215
  Reduce input records=20
  Reduce output records=10
  Spilled Records=20
  Shuffled Maps=1
  Failed Shuffles=0
  Merged Map outputs=1
  GC time elapsed (ms)=0
  Total committed heap usage (bytes)=633339904
Shuffle Errors
  BAD_ID=0
  CONNECTION=0
  IO_ERROR=0
  WRONG_LENGTH=0
  WRONG_MAP=0
  WRONG_REDUCE=0
File Output Format Counters
  Bytes Written=69
2025-04-29 15:29:15,424 INFO mapred.LocalJobRunner: Finishing task: attempt_local537170479_0001_r_000000_0
2025-04-29 15:29:15,424 INFO mapred.LocalJobRunner: reduce task executor complete.
2025-04-29 15:29:16,185 INFO mapreduce.Job: Job job_local537170479_0001 running in uber mode : false
2025-04-29 15:29:16,186 INFO mapreduce.Job: map 100% reduce 100%
2025-04-29 15:29:16,186 INFO mapreduce.Job: Job job_local537170479_0001 completed successfully
2025-04-29 15:29:16,190 INFO mapreduce.Job: Counters: 36
  File System Counters
    FILE: Number of bytes read=8774
    FILE: Number of bytes written=1285157
    FILE: Number of read operations=0
    FILE: Number of large read operations=0
    FILE: Number of write operations=0
    HDFS: Number of bytes read=180
    HDFS: Number of bytes written=69
    HDFS: Number of read operations=15
    HDFS: Number of large read operations=0
    HDFS: Number of write operations=4
    HDFS: Number of bytes read erasure-coded=0
Map-Reduce Framework
  Map input records=6
  Map output records=20
  Map output bytes=169
  Map output materialized bytes=215
  Input split bytes=86
  Combine input records=0
  Combine output records=0
  Reduce input groups=10
  Reduce shuffle bytes=215
  Reduce input records=20
  Reduce output records=10
  Spilled Records=40
  Shuffled Maps=1
  Failed Shuffles=0
  Merged Map outputs=1
  GC time elapsed (ms)=0
  Total committed heap usage (bytes)=1266679808
Shuffle Errors
  BAD_ID=0
  CONNECTION=0
  IO_ERROR=0
  WRONG_LENGTH=0
  WRONG_MAP=0
  WRONG_REDUCE=0
File Input Format Counters
  Bytes Read=60
File Output Format Counters
  Bytes Written=69
0
hadop@bmscscce-HP-Elite-Tower-800-G9-Desktop-PC: ~ $ hadoop fs -ls /output
Found 2 items
-rw-r--r-- 1 hadoop supergroup 0 2025-04-29 15:29 /output/_SUCCESS
-rw-r--r-- 1 hadoop supergroup 69 2025-04-29 15:29 /output/part-000000
hadop@bmscscce-HP-Elite-Tower-800-G9-Desktop-PC: ~ $ hadoop fs -cat /output/part-000000
are 1
brother 1
family 1
```

```
FILE: Number of read operations=0
FILE: Number of large read operations=0
FILE: Number of write operations=0
HDFS: Number of bytes read=180
HDFS: Number of bytes written=69
HDFS: Number of read operations=15
HDFS: Number of large read operations=0
HDFS: Number of write operations=4
HDFS: Number of bytes read erasure-coded=0

Map-Reduce Framework
  Map input records=6
  Map output records=20
  Map output bytes=169
  Map output materialized bytes=215
  Input split bytes=40
  Combine input records=0
  Combine output records=0
  Reduce input groups=16
  Reduce shuffle bytes=215
  Reduce input records=20
  Reduce output records=10
  Spilled Records=40
  Shuffled Maps =1
  Failed Shuffles=0
  Merged Map outputs=1
  GC time elapsed (ms)=8
  Total committed heap usage (bytes)=1266679808

Shuffle Errors
BAD_ID=0
CONNECTION=0
IO_ERROR=0
WRONG_LENGTH=0
WRONG_MAP=0
WRONG_REDUCE=0

File Input Format Counters
  Bytes Read=90
File Output Format Counters
  Bytes Written=69

0
hadoop@bmscscse-HP-Elite-Tower-800-G9-Desktop-PC: $ hadoop fs -ls /output
Found 2 items
-rw-r--r-- 1 hadoop supergroup          0 2025-04-29 15:29 /output/_SUCCESS
-rw-r--r-- 1 hadoop supergroup    69 2025-04-29 15:29 /output/part-00000
hadoop@bmscscse-HP-Elite-Tower-800-G9-Desktop-PC: $ hadoop fs -cat /output/part-00000
are 1
brother 1
fantly 1
ht 1
how 5
ls 4
job 1
sister 1
you 1
your 4
hadoop@bmscscse-HP-Elite-Tower-800-G9-Desktop-PC: $
```

```
Activities  Eclipse  Apr 29 15:34
eclipse-workspace - WordCount/src/MyDriver.java - Eclipse IDE

File Edit Source Refactor Navigate Search Project Run Window Help

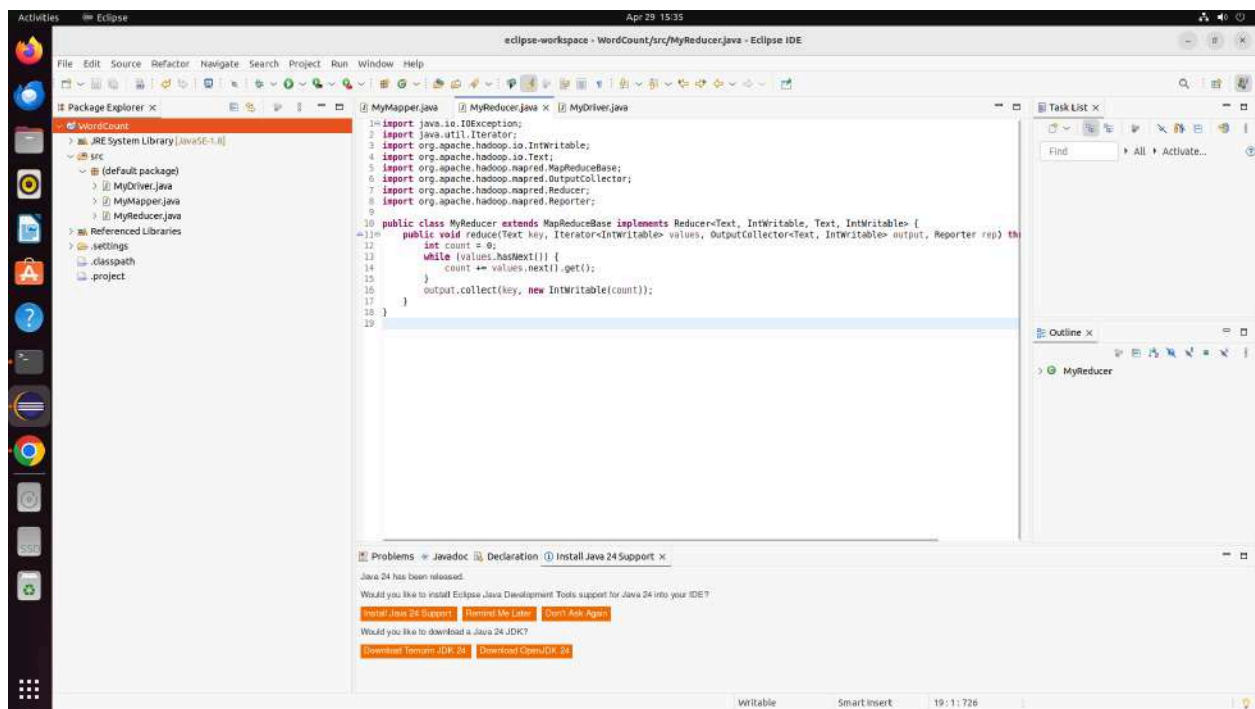
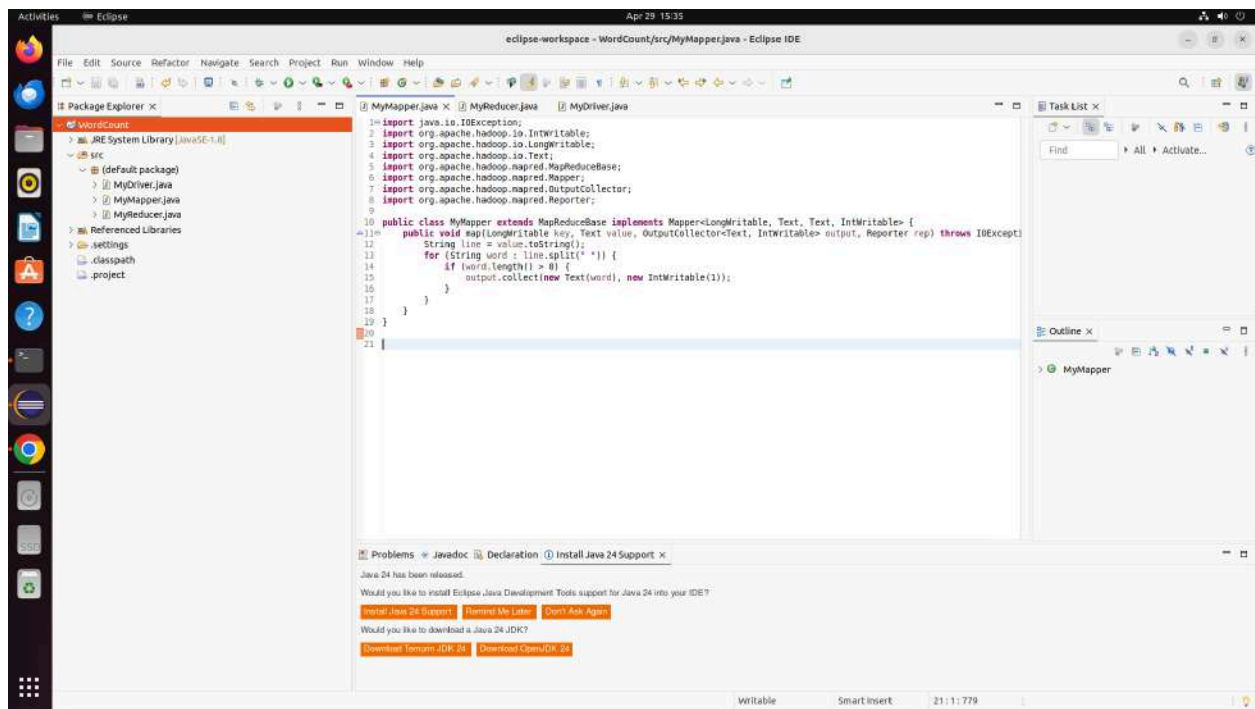
# Package Explorer x  # MyMapper.java  # MyReducer.java  # MyDriver.java x
> JRE System Library [jre56-1.8]
> src
  (default package)
  > MyDriver.java
  > MyMapper.java
  > MyReducer.java
> Referenced Libraries
> .settings
> .classpath
> .project

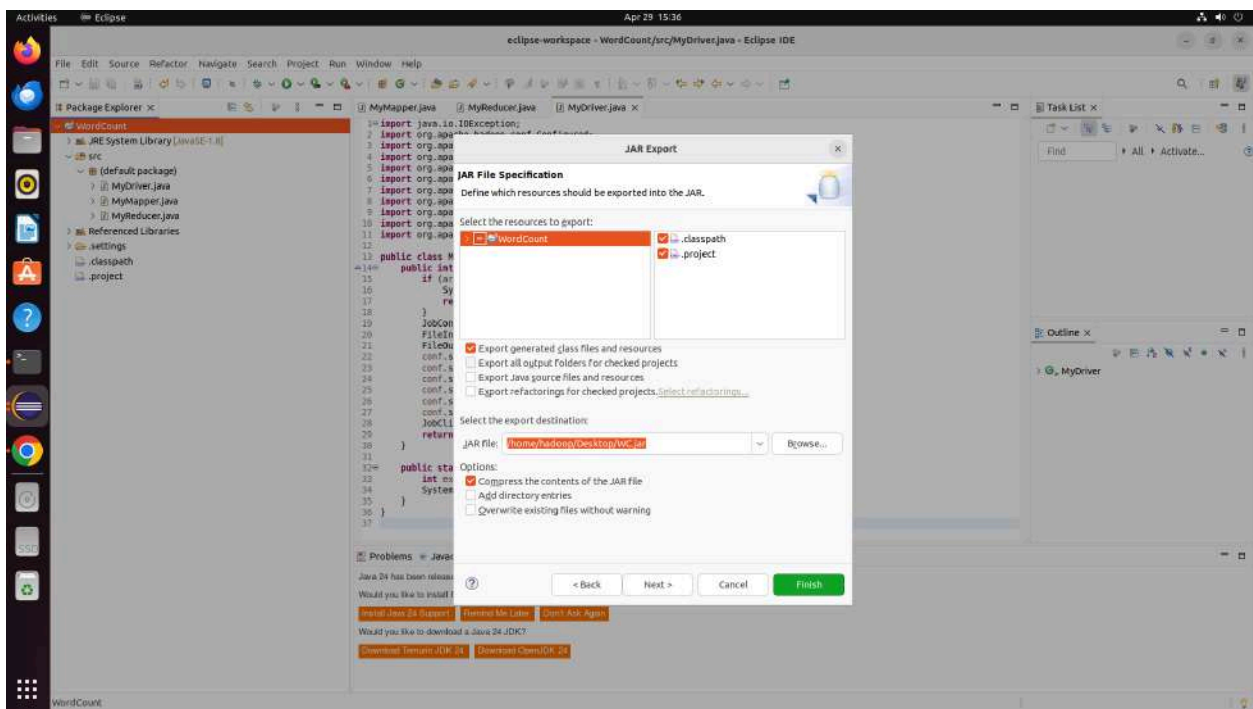
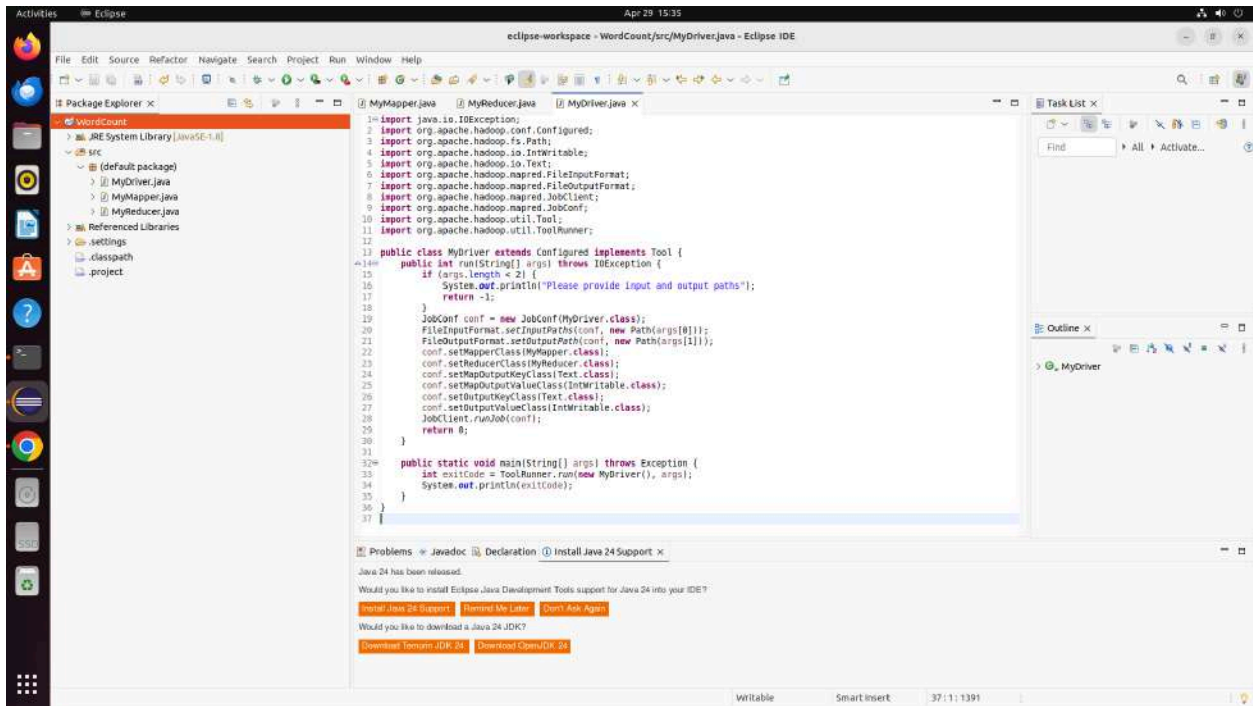
Task List x
Find  All  Activate...

Outline x
MyDriver

Problems x  Javadoc  Declaration  Install Java 24 Support x
Java 24 has been released.
Would you like to install Eclipse Java Development Tools support for Java 24 into your IDE?
Install Java 24 Support  Remind Me Later  Don't Ask Again
Would you like to download a Java 24 JDK?
Download Temurin JDK 24  Download OpenJDK 24

WordCount
```





## EXPERIMENT-7

From the following link extract the weather data

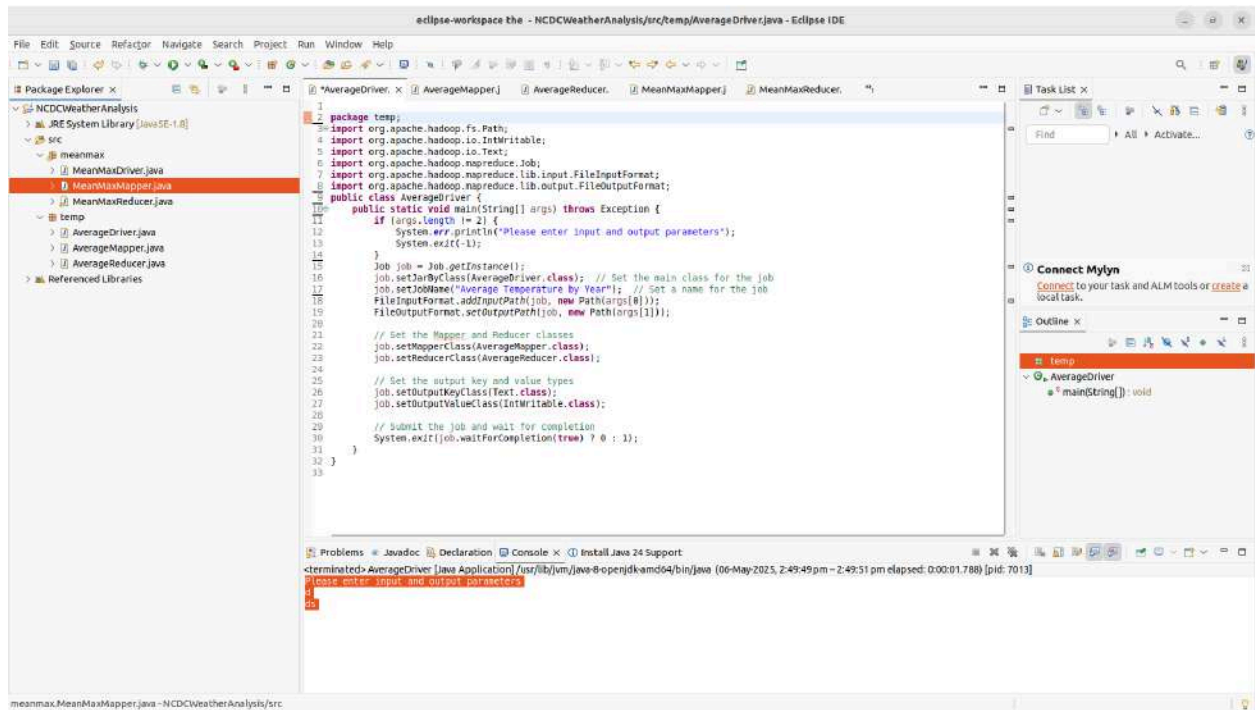
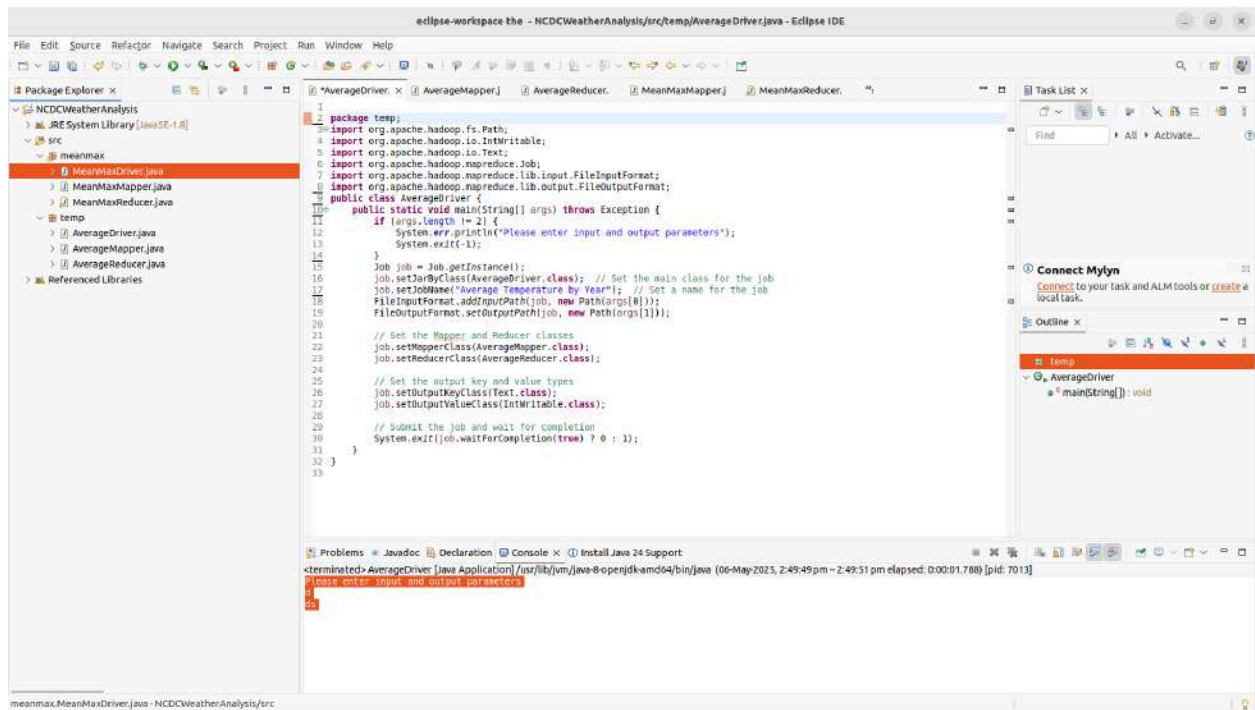
<https://github.com/tomwhite/hadoop-book/tree/master/input/ncdc/all>

a) Create a MapReduce program to find average temperature for each year from NCDC data set.

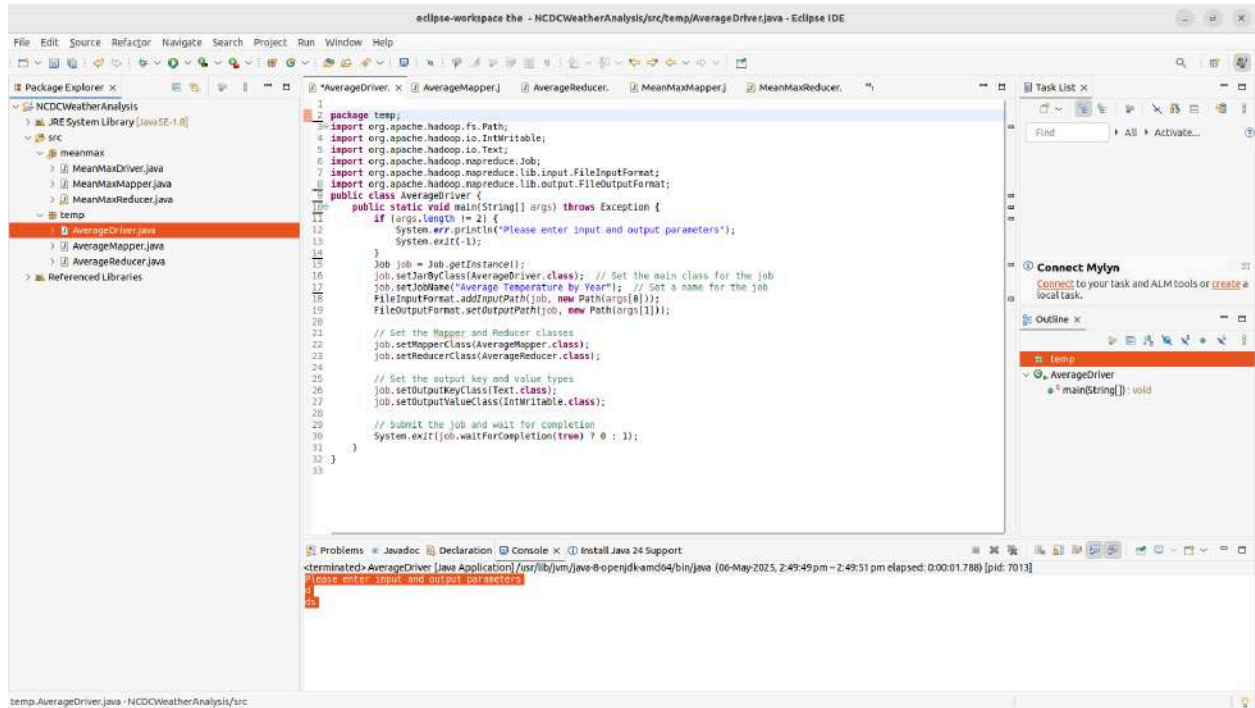
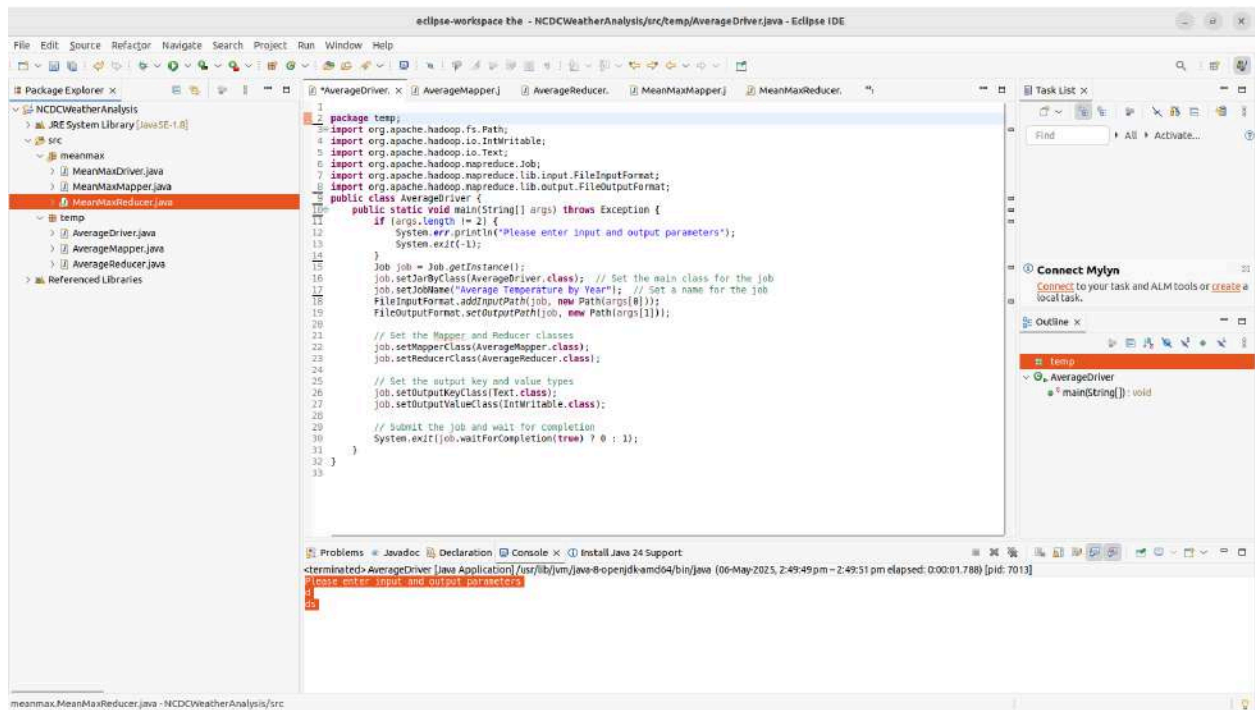
b) find the mean max temperature for every month.

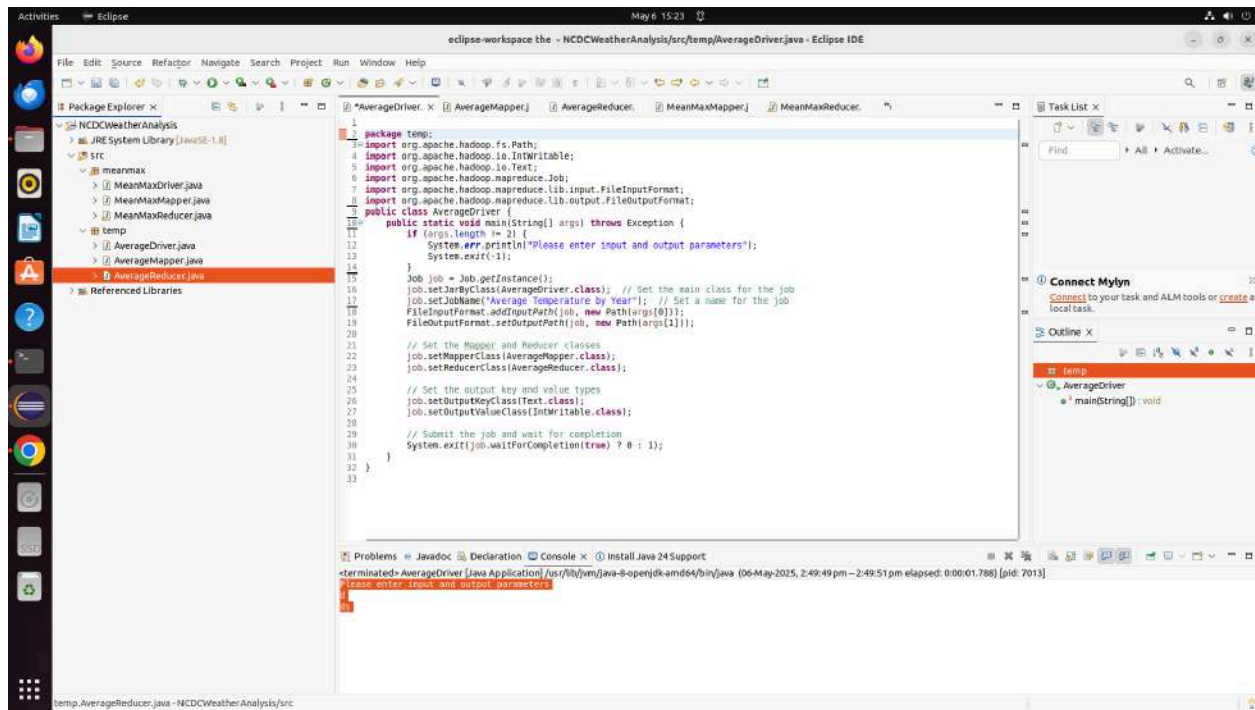
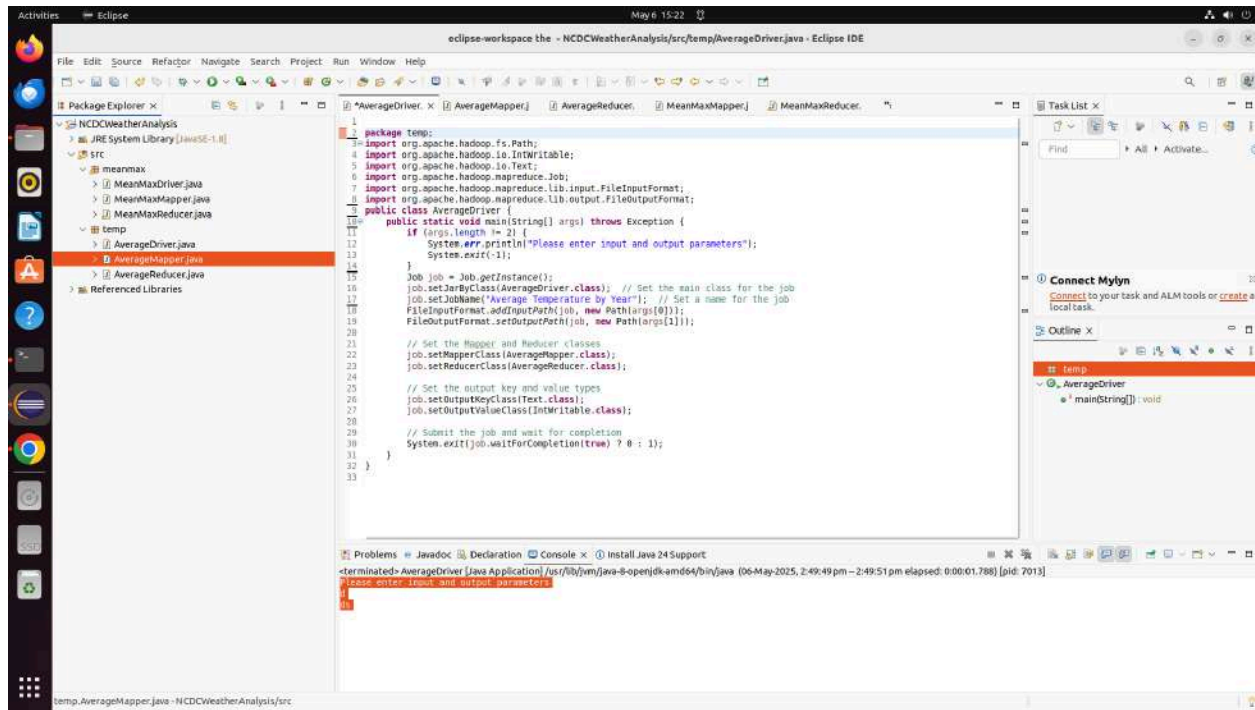
Code with screenshots:











```
Activities Terminal May 6 15:25
hadoop@bmsccese-HP-Elite-Tower-600-G9-Desktop-PC: ~/Desktop

hadoop@bmsccese-HP-Elite-Tower-600-G9-Desktop-PC: $ hdfs dfs -mkdir /ncdc_input
mkdir: call from bmsccese-HP-Elite-Tower-600-G9-Desktop-PC/127.0.0.1 to localhost:9000 failed on connection exception: java.net.ConnectException: Connection refused; For more details see: http://wiki.apa
che.org/hadoop/connectionRefused

hadoop@bmsccese-HP-Elite-Tower-600-G9-Desktop-PC: $ start-dfs.sh
start-yarn.sh
Starting namenodes on [localhost]
Starting datanodes
Starting secondary namenodes [bmsccese-HP-Elite-Tower-600-G9-Desktop-PC]
Starting resourcemanager
Starting nodemanagers
hadoop@bmsccese-HP-Elite-Tower-600-G9-Desktop-PC: $ jps
16032 SecondaryNameNode
15745 DataNode
15556 NameNode
7093 org.eclipse.equinox.launcher_1.6.1000.v20250227-1734-Jar
16504 NodeManager
16937 Jps
16318 ResourceManager

hadoop@bmsccese-HP-Elite-Tower-600-G9-Desktop-PC: $ hdfs dfs -mkdir /ncdc_input
hadoop@bmsccese-HP-Elite-Tower-600-G9-Desktop-PC: $ git clone https://github.com/tomwhite/hadoop-book.git
fatal: destination path 'hadoop-book' already exists and is not an empty directory.
hadoop@bmsccese-HP-Elite-Tower-600-G9-Desktop-PC: $ cd /Desktop/
hdfs dfs -put *.ncdc_input
hadoop@bmsccese-HP-Elite-Tower-600-G9-Desktop-PC: ~/hadoop-book/input/ncdc/all $ hdfs dfs -ls /ncdc_input
Found 2 items
-rw-r--r-- 1 hadoop supergroup 73067 2025-05-06 15:11 /ncdc_input/1902.gz
-rw-r--r-- 1 hadoop supergroup 74105 2025-05-06 15:11 /ncdc_input/1902.qz
hadoop@bmsccese-HP-Elite-Tower-600-G9-Desktop-PC: ~/hadoop-book/input/ncdc/all $ hadoop jar AverageTemp.jar temp.AverageDriver /ncdc_input /output_avg
JAR does not exist or is not a normal file: /home/hadoop/hadoop-book/input/ncdc/all/AverageTemp.jar
hadoop@bmsccese-HP-Elite-Tower-600-G9-Desktop-PC: ~/Desktop $ hadoop jar AverageTemp.jar temp.AverageDriver /ncdc_input /output_avg
JAR does not exist or is not a normal file: /home/hadoop/Desktop/AverageTemp.jar
hadoop@bmsccese-HP-Elite-Tower-600-G9-Desktop-PC: ~/Desktop $ is
ACR0C611800K291vblxb0_3-Pyt4Q3o2ap_tec250505hdtcf_rvnmI-d7XMX20B49nx4HgZfK256Ak8uTkhYxCYN8Kew7y0GF1117d183C8kdjx9GAsR23r3p1GQyTrpg5UC5U.pdf
file.txt
ncdcweather.jar
sample.txt
Tajawout.jar
ncdcCount.jar
hadoop@bmsccese-HP-Elite-Tower-600-G9-Desktop-PC: ~/Desktop $ hadoop jar ncdcweather.jar temp.AverageDriver /ncdc_input /output_avg
2025-05-06 15:14:57.391 INFO Impl.MetricsConfig: Loaded properties from hadoop-metrics2.properties
2025-05-06 15:14:57.428 INFO Impl.MetricsSystemImpl: Scheduled Metric snapshot period at 10 second(s).
2025-05-06 15:14:57.430 INFO Impl.MetricsSystemImpl: JobTracker metrics system started
2025-05-06 15:14:57.485 WARN mapreduce.JobResourceUploader: Hadoop command-line option parsing not performed. Implement the Tool Interface and execute your application with ToolRunner to remedy this.
2025-05-06 15:14:57.537 INFO Input.FileInputFormat: Total input files to process : 2
2025-05-06 15:14:57.553 INFO mapreduce.JobSubmitter: number of splits:2
2025-05-06 15:14:57.607 INFO mapreduce.JobSubmitter: Submitting tokens for job: job_local1944593155_0001
2025-05-06 15:14:57.607 INFO mapreduce.JobSubmitter: Executing with tokens: []
2025-05-06 15:14:57.656 INFO mapreduce.Job: The url to track the job: http://localhost:8080/
2025-05-06 15:14:57.657 INFO mapreduce.Job: Running job: job_local1944593155_0001
2025-05-06 15:14:57.657 INFO mapreduce.LocalJobRunner: OutputCommitter set to config null
2025-05-06 15:14:57.661 INFO output.FileOutputCommitter: FileOutputCommitter Algorithm version is 2
2025-05-06 15:14:57.661 INFO output.FileOutputCommitter: FileOutputCommitter skip cleanup_temporary folders under output directory:false, ignore cleanup failures: false
2025-05-06 15:14:57.661 INFO mapreduce.LocalJobRunner: OutputCommitter is org.apache.hadoop.mapreduce.lib.output.FileOutputCommitter
hadoop@bmsccese-HP-Elite-Tower-600-G9-Desktop-PC: ~/Desktop $ hadoop jar ncdcCount.jar temp.AverageDriver /ncdc_input /output_avg
2025-05-06 15:14:57.391 INFO Impl.MetricsConfig: Loaded properties from hadoop-metrics2.properties
2025-05-06 15:14:57.428 INFO Impl.MetricsSystemImpl: Scheduled Metric snapshot period at 10 second(s).
2025-05-06 15:14:57.430 INFO Impl.MetricsSystemImpl: JobTracker metrics system started
2025-05-06 15:14:57.485 WARN mapreduce.JobResourceUploader: Hadoop command-line option parsing not performed. Implement the Tool Interface and execute your application with ToolRunner to remedy this.
2025-05-06 15:14:57.537 INFO Input.FileInputFormat: Total input files to process : 2
2025-05-06 15:14:57.553 INFO mapreduce.JobSubmitter: number of splits:2
2025-05-06 15:14:57.607 INFO mapreduce.JobSubmitter: Submitting tokens for job: job_local1944593155_0001
2025-05-06 15:14:57.607 INFO mapreduce.JobSubmitter: Executing with tokens: []
2025-05-06 15:14:57.656 INFO mapreduce.Job: The url to track the job: http://localhost:8080/
2025-05-06 15:14:57.657 INFO mapreduce.Job: Running job: job_local1944593155_0001
2025-05-06 15:14:57.657 INFO mapreduce.LocalJobRunner: OutputCommitter set to config null
2025-05-06 15:14:57.661 INFO output.FileOutputCommitter: FileOutputCommitter Algorithm version is 2
2025-05-06 15:14:57.661 INFO mapreduce.LocalJobRunner: OutputCommitter is org.apache.hadoop.mapreduce.lib.output.FileOutputCommitter
2025-05-06 15:14:57.661 INFO mapreduce.LocalJobRunner: Waiting for map tasks
2025-05-06 15:14:57.692 INFO mapreduce.LocalJobRunner: Starting task: attempt_local1944593155_0001_m_000000_0
2025-05-06 15:14:57.701 INFO output.FileOutputCommitter: FileOutputCommitter Algorithm version is 2
2025-05-06 15:14:57.701 INFO output.FileOutputCommitter: FileOutputCommitter skip cleanup_temporary folders under output directory:false, ignore cleanup failures: false
2025-05-06 15:14:57.701 INFO mapreduce.Task: using ResourceCalculatorProcessTree [ ]
2025-05-06 15:14:57.706 INFO mapreduce.MapTask: Processing split: hdfs://localhost:9000/ncdc_input/1902.gz:0-74105
2025-05-06 15:14:57.739 INFO mapreduce.MapTask: (EQUATOR) 0 kv: 26214396(104857584)
2025-05-06 15:14:57.739 INFO mapreduce.MapTask: mapreduce.task.io.sort.mb: 100
2025-05-06 15:14:57.739 INFO mapreduce.MapTask: soft limit at: 8386080
2025-05-06 15:14:57.739 INFO mapreduce.MapTask: bufstart = 0; bufvold = 104857600
2025-05-06 15:14:57.739 INFO mapreduce.MapTask: kvstart = 26214396; length = 6553600
2025-05-06 15:14:57.741 INFO mapreduce.MapTask: Map output collector class = org.apache.hadoop.mapreduce.MapTask$MapOutputBuffer
2025-05-06 15:14:57.754 INFO zlib.ZlibFactory: Successfully loaded & initialized native-zlib library
2025-05-06 15:14:57.754 INFO compress.CodecPool: Got brand-new decompressor [.gz]
2025-05-06 15:14:57.815 INFO mapreduce.LocalJobRunner:
2025-05-06 15:14:57.816 INFO mapreduce.MapTask: Starting flush of map output
2025-05-06 15:14:57.816 INFO mapreduce.MapTask: Spilling map output
2025-05-06 15:14:57.816 INFO mapreduce.MapTask: bufstart = 0; bufend = 59885; bufvold = 104857600
2025-05-06 15:14:57.822 INFO mapreduce.MapTask: kvstart = 26214396(104857584); kvend = 26188140(104752508); length = 26257/6553600
2025-05-06 15:14:57.822 INFO mapreduce.MapTask: Finished spill 0
2025-05-06 15:14:57.826 INFO mapreduce.Task: Task:attempt_local1944593155_0001_m_000000_0 is done. And is in the process of committing
2025-05-06 15:14:57.828 INFO mapreduce.LocalJobRunner: map
2025-05-06 15:14:57.829 INFO mapreduce.Task: Task:attempt_local1944593155_0001_m_000000_0 done.
2025-05-06 15:14:57.831 INFO mapreduce.Task: Final counters for attempt_local1944593155_0001_m_000000_0: Counters: 23
File System Counters:
FILE: Number of bytes read=7873
FILE: Number of bytes written=719820
FILE: Number of read operations=0
FILE: Number of large read operations=0
FILE: Number of write operations=0
HDFS: Number of bytes read=74105
HDFS: Number of bytes written=0
HDFS: Number of read operations=5
HDFS: Number of large read operations=0
HDFS: Number of write operations=1
HDFS: Number of bytes read erasure-coded=0
Map-Reduce Framework
Map input records=6565
Map output records=6565
Map output bytes=59885
```



```
Activities Terminal May 6 15:26 hadoop@bmsccse-HP-Elite-Tower-600-G9-Desktop-PC: ~/Desktop

Map-Reduce Framework
  Map input records=6505
  Map output records=6505
  Map output bytes=59085
  Map output materialized bytes=72221
  Input split bytes=105
  Combine input records=0
  Spilled Records=6505
  Failed Shuffles=0
  Merged Map outputs=0
  GC time elapsed (ms)=0
  Total committed heap usage (bytes)=526385152
  File Input Format Counters
    Bytes Read=74105
2025-05-06 15:14:57.831 INFO mapred.LocalJobRunner: Finishing task: attempt_local1944593155_0001_r_000000_0
2025-05-06 15:14:57.831 INFO mapred.LocalJobRunner: Starting task: attempt_local1944593155_0001_r_000000_0
2025-05-06 15:14:57.832 INFO output.FileOutputCommitter: File Output Committer Algorithm version is 2
2025-05-06 15:14:57.832 INFO output.FileOutputCommitter: FileOutputCommitter skip cleanup _temporary folders under output directory:false, ignore cleanup failures: false
2025-05-06 15:14:57.832 INFO mapred.Task: Using ResourceCalculatorProcessTree : [ ]
2025-05-06 15:14:57.833 INFO mapred.MapTask: Processing split: hdfs://localhost:9000/ncdc_input/1901.gz:0:73807
2025-05-06 15:14:57.859 INFO mapred.MapTask: (EQUATOR) 0 kvl 26214396(104857584)
2025-05-06 15:14:57.859 INFO mapred.MapTask: mapreduce.task.io.sort.mb: 100
2025-05-06 15:14:57.859 INFO mapred.MapTask: soft limit at: 83860800
2025-05-06 15:14:57.859 INFO mapred.MapTask: bufstart = 0; bufvoid = 104857600
2025-05-06 15:14:57.859 INFO mapred.MapTask: kvstart = 26214396; length = 6553600
2025-05-06 15:14:57.860 INFO mapred.MapTask: Map output collector class = org.apache.hadoop.mapred.MapTask$MapOutputBuffer
2025-05-06 15:14:57.887 INFO mapred.LocalJobRunner: map
2025-05-06 15:14:57.887 INFO mapred.MapTask: Starting flush of map output
2025-05-06 15:14:57.887 INFO mapred.MapTask: Spilling map output
2025-05-06 15:14:57.887 INFO mapred.MapTask: bufstart = 0; bufend = 59076; bufvoid = 104857600
2025-05-06 15:14:57.887 INFO mapred.MapTask: kvstart = 26214396(104857584); kvend = 26188144(104752576); length = 26253/6553600
2025-05-06 15:14:57.890 INFO mapred.MapTask: Finished spill 0
2025-05-06 15:14:57.892 INFO mapred.Task: Task:attempt_local1944593155_0001_r_000000_0 is done. And is in the process of committing
2025-05-06 15:14:57.894 INFO mapred.LocalJobRunner: map
2025-05-06 15:14:57.894 INFO mapred.Task: Task 'attempt_local1944593155_0001_r_000000_0' done.
2025-05-06 15:14:57.895 INFO mapred.Task: Final Counters for attempt_local1944593155_0001_r_000000_0: Counters: 23
  File System Counters
    FILE: Number of bytes read=8102
    FILE: Number of bytes written=792860
    FILE: Number of read operations=0
    FILE: Number of large read operations=0
    FILE: Number of write operations=4
    HDFS: Number of bytes read=147972
    HDFS: Number of bytes written=0
    HDFS: Number of read operations=7
    HDFS: Number of large read operations=0
    HDFS: Number of write operations=1
    HDFS: Number of bytes read erasure-coded=0
  Map-Reduce Framework
    Map input records=6505
    Map output records=6504
    Map output bytes=59076
    Map output materialized bytes=72210
    Input split bytes=105
    Combine input records=0
    Spilled Records=6504
    Failed Shuffles=0
    Merged Map outputs=0
    GC time elapsed (ms)=0
    Total committed heap usage (bytes)=526385152
    File Input Format Counters
      Bytes Read=73807
2025-05-06 15:14:57.895 INFO mapred.LocalJobRunner: Finishing task: attempt_local1944593155_0001_r_000000_0
2025-05-06 15:14:57.895 INFO mapred.LocalJobRunner: map task executor complete.
2025-05-06 15:14:57.897 INFO mapred.LocalJobRunner: Waiting for reduce tasks
2025-05-06 15:14:57.897 INFO mapred.LocalJobRunner: Starting task: attempt_local1944593155_0001_r_000000_0
2025-05-06 15:14:57.904 INFO output.FileOutputCommitter: File Output Committer Algorithm version is 2
2025-05-06 15:14:57.904 INFO output.FileOutputCommitter: FileOutputCommitter skip cleanup _temporary folders under output directory:false, ignore cleanup failures: false
2025-05-06 15:14:57.904 INFO mapred.Task: Using ResourceCalculatorProcessTree : [ ]
2025-05-06 15:14:57.905 INFO mapred.ReduceTask: org.apache.hadoop.mapreduce.task.reduce.Shuffle@877f0e5
2025-05-06 15:14:57.906 WARN Impl: MetricsSystemImpl: JobTracker metrics system already initialized
2025-05-06 15:14:57.915 INFO reduce.MergeManagerImpl: MergeManager: memoryLimit=5827985488, maxSingleShuffleLimit=1456996352, mergeThreshold=3846478400, ioSortFactor=10, memToMemMergeOutputsThreshold=10
2025-05-06 15:14:57.916 INFO reduce.EventFetcher: attempt_local1944593155_0001_r_000000_0 Thread started: EventFetcher for fetching Map Completion Events
2025-05-06 15:14:57.927 INFO reduce.LocalFetcher: localfetcher1 about to shuffle output of map attempt_local1944593155_0001_r_000000_0 decomp: 72200 len: 72210 to MEMORY
2025-05-06 15:14:57.928 INFO reduce.InMemoryMapOutput: Read 72200 bytes from map-output for attempt_local1944593155_0001_r_000000_0
2025-05-06 15:14:57.928 INFO reduce.MergeManagerImpl: closeInMemoryFile -> map-output of size: 72200, inMemoryMapOutputs.size() -> 1, commitMemory -> 0, usedMemory -> 72200
2025-05-06 15:14:57.929 INFO reduce.LocalFetcher: localfetcher1 about to shuffle output of map attempt_local1944593155_0001_r_000000_0 decomp: 72217 len: 72221 to MEMORY
2025-05-06 15:14:57.930 INFO reduce.InMemoryMapOutput: Read 72217 bytes from map-output for attempt_local1944593155_0001_r_000000_0
2025-05-06 15:14:57.930 INFO reduce.MergeManagerImpl: closeInMemoryFile -> map-output of size: 72217, inMemoryMapOutputs.size() -> 2, commitMemory -> 72200, usedMemory -> 144423
2025-05-06 15:14:57.930 INFO reduce.EventFetcher: EventFetcher is interrupted.. Returning
2025-05-06 15:14:57.930 WARN io.ReadAheadPool: Failed readahead on file
  ERROR: bad file descriptor
  at org.apache.hadoop.io.nativeio.NativeIO$POSIX.posix_fadvise(Native Method)
  at org.apache.hadoop.io.nativeio.NativeIO$POSIX.posix_fadviseIfPossible(NativeIO.java:419)
  at org.apache.hadoop.io.nativeio.NativeIO$POSIX$CacheManipulator.posix_fadviseIfPossible(NativeIO.java:296)
  at org.apache.hadoop.io.ReadAheadPool$ReadAheadRequestImpl.run(ReadAheadPool.java:226)
  at java.base/java.util.concurrent.ThreadPoolExecutor.runWorker(ThreadPoolExecutor.java:1128)
  at java.base/java.util.concurrent.ThreadPoolExecutor$Worker.run(ThreadPoolExecutor.java:628)
  at java.base/java.lang.Thread.run(Thread.java:829)
2025-05-06 15:14:57.930 INFO mapred.LocalJobRunner: 2 / 2 copied.
2025-05-06 15:14:57.931 INFO reduce.MergeManagerImpl: finalMerge called with 2 in-memory map-outputs and 0 on-disk map-outputs
2025-05-06 15:14:57.933 INFO mapred.Merger: Merging 2 sorted segments
2025-05-06 15:14:57.933 INFO mapred.Merger: Down to the last merge-pass, with 2 segments left of total size: 144400 bytes
2025-05-06 15:14:57.938 INFO reduce.MergeManagerImpl: Merged 2 segments, 144423 bytes to disk to satisfy reduce memory limit
2025-05-06 15:14:57.938 INFO reduce.MergeManagerImpl: Merging 1 files, 144423 bytes from disk
2025-05-06 15:14:57.938 INFO reduce.MergeManagerImpl: Merging 0 segments, 0 bytes from memory into reduce
2025-05-06 15:14:57.938 INFO mapred.Merger: Merging 1 sorted segments
2025-05-06 15:14:57.939 INFO mapred.Merger: Down to the last merge-pass, with 1 segments left of total size: 144414 bytes
2025-05-06 15:14:57.939 INFO mapred.LocalJobRunner: 2 / 2 copied.
2025-05-06 15:14:57.968 INFO Configuration.deprecation: mapred.skip.on is deprecated. Instead, use mapreduce.job.skiprecords
2025-05-06 15:14:58.022 INFO mapred.Task: Task:attempt_local1944593155_0001_r_000000_0 is done. And is in the process of committing
2025-05-06 15:14:58.023 INFO mapred.LocalJobRunner: reduce - reduce
2025-05-06 15:14:58.023 INFO mapred.Task: Task:attempt_local1944593155_0001_r_000000_0 is allowed to commit now
2025-05-06 15:14:58.038 INFO output.FileOutputCommitter: Saved output of task 'attempt_local1944593155_0001_r_000000_0' to hdfs://localhost:9000/output_avg
2025-05-06 15:14:58.038 INFO mapred.LocalJobRunner: reduce - reduce
2025-05-06 15:14:58.038 INFO mapred.Task: Task 'attempt_local1944593155_0001_r_000000_0' done.
2025-05-06 15:14:58.038 INFO mapred.Task: Final Counters for attempt_local1944593155_0001_r_000000_0: Counters: 36
  File System Counters
    FILE: Number of bytes read=297822
```

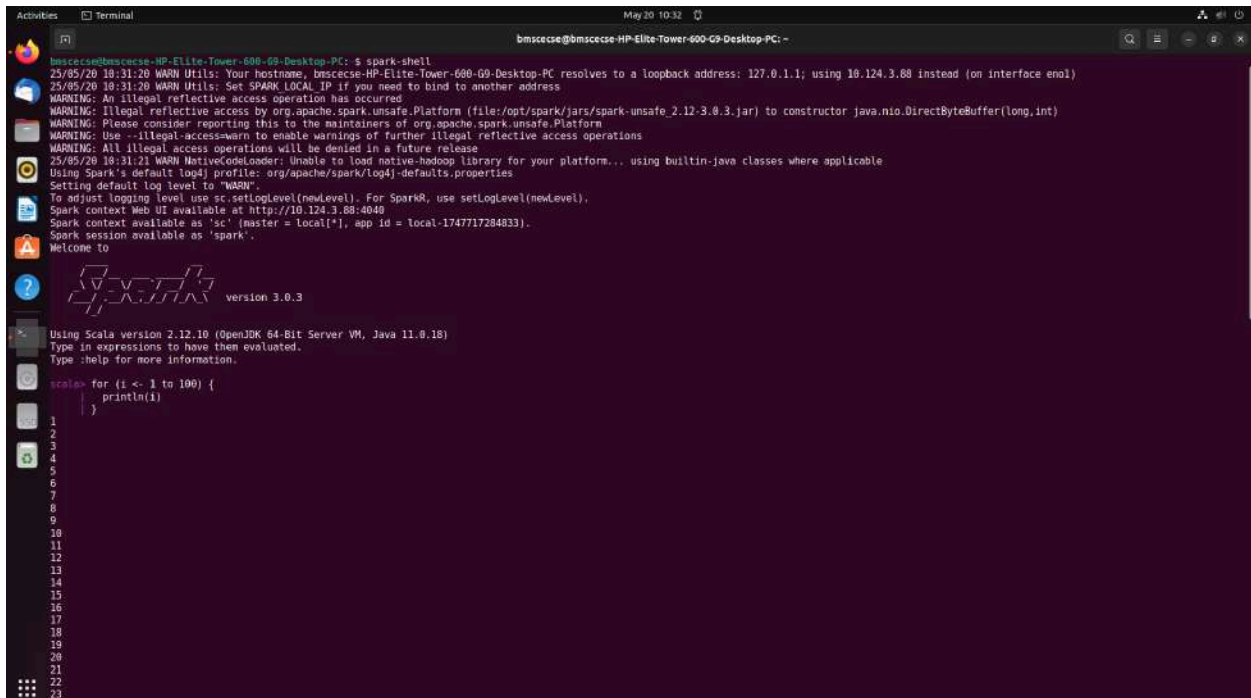
```
Activities Terminal May 6 15:27
hadoop@bmsccse-HP-Elite-Tower-600-G9-Desktop-PC: ~/Desktop


2025-05-06 15:14:58,038 INFO mapred.Task: Final Counters for attempt_local1944593155_0001_r_000000_0: Counters: 30
File System Counters
  FILE: Number of bytes read=297022
  FILE: Number of bytes written=36493
  FILE: Number of read operations=0
  FILE: Number of large read operations=0
  FILE: Number of write operations=0
  HDFS: Number of bytes read=147972
  HDFS: Number of bytes written=16
  HDFS: Number of read operations=12
  HDFS: Number of large read operations=0
  HDFS: Number of write operations=1
  HDFS: Number of bytes read erasure-coded=0
Map-Reduce Framework
  Combine Input records=0
  Combine output records=0
  Reduce input groups=2
  Reduce shuffle bytes=144431
  Reduce input records=13129
  Reduce output records=2
  Spilled Records=13129
  Shuffled Maps =2
  Failed Shuffles=0
  Merged Map outputs=2
  GC time elapsed (ms)=2
  Total committed heap usage (bytes)=633339904
Shuffle Errors
  BAD_ID=0
  CONNECTION=0
  IO_ERROR=0
  WRONG_LENGTH=0
  WRONG_MAP=0
  WRONG_REDUCE=0
File Output Format Counters
  Bytes Written=10
2025-05-06 15:14:58,038 INFO mapred.LocalJobRunner: Finishing task: attempt_local1944593155_0001_r_000000_0
2025-05-06 15:14:58,039 INFO mapred.LocalJobRunner: reduce task executor complete.
2025-05-06 15:14:58,059 INFO mapreduce.Job: Job: job_local1944593155_0001 running in uber mode : false
2025-05-06 15:14:58,559 INFO mapreduce.Job: map 100% reduce 100%
2025-05-06 15:14:58,660 INFO mapreduce.Job: Job: job_local1944593155_0001 completed successfully
2025-05-06 15:14:58,666 INFO mapreduce.Job: Counters: 36
File System Counters
  FILE: Number of bytes read=312997
  FILE: Number of bytes written=2448387
  FILE: Number of read operations=0
  FILE: Number of large read operations=0
  FILE: Number of write operations=0
  HDFS: Number of bytes read=370049
  HDFS: Number of bytes written=16
  HDFS: Number of read operations=24
  HDFS: Number of large read operations=0
  HDFS: Number of write operations=3
  HDFS: Number of bytes read erasure-coded=0
Map-Reduce Framework
  Map input records=13130
  Map output records=13129
  Reduce output records=12
  Spilled Records=20258
  Shuffled Maps =2
  Failed Shuffles=0
  Merged Map outputs=2
  GC time elapsed (ms)=3
  Total committed heap usage (bytes)=1686110288
Shuffle Errors
  BAD_ID=0
  CONNECTION=0
  IO_ERROR=0
  WRONG_LENGTH=0
  WRONG_MAP=0
  WRONG_REDUCE=0
File Input Format Counters
  Bytes Read=147972
File Output Format Counters
  Bytes Written=2
hadoop@bmsccse-HP-Elite-Tower-600-G9-Desktop-PC:~/Desktop$ hdfs dfs -ls /output_avg
Found 2 items
-rw-r--r-- 1 hadoop supergroup 0 2025-05-06 15:14 /output_avg/_SUCCESS
-rw-r--r-- 1 hadoop supergroup 16 2025-05-06 15:14 /output_avg/part-r-00000
hadoop@bmsccse-HP-Elite-Tower-600-G9-Desktop-PC:~/Desktop$ hdfs dfs -cat /output_avg/part-r-00000
1901 46
1902 21
hadoop@bmsccse-HP-Elite-Tower-600-G9-Desktop-PC:~/Desktop$ hdfs dfs -ls /output_avg
Found 2 items
-rw-r--r-- 1 hadoop supergroup 0 2025-05-06 15:14 /output_avg/_SUCCESS
-rw-r--r-- 1 hadoop supergroup 16 2025-05-06 15:14 /output_avg/part-r-00000
hadoop@bmsccse-HP-Elite-Tower-600-G9-Desktop-PC:~/Desktop$ hdfs dfs -cat /output_avg/part-r-00000
1901 46
1902 21
hadoop@bmsccse-HP-Elite-Tower-600-G9-Desktop-PC:~/Desktop$ hdfs dfs -cat /output_avg/part-r-00000
01 4
02 1
03 6
04 34
05 69
06 143
07 182
08 172
09 123
10 73
11 21
12 3
hadoop@bmsccse-HP-Elite-Tower-600-G9-Desktop-PC:~/Desktop$ hadoop@bmsccse-HP-Elite-Tower-600-G9-Desktop-PC:~/Desktop$ hdfs dfs -cat /output_avg/part-r-00000
1901 46
1902 21
hadoop@bmsccse-HP-Elite-Tower-600-G9-Desktop-PC:~/Desktop$ hdfs dfs -cat /output_avg/part-r-00000
1901 46
1902 21
```

## EXPERIMENT-8



Write a Scala program to print numbers from 1 to 100 using for loop.

Code with screenshots:



```
bmsccse@bmsccse-HP-Elite-Tower-600-G9-Desktop-PC: ~  
$ spark-shell  
25/05/20 18:31:20 WARN Utils: Your hostname, bmsccse-HP-Elite-Tower-600-G9-Desktop-PC resolves to a loopback address: 127.0.1.1; using 10.124.3.88 instead (on interface eno1)  
25/05/20 18:31:20 WARN Utils: Set SPARK_LOCAL_IP if you need to bind to another address  
WARNING: An illegal reflective access operation has occurred  
WARNING: Illegal reflective access by org.apache.spark.unsafe.Platform (file:/opt/spark/jars/spark-unsafe_2.12-3.0.3.jar) to constructor java.nio.DirectByteBuffer(long,int)  
WARNING: Please consider reporting this to the maintainers of org.apache.spark.unsafe.Platform  
WARNING: Use --illegal-access=warn to enable warnings of further illegal reflective access operations  
WARNING: All illegal access operations will be denied in a future release  
25/05/20 18:31:21 WARN NativeCodeLoader: Unable to load native-hadoop library for your platform... using builtin-java classes where applicable  
Using Spark's default log4j profile: org/apache/spark/log4j-defaults.properties  
Setting default log level to "WARN".  
To adjust logging level use sc.setLogLevel(newLevel). For SparkR, use setLogLevel(newLevel).  
Spark context Web UI available at http://10.124.3.88:4040  
Spark context available as 'sc' (master = local[*], app id = local-1747717284833).  
Spark session available as 'spark'.  
Welcome to  
 version 3.0.3  
Using Scala version 2.12.10 (OpenJDK 64-Bit Server VM, Java 11.0.18)  
Type in expressions to have them evaluated.  
Type :help for more information.  
scala> for (i <- 1 to 100) {  
  |   println(i)  
  | }  
1  
2  
3  
4  
5  
6  
7  
8  
9  
10  
11  
12  
13  
14  
15  
16  
17  
18  
19  
20  
21  
22  
23
```





Using RDD and FlatMap count how many times each word appears in a file and write out a list of words whose count is strictly greater than 4 using Spark.

```
bmsccec@bmsccece-HP-Ellite-Tower-B80-G9-Desktop-PC:~$ spark-shell
25/05/20 15:32:38 WARN Utils: Your hostname, bmscccece-HP-Ellite-Tower-800-G9-Desktop-PC resolves to a loopback address: 127.0.1.1; using 10.124.2.8 instead (on interface eno1)
25/05/20 15:32:38 WARN Utils: Set SPARK_LOCAL_IP if you need to bind to another address
WARNING: An illegal reflective access operation has occurred
WARNING: Illegal reflective access by org.apache.spark.unsafe.Platform (file:/opt/spark/jars/spark-unsafe_2.12-3.0.3.jar) to constructor java.nio.DirectByteBuffer(long,int)
WARNING: Please consider reporting this to the maintainers of org.apache.spark.unsafe.Platform
WARNING: Use --illegal-access=warn to enable warnings of further illegal reflective access operations
WARNING: All illegal access operations will be denied in a future release
25/05/20 15:32:38 WARN NativeCodeLoader: Unable to load native-hadoop library for your platform... using builtin-java classes where applicable
Using Spark's default log4j profile: org/apache/spark/log4j-defaults.properties
Setting default log level to "WARN".
To adjust logging level use sc.setLogLevel(newLevel). For SparkR, use setLogLevel(newLevel).
Spark context Web UI available at http://10.124.2.8:4040
Spark context available as 'sc' (master = local[*], app id = local-174773561481).
Spark session available as 'spark'.
Welcome to

      ____ _ 
     / ___/_ __| | |_ _|
    / /___\__ \| |_| |_) |
   /_____/___/\__ \__,_|_|
                          version 3.0.3


Using Scala version 2.12.10 (OpenJDK 64-Bit Server VM, Java 11.0.26)
Type in expressions to have them evaluated.
Type :help for more information.

scala> val textFile = sc.textFile("/home/bmsccece/Desktop/sparkdata.txt")
textFile: org.apache.spark.rdd.RDD[String] = /home/bmsccece/Desktop/sparkdata.txt MapPartitionsRDD[1] at textFile at <console>;2
4

scala> 

scala> val counts = textFile
counts: org.apache.spark.rdd.RDD[String] = /home/bmsccece/Desktop/sparkdata.txt MapPartitionsRDD[1] at textFile at <console>;24

scala> .flatMap(line => line.split(" "))
res0: org.apache.spark.rdd.RDD[String] = MapPartitionsRDD[2] at flatMap at <console>;26

scala> .map(word => (word, 1))

scala> 

scala> val data = sc.textFile("sparkdata.txt")
data: org.apache.spark.rdd.RDD[String] = sparkdata.txt MapPartitionsRDD[1] at textFile at <console>;25

scala> val splitdata = data.flatMap(line => line.split(" "))
splitdata: org.apache.spark.rdd.RDD[String] = MapPartitionsRDD[2] at flatMap at <console>;26

scala> val mapdata = splitdata.map(word => (word, 1))
mapdata: org.apache.spark.rdd.RDD[(String, Int)] = MapPartitionsRDD[3] at map at <console>;26

scala> val reducedata = mapdata.reduceByKey(_ + _)
reducedata: org.apache.spark.rdd.RDD[(String, Int)] = ShuffledRDD[4] at reduceByKey at <console>;26

scala> reducedata.collect.foreach(println)
(,1)
(hello,2)
(world,1)
(spark,1)
```

## EXPERIMENT-10

Write a simple streaming program in Spark to receive text data streams on a particular port, perform basic text cleaning (like white space removal, stop words removal, lemmatization, etc.), and print the cleaned text on the screen. (Open Ended Question).

Code with Screenshots:

```
import org.apache.spark.SparkConf

import org.apache.spark.streaming.{Seconds, StreamingContext}

import org.apache.spark.streaming.dstream.DStream

object StreamingTextCleaner {

    val stopWords = Set(
        "a", "an", "the", "is", "are", "was", "were", "this", "that",
        "on", "in", "and", "or", "of", "to", "for", "with", "as", "by"
    )

    def simpleLemmatize(word: String): String = {
        if (word.endsWith("ing")) word.dropRight(3)
        else if (word.endsWith("ed")) word.dropRight(2)
        else word
    }
}
```

```

def main(args: Array[String]): Unit = {

  if (args.length < 2) {

    System.err.println("Usage: StreamingTextCleaner <hostname> <port>")

    System.exit(1)

  }

  val hostname = args(0)

  val port = args(1).toInt

  val conf = new SparkConf().setAppName("StreamingTextCleaner").setMaster("local[*]")

  val ssc = new StreamingContext(conf, Seconds(5))

  val lines = ssc.socketTextStream(hostname, port)

  val cleanedLines: DStream[String] = lines.map(_._toLowerCase)

    .map(_._trim)

    .map(line => line.replaceAll("[^a-zA-Z\\s]", ""))

    .flatMap(_._split("\\s+"))

    .filter(word => word.nonEmpty && !stopWords.contains(word))

    .map(simpleLemmatize)

    .filter(_._nonEmpty)

    .reduceByWindow(

```

```

(a: String, b: String) => a + " " + b,

Seconds(5),

Seconds(5)

)

cleanedLines.print()

ssc.start()

ssc.awaitTermination()

}

}

```

```

Requirement already satisfied: nltk in /usr/local/lib/python3.11/dist-packages (3.9.1)
Requirement already satisfied: click in /usr/local/lib/python3.11/dist-packages (from nltk) (8.2.0)
Requirement already satisfied: joblib in /usr/local/lib/python3.11/dist-packages (from nltk) (1.5.0)
Requirement already satisfied: regex>=2021.8.3 in /usr/local/lib/python3.11/dist-packages (from nltk) (2024.11.6)
Requirement already satisfied: tqdm in /usr/local/lib/python3.11/dist-packages (from nltk) (4.67.1)
[nltk_data] Downloading package punkt to /root/nltk_data...
[nltk_data]   Unzipping tokenizers/punkt.zip.
[nltk_data] Downloading package stopwords to /root/nltk_data...
[nltk_data]   Unzipping corpora/stopwords.zip.
[nltk_data] Downloading package wordnet to /root/nltk_data...
+-----+
|word |
+-----+

|hello |
|hate  |
|hate  |
|love  |
|dont  |
|want  |
|cant  |
|put   |
|nobody|
|else  |
+-----+

```