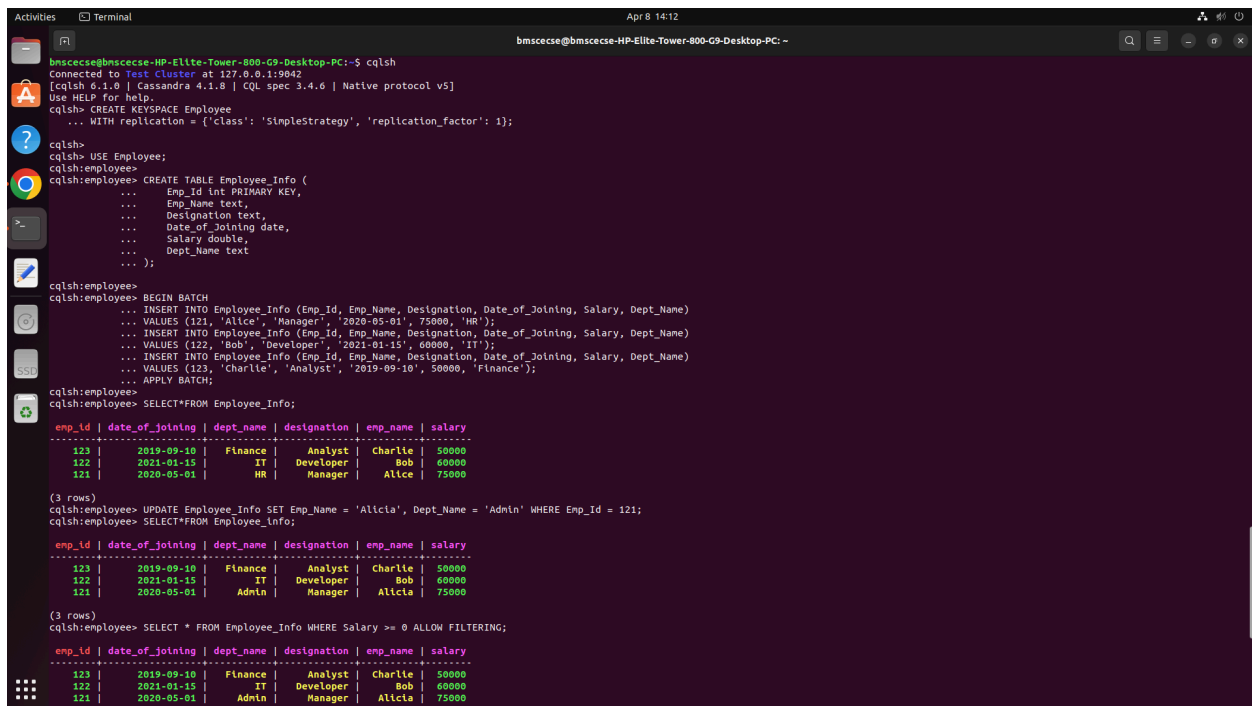


## EXPERIMENT-3

Perform the following DB operations using Cassandra.

- Create a keyspace by name Employee
- Create a column family by name Employee-Info with attributes Emp\_Id Primary Key, Emp\_Name, Designation, Date\_of\_Joining, Salary, Dept\_Name
- Insert the values into the table in batch
- Update Employee name and Department of Emp-Id 121
- Sort the details of Employee records based on salary
- Alter the schema of the table Employee\_Info to add a column Projects which stores a set of Projects done by the corresponding Employee.
- Update the altered table to add project names.
- Create a TTL of 15 seconds to display the values of Employees.

CODE WITH SCREENSHOTS:



```
bmsccese@bmsccese-HP-Ellite-Tower-800-G9-Desktop-PC:~$ cqlsh
Connected to Test cluster at 127.0.0.1:9042
[cqlsh 6.1.0 | Cassandra 4.1.8 | CQL spec 3.4.6 | Native protocol v5]
Use HELP for help.
cqlsh> CREATE KEYSPACE Employee
... WITH replication = {'class': 'SimpleStrategy', 'replication_factor': 1};

cqlsh> USE Employee;
cqlsh:employee>
cqlsh:employee> CREATE TABLE Employee_Info (
...     Emp_Id int PRIMARY KEY,
...     Emp_Name text,
...     Designation text,
...     Date_of_Joining date,
...     Salary double,
...     Dept_Name text
... );

cqlsh:employee>
cqlsh:employee> BEGIN BATCH
... INSERT INTO Employee_Info (Emp_Id, Emp_Name, Designation, Date_of_Joining, Salary, Dept_Name)
... VALUES (121, 'Alice', 'Manager', '2020-05-01', 75000, 'HR');
... INSERT INTO Employee_Info (Emp_Id, Emp_Name, Designation, Date_of_Joining, Salary, Dept_Name)
... VALUES (122, 'Bob', 'Developer', '2021-01-15', 60000, 'IT');
... INSERT INTO Employee_Info (Emp_Id, Emp_Name, Designation, Date_of_Joining, Salary, Dept_Name)
... VALUES (123, 'Charlie', 'Analyst', '2019-09-10', 50000, 'Finance');
... APPLY BATCH;

cqlsh:employee> SELECT * FROM Employee_Info;

emp_id | date_of_joining | dept_name | designation | emp_name | salary
-----|-----|-----|-----|-----|-----
123    | 2019-09-10     | Finance  | Analyst     | Charlie  | 50000
122    | 2021-01-15     | IT       | Developer   | Bob      | 60000
121    | 2020-05-01     | HR       | Manager     | Alice    | 75000
(3 rows)

cqlsh:employee> UPDATE Employee_Info SET Emp_Name = 'Alicia', Dept_Name = 'Admin' WHERE Emp_Id = 121;
cqlsh:employee> SELECT * FROM Employee_Info;

emp_id | date_of_joining | dept_name | designation | emp_name | salary
-----|-----|-----|-----|-----|-----
123    | 2019-09-10     | Finance  | Analyst     | Charlie  | 50000
122    | 2021-01-15     | IT       | Developer   | Bob      | 60000
121    | 2020-05-01     | Admin    | Manager     | Alicia   | 75000
(3 rows)

cqlsh:employee> SELECT * FROM Employee_Info WHERE Salary >= 0 ALLOW FILTERING;

emp_id | date_of_joining | dept_name | designation | emp_name | salary
-----|-----|-----|-----|-----|-----
123    | 2019-09-10     | Finance  | Analyst     | Charlie  | 50000
122    | 2021-01-15     | IT       | Developer   | Bob      | 60000
121    | 2020-05-01     | Admin    | Manager     | Alicia   | 75000
```

```
Activities Terminal Apr 8 14:12 bmscsece@bmscsece-HP-Ellite-Tower-800-G9-Desktop-PC: ~  
  
121 | 2020-05-01 | HR | Manager | Alice | 75000  
(3 rows)  
cqlsh:employee> UPDATE Employee_Info SET Emp_Name = 'Alicia', Dept_Name = 'Admn' WHERE Emp_Id = 121;  
cqlsh:employee> SELECT*FROM Employee_Info;  
  
emp_id | date_of_joining | dept_name | designation | emp_name | salary  
-----  
123 | 2019-09-10 | Finance | Analyst | Charlie | 50000  
122 | 2021-01-15 | IT | Developer | Bob | 60000  
121 | 2020-05-01 | Admn | Manager | Alicia | 75000  
(3 rows)  
cqlsh:employee> SELECT * FROM Employee_Info WHERE Salary >= 0 ALLOW FILTERING;  
  
emp_id | date_of_joining | dept_name | designation | emp_name | salary  
-----  
123 | 2019-09-10 | Finance | Analyst | Charlie | 50000  
122 | 2021-01-15 | IT | Developer | Bob | 60000  
121 | 2020-05-01 | Admn | Manager | Alicia | 75000  
(3 rows)  
cqlsh:employee> ALTER TABLE Employee_Info ADD Projects set<text>;  
cqlsh:employee> SELECT*FROM Employee_Info;  
  
emp_id | date_of_joining | dept_name | designation | emp_name | projects | salary  
-----  
123 | 2019-09-10 | Finance | Analyst | Charlie | null | 50000  
122 | 2021-01-15 | IT | Developer | Bob | null | 60000  
121 | 2020-05-01 | Admn | Manager | Alicia | null | 75000  
(3 rows)  
cqlsh:employee> UPDATE Employee_Info SET Projects = ('Payroll System', 'Attendance Tracker') WHERE Emp_Id = 121;  
cqlsh:employee> SELECT*FROM Employee_Info;  
  
emp_id | date_of_joining | dept_name | designation | emp_name | projects | salary  
-----  
123 | 2019-09-10 | Finance | Analyst | Charlie | null | 50000  
122 | 2021-01-15 | IT | Developer | Bob | null | 60000  
121 | 2020-05-01 | Admn | Manager | Alicia | {'Attendance Tracker', 'Payroll System'} | 75000  
(3 rows)  
cqlsh:employee> INSERT INTO Employee_Info (Emp_Id, Emp_Name, Designation, Date_of_Joining, Salary, Dept_Name)  
... VALUES (124, 'Tester', 'Tester', '2022-03-05', 45000, 'QA') USING TTL 15;  
cqlsh:employee> select*from employee_Info;  
  
emp_id | date_of_joining | dept_name | designation | emp_name | projects | salary  
-----  
123 | 2019-09-10 | Finance | Analyst | Charlie | null | 50000  
122 | 2021-01-15 | IT | Developer | Bob | null | 60000  
121 | 2020-05-01 | Admn | Manager | Alicia | {'Attendance Tracker', 'Payroll System'} | 75000  
(3 rows)  
cqlsh:employee> |
```

```
Activities Terminal Apr 8 14:24 bmscsece@bmscsece-HP-Ellite-Tower-800-G9-Desktop-PC: ~  
  
bmscsece@bmscsece-HP-Ellite-Tower-800-G9-Desktop-PC:~$ cqlsh  
Connected to Test Cluster at 127.0.0.1:9042  
[cqlsh 4.1.0 | Cassandra 4.1.0 | CQL spec 3.4.6 | Native protocol v5]  
Use HELP for help.  
cqlsh> CREATE KEYSPACE Library  
... WITH replication = {'class': 'SimpleStrategy', 'replication_factor': 1};  
cqlsh> USE Library;  
cqlsh:library>  
cqlsh:library> CREATE TABLE Library_Info (  
... Stud_Id int PRIMARY KEY,  
... Counter_value counter,  
... Stud_Name text,  
... Book_Name text,  
... Book_Id int,  
... Date_of_issue date  
... );  
InvalidRequest: Error from server: code=2200 [Invalid query] message="Cannot mix counter and non counter columns in the same table"  
cqlsh:library> CREATE TABLE Library_Info (  
... Stud_Id int PRIMARY KEY,  
... Stud_Name text,  
... Book_Name text,  
... Book_Id int,  
... Date_of_issue date  
... );  
cqlsh:library>  
cqlsh:library> CREATE TABLE Library_Counter (  
... Stud_Id int PRIMARY KEY,  
... Counter_value counter  
... );  
cqlsh:library>  
cqlsh:library> SELECT*FROM Library_Info;  
  
stud_id | book_id | book_name | date_of_issue | stud_name  
-----  
(0 rows)  
cqlsh:library> CREATE TABLE Library_Info1(  
... Stud_Id int,  
... Book_Name text,  
... Counter_Value counter,  
... PRIMARY KEY (Stud_Id, Book_Name)  
... );  
cqlsh:library> DROP TABLE LIBRARY_INFO;  
  
cqlsh:library>  
cqlsh:library>  
cqlsh:library> DROP TABLE LIBRARY_INFO1;  
cqlsh:library> CREATE TABLE Library_Info (  
... Stud_Id int,  
... Book_Name text,  
... Counter_Value counter,  
... PRIMARY KEY (Stud_Id, Book_Name)  
... );  
cqlsh:library>
```

```
Activities Terminal Apr 8 14:25 bmscsece@bmscsece-HP-Elite-Tower-800-G9-Desktop-PC: ~

... Counter_Value counter,
... PRIMARY KEY (Stud_Id, Book_Name)
... );

cqlsh:llibrary> BEGIN BATCH
... UPDATE llibrary_Info SET Counter_Value = Counter_Value + 1 WHERE Stud_Id = 112 AND Book_Name = 'BDA';
... UPDATE llibrary_Info SET Counter_Value = Counter_Value + 1 WHERE Stud_Id = 113 AND Book_Name = 'DSA';
... UPDATE llibrary_Info SET Counter_Value = Counter_Value + 1 WHERE Stud_Id = 112 AND Book_Name = 'BDA';
... APPLY BATCH;
InvalidRequest: Error from server: code=2200 [Invalid query] message="Cannot include a counter statement in a logged batch"
cqlsh:llibrary> BEGIN UNLOGGED BATCH
... UPDATE llibrary_Info SET Counter_Value = Counter_Value + 1 WHERE Stud_Id = 112 AND Book_Name = 'BDA';
... UPDATE llibrary_Info SET Counter_Value = Counter_Value + 1 WHERE Stud_Id = 113 AND Book_Name = 'DSA';
... UPDATE llibrary_Info SET Counter_Value = Counter_Value + 1 WHERE Stud_Id = 112 AND Book_Name = 'BDA';
... APPLY BATCH;
cqlsh:llibrary> SELECT*FROM llibrary_Info;

stud_id | book_name | counter_value
-----
113 | DSA | 1
112 | BDA | 2
(2 rows)
cqlsh:llibrary> SELECT * FROM Library_Info WHERE Stud_Id = 112 AND Book_Name = 'BDA';

stud_id | book_name | counter_value
-----
112 | BDA | 2
(1 rows)
cqlsh:llibrary> COPY Library_Info TO '/home/your_username/library_info.csv' WITH HEADER = TRUE;
Using 16 child processes
Can't open '/home/your_username/library_info.csv' for writing: [Errno 2] No such file or directory: '/home/your_username/library_info.csv'
cqlsh:llibrary> COPY library_info (stud_id, book_name, counter_value) TO '/home/bmscsece/library_info.csv';
Using 16 child processes

Starting copy of library.llibrary_info with columns [stud_id, book_name, counter_value].
Processed: 2 rows; Rate: 22 rows/s; Avg. rate: 22 rows/s
2 rows exported to 1 files in 0.105 seconds.
cqlsh:llibrary> COPY library_info (stud_id, book_name, counter_value) FROM '/home/bmscsece/library_info.csv' WITH HEADER = true;
Using 16 child processes

Starting copy of library.llibrary_info with columns [stud_id, book_name, counter_value].
Processed: 1 rows; Rate: 2 Rows/s; Avg. rate: 3 rows/s
1 rows imported from 1 files in 0.358 seconds (0 skipped).
cqlsh:llibrary> select*from library.library_info;

stud_id | book_name | counter_value
-----
113 | DSA | 1
112 | BDA | 4
(2 rows)
cqlsh:llibrary> 
```