

## EXPERIMENT-6

### Implement Wordcount program on Hadoop framework

#### Code with Screenshots:

```
Activities Terminal Apr 29 15:30
hadoop@bmscsece-HP-Ellite-Tower-800-G9-Desktop-PC: ~
hadoop@bmscsece-HP-Ellite-Tower-800-G9-Desktop-PC:~$ hadoop fs -mkdir /input
hadoop fs -copyFromLocal /home/hadoop/Desktop/sample.txt /input/
mkdir: '/input': File exists
copyFromLocal: '/home/hadoop/Desktop/sample.txt': No such file or directory
hadoop@bmscsece-HP-Ellite-Tower-800-G9-Desktop-PC:~$ hadoop fs -rm -r /input
hadoop fs -mkdir /input
Deleted /input
hadoop@bmscsece-HP-Ellite-Tower-800-G9-Desktop-PC:~$ ls /home/hadoop/Desktop/
HC.jar
hadoop@bmscsece-HP-Ellite-Tower-800-G9-Desktop-PC:~$ nano /home/hadoop/Desktop/sample.txt
hadoop@bmscsece-HP-Ellite-Tower-800-G9-Desktop-PC:~$ hadoop fs -copyFromLocal /home/hadoop/Desktop/sample.txt /input/
hadoop@bmscsece-HP-Ellite-Tower-800-G9-Desktop-PC:~$ hadoop fs -mkdir /input
mkdir: '/input': File exists
hadoop@bmscsece-HP-Ellite-Tower-800-G9-Desktop-PC:~$ hadoop fs -copyFromLocal /home/hadoop/Desktop/sample.txt /input/
copyFromLocal: '/input/sample.txt': File exists
hadoop@bmscsece-HP-Ellite-Tower-800-G9-Desktop-PC:~$ hadoop fs -ls /input
Found 1 items
-rw-r--r-- 1 hadoop supergroup 89 2025-04-29 15:21 /input/sample.txt
hadoop@bmscsece-HP-Ellite-Tower-800-G9-Desktop-PC:~$ hadoop jar /home/hadoop/Desktop/HC.jar MyDriver /input/sample.txt /output
2025-04-29 15:22:52.183 INFO client.DefaultHadoopFallbackProxyProvider: Connecting to ResourceManager at /0.0.0.0:8032
2025-04-29 15:22:52.243 INFO client.DefaultHadoopFallbackProxyProvider: Connecting to ResourceManager at /0.0.0.0:8032
2025-04-29 15:22:52.320 WARN mapreduce.JobResourceUploader: Hadoop command-line option parsing not performed. Implement the Tool interface and execute your application with ToolRunner to remedy this.
2025-04-29 15:22:52.345 INFO mapreduce.JobResourceUploader: Disabling Erasure Coding for path: /tmp/hadoop-yarn/staging/hadoop/.staging/job_1745916871358_0002
2025-04-29 15:22:52.580 INFO mapred.FileInputFormat: Total input files to process : 1
2025-04-29 15:22:52.578 INFO mapreduce.JobSubmitter: number of splits:2
2025-04-29 15:22:52.670 INFO mapreduce.JobSubmitter: Submitting tokens for job: job_1745916871358_0002
2025-04-29 15:22:52.670 INFO mapreduce.JobSubmitter: Executing with tokens: []
2025-04-29 15:22:52.740 INFO conf.Configuration: resource-types.xml not found
2025-04-29 15:22:52.740 INFO resource.ResourceUtils: Unable to find 'resource-types.xml'.
2025-04-29 15:22:52.765 INFO Impl.YarnClientImpl: Submitted application application_1745916871358_0002
2025-04-29 15:22:52.781 INFO mapreduce.Job: The url to track the job: http://bmscsece-HP-Ellite-Tower-800-G9-Desktop-PC:8088/proxy/application_1745916871358_0002/
2025-04-29 15:22:55.795 INFO mapreduce.Job: Job job_1745916871358_0002 running in uber mode : false
2025-04-29 15:22:55.796 INFO mapreduce.Job: map 0% reduce 0%
2025-04-29 15:22:55.802 INFO mapreduce.Job: Job job_1745916871358_0002 failed with state FAILED due to: Application application_1745916871358_0002 failed 2 times due to AM Container for appattempt1745916871358_0002_000002 exited with exitcode: 1
Falling this attempt.Diagnostics: [2025-04-29 15:22:55.064]Exception from container-launch.
Container id: container_1745916871358_0002_02_000001
Exit code: 1
[2025-04-29 15:22:55.066]Container exited with a non-zero exit code 1. Error file: prelaunch.err.
Last 4096 bytes of prelaunch.err :
Last 4096 bytes of stderr :
Error: Could not find or load main class org.apache.hadoop.mapreduce.v2.app.MRAppMaster
Caused by: java.lang.ClassNotFoundException: org.apache.hadoop.mapreduce.v2.app.MRAppMaster

Please check whether your <HADOOP_HOME>/etc/hadoop/mapred-site.xml contains the below configuration:
<property>
<name>yarn.app.mapreduce.am.env</name>
<value>HADOOP_MAPRED_HOME=${full path of your hadoop distribution directory}</value>
</property>
<property>
<name>mapreduce.map.env</name>
<value>HADOOP_MAPRED_HOME=${full path of your hadoop distribution directory}</value>
</property>
<property>
<name>mapreduce.reduce.env</name>
<value>HADOOP_MAPRED_HOME=${full path of your hadoop distribution directory}</value>
</property>
</property>
For more detailed output, check the application tracking page: http://bmscsece-HP-Ellite-Tower-800-G9-Desktop-PC:8088/cluster/app/application_1745916871358_0002 Then click on links to logs of each attempt.
. Falling the application.
2025-04-29 15:22:55.810 INFO mapreduce.Job: Counters: 0
Exception in thread "main" java.io.IOException: Job failed!
at org.apache.hadoop.mapred.JobClient.runJob(JobClient.java:876)
at MyDriver.run(MyDriver.java:28)
at org.apache.hadoop.util.ToolRunner.run(ToolRunner.java:82)
at org.apache.hadoop.util.ToolRunner.run(ToolRunner.java:97)
at MyDriver.main(MyDriver.java:33)
at java.base/jdk.internal.reflect.NativeMethodAccessorImpl.invoke(Native Method)
at java.base/jdk.internal.reflect.NativeMethodAccessorImpl.invoke(NativeMethodAccessorImpl.java:62)
at java.base/jdk.internal.reflect.DelegatingMethodAccessorImpl.invoke(DelegatingMethodAccessorImpl.java:43)
at java.base/java.lang.reflect.Method.invoke(Method.java:566)
at org.apache.hadoop.util.RunJar.run(RunJar.java:328)
at org.apache.hadoop.util.RunJar.main(RunJar.java:243)
hadoop@bmscsece-HP-Ellite-Tower-800-G9-Desktop-PC:~$ hadoop fs -ls /output
ls: '/output': No such file or directory
hadoop@bmscsece-HP-Ellite-Tower-800-G9-Desktop-PC:~$ nano /home/hadoop/hadoop/etc/hadoop/mapred-site.xml
hadoop@bmscsece-HP-Ellite-Tower-800-G9-Desktop-PC:~$ start-all.sh
WARNING: Attempting to start all Apache Hadoop daemons as hadoop in 10 seconds.
WARNING: This is not a recommended production deployment configuration.
WARNING: Use CTRL-C to abort.
Starting namenodes on [localhost]
localhost: namenode is running as process 7941. Stop it first and ensure /tmp/hadoop-hadoop-namenode.pid file is empty before retry.
Starting datanodes
localhost: datanode is running as process 8118. Stop it first and ensure /tmp/hadoop-hadoop-datanode.pid file is empty before retry.
Starting secondary namenodes [bmscsece-HP-Ellite-Tower-800-G9-Desktop-PC]
bmscsece-HP-Ellite-Tower-800-G9-Desktop-PC: secondarynamenode is running as process 8400. Stop it first and ensure /tmp/hadoop-hadoop-secondarynamenode.pid file is empty before retry.
Starting resourcemanager
resourcemanager is running as process 8687. Stop it first and ensure /tmp/hadoop-hadoop-resourcemanager.pid file is empty before retry.
```

```
Activities Terminal Apr 29 15:31
hadoop@bmsccscse-HP-Elite-Tower-800-G9-Desktop-PC: ~
resourceManager is running as process 8687. Stop it first and ensure /tmp/hadoop-hadoop-resourceManager.pid file is empty before retry.
Starting namenangers
localhost: namenanager is running as process 8852. Stop it first and ensure /tmp/hadoop-hadoop-namenanager.pid file is empty before retry.
hadoop@bmsccscse-HP-Elite-Tower-800-G9-Desktop-PC: ~$ jps
8400 SecondaryNameNode
15538 Jps
8852 NodeManager
7941 NameNode
8118 DataNode
10892 org.eclipse.equinox.launcher_1.6.1000.v20250227-1734.jar
8687 ResourceManager
hadoop@bmsccscse-HP-Elite-Tower-800-G9-Desktop-PC: ~$ cd ~/Desktop # or navigate to any other directory where you want to create the file
hadoop@bmsccscse-HP-Elite-Tower-800-G9-Desktop-PC: ~/Desktop$ nano sample.txt
hadoop@bmsccscse-HP-Elite-Tower-800-G9-Desktop-PC: ~/Desktop$ cd ~
hadoop@bmsccscse-HP-Elite-Tower-800-G9-Desktop-PC: ~$ jps
8400 SecondaryNameNode
8852 NodeManager
7941 NameNode
8118 DataNode
15593 Jps
10892 org.eclipse.equinox.launcher_1.6.1000.v20250227-1734.jar
8687 ResourceManager
hadoop@bmsccscse-HP-Elite-Tower-800-G9-Desktop-PC: ~$ hadoop fs -mkdir /rgs
mkdir: '/rgs': File exists
hadoop@bmsccscse-HP-Elite-Tower-800-G9-Desktop-PC: ~$ hadoop fs -copyFromLocal ~/Desktop/sample.txt /rgs/test.txt
copyFromLocal: '/rgs/test.txt': File exists
hadoop@bmsccscse-HP-Elite-Tower-800-G9-Desktop-PC: ~$ hadoop fs -ls /rgs
Found 1 items
-rw-r--r-- 1 hadoop supergroup 90 2025-04-15 14:19 /rgs/test.txt
hadoop@bmsccscse-HP-Elite-Tower-800-G9-Desktop-PC: ~$ hadoop fs -cat /rgs/test.txt
ht how are you
how is your job
how is your family
how is your brother
how is your sister
hadoop@bmsccscse-HP-Elite-Tower-800-G9-Desktop-PC: ~$ hadoop jar /home/hadoop/Desktop/MC.jar MyDriver /rgs /output
2025-04-29 15:29:14,897 INFO Impl.MetricsConfig: Loaded properties from hadoop-metrics2.properties
2025-04-29 15:29:14,935 INFO Impl.MetricsSystemImpl: Scheduled Metric snapshot period at 10 second(s).
2025-04-29 15:29:14,935 INFO Impl.MetricsSystemImpl: JobTracker metrics system started
2025-04-29 15:29:14,941 WARN Impl.MetricsSystemImpl: JobTracker metrics system already initialized!
2025-04-29 15:29:14,996 WARN mapreduce.Job$ResourceUploader: Hadoop command-line option parsing not performed. Implement the Tool interface and execute your application with ToolRunner to remedy this.
2025-04-29 15:29:15,059 INFO mapred.FileInputFormat: Total input files to process : 1
2025-04-29 15:29:15,064 INFO mapreduce.JobSubmitter: number of splits:1
2025-04-29 15:29:15,126 INFO mapreduce.JobSubmitter: Submitting tokens for job: job_local537170479_0001
2025-04-29 15:29:15,126 INFO mapreduce.JobSubmitter: Executing with tokens: []
2025-04-29 15:29:15,182 INFO mapreduce.Job: The url to track the job: http://localhost:8080/
2025-04-29 15:29:15,183 INFO mapreduce.Job: Running job: job_local537170479_0001
2025-04-29 15:29:15,183 INFO mapred.LocalJobRunner: OutputCommitter set in config null
2025-04-29 15:29:15,184 INFO mapred.LocalJobRunner: OutputCommitter is org.apache.hadoop.mapred.FileOutputCommitter
2025-04-29 15:29:15,186 INFO output.FileOutputCommitter: File Output Committer Algorithm version is 2
2025-04-29 15:29:15,187 INFO output.FileOutputCommitter: FileOutputCommitter skip cleanup_temporary folders under output directory:false, ignore cleanup failures: false
2025-04-29 15:29:15,221 INFO mapred.LocalJobRunner: Waiting for map tasks
2025-04-29 15:29:15,222 INFO mapred.LocalJobRunner: Starting task: attempt_local537170479_0001_m_000000_0
2025-04-29 15:29:15,233 INFO output.FileOutputCommitter: File Output Committer Algorithm version is 2
2025-04-29 15:29:15,233 INFO output.FileOutputCommitter: FileOutputCommitter skip cleanup_temporary folders under output directory:false, ignore cleanup failures: false
```

```
Activities Terminal Apr 29 15:31
hadoop@bmsccscse-HP-Elite-Tower-800-G9-Desktop-PC: ~$ hadoop jar /home/hadoop/Desktop/MC.jar MyDriver /rgs /output
2025-04-29 15:29:14,897 INFO Impl.MetricsConfig: Loaded properties from hadoop-metrics2.properties
2025-04-29 15:29:14,935 INFO Impl.MetricsSystemImpl: Scheduled Metric snapshot period at 10 second(s).
2025-04-29 15:29:14,935 INFO Impl.MetricsSystemImpl: JobTracker metrics system started
2025-04-29 15:29:14,941 WARN Impl.MetricsSystemImpl: JobTracker metrics system already initialized!
2025-04-29 15:29:14,996 WARN mapreduce.Job$ResourceUploader: Hadoop command-line option parsing not performed. Implement the Tool interface and execute your application with ToolRunner to remedy this.
2025-04-29 15:29:15,059 INFO mapred.FileInputFormat: Total input files to process : 1
2025-04-29 15:29:15,064 INFO mapreduce.JobSubmitter: number of splits:1
2025-04-29 15:29:15,126 INFO mapreduce.JobSubmitter: Submitting tokens for job: job_local537170479_0001
2025-04-29 15:29:15,126 INFO mapreduce.JobSubmitter: Executing with tokens: []
2025-04-29 15:29:15,182 INFO mapreduce.Job: The url to track the job: http://localhost:8080/
2025-04-29 15:29:15,183 INFO mapreduce.Job: Running job: job_local537170479_0001
2025-04-29 15:29:15,183 INFO mapred.LocalJobRunner: OutputCommitter set in config null
2025-04-29 15:29:15,184 INFO mapred.LocalJobRunner: OutputCommitter is org.apache.hadoop.mapred.FileOutputCommitter
2025-04-29 15:29:15,186 INFO output.FileOutputCommitter: File Output Committer Algorithm version is 2
2025-04-29 15:29:15,187 INFO output.FileOutputCommitter: FileOutputCommitter skip cleanup_temporary folders under output directory:false, ignore cleanup failures: false
2025-04-29 15:29:15,221 INFO mapred.LocalJobRunner: Waiting for map tasks
2025-04-29 15:29:15,222 INFO mapred.LocalJobRunner: Starting task: attempt_local537170479_0001_m_000000_0
2025-04-29 15:29:15,233 INFO output.FileOutputCommitter: File Output Committer Algorithm version is 2
2025-04-29 15:29:15,239 INFO mapred.Task: Using ResourceCalculatorProcessFree : [ ]
2025-04-29 15:29:15,242 INFO mapred.MapTask: Processing split: hdfs://localhost:9000/rgs/test.txt:0+90
2025-04-29 15:29:15,252 INFO mapred.MapTask: numReduceTasks: 1
2025-04-29 15:29:15,281 INFO mapred.MapTask: (EQUATOR) 0 kvl 26214396(104857584)
2025-04-29 15:29:15,281 INFO mapred.MapTask: mapreduce.task.io.sort.mb: 100
2025-04-29 15:29:15,281 INFO mapred.MapTask: soft limit at 83886080
2025-04-29 15:29:15,281 INFO mapred.MapTask: bufstart = 0; bufvoid = 104857600
2025-04-29 15:29:15,281 INFO mapred.MapTask: kvstart = 26214396; length = 6553600
2025-04-29 15:29:15,283 INFO mapred.MapTask: Map output collector class = org.apache.hadoop.mapred.MapTask$MapOutputBuffer
2025-04-29 15:29:15,316 INFO mapred.LocalJobRunner:
2025-04-29 15:29:15,316 INFO mapred.MapTask: Starting flush of map output
2025-04-29 15:29:15,316 INFO mapred.MapTask: Spilling map output
2025-04-29 15:29:15,316 INFO mapred.MapTask: bufstart = 0; bufend = 169; bufvoid = 104857600
2025-04-29 15:29:15,316 INFO mapred.MapTask: kvstart = 26214396(104857584); kvend = 26214320(104857280); length = 77/6553600
2025-04-29 15:29:15,319 INFO mapred.MapTask: Finished spill 0
2025-04-29 15:29:15,325 INFO mapred.Task: Task:attempt_local537170479_0001_m_000000_0 is done. And is in the process of committing
2025-04-29 15:29:15,327 INFO mapred.LocalJobRunner: hdfs://localhost:9000/rgs/test.txt:0+90
2025-04-29 15:29:15,327 INFO mapred.Task: Task 'attempt_local537170479_0001_m_000000_0' done.
2025-04-29 15:29:15,329 INFO mapred.Task: Final Counters for attempt_local537170479_0001_m_000000_0: Counters: 23
File System Counters
FILE: Number of bytes read=4156
FILE: Number of bytes written=642471
FILE: Number of read operations=0
FILE: Number of large read operations=0
FILE: Number of write operations=0
HDFS: Number of bytes read=90
HDFS: Number of bytes written=0
HDFS: Number of read operations=5
HDFS: Number of large read operations=0
HDFS: Number of write operations=1
HDFS: Number of bytes read erasure-coded=0
Map-Reduce Framework
Map input records=6
Map output records=20
Map output bytes=169
Map output materialized bytes=215
```

```
Activities Terminal Apr 29 15:32
hadoop@bmsccse-HP-Elite-Tower-800-G9-Desktop-PC: ~
2025-04-29 15:29:15.316 INFO mapred.MapTask: bufstart = 0; bufend = 169; bufvoid = 104857600
2025-04-29 15:29:15.316 INFO mapred.MapTask: kvstart = 26214396(104857584); kvend = 26214320(104857280); length = 77/655360
2025-04-29 15:29:15.319 INFO mapred.MapTask: Finished split 0
2025-04-29 15:29:15.325 INFO mapred.Task: Task:attempt_local537170479_0001_m_000000_0 is done. And is in the process of committing
2025-04-29 15:29:15.327 INFO mapred.LocalJobRunner: hdfs://localhost:9000/rq/test.txt:0-90
2025-04-29 15:29:15.327 INFO mapred.Task: Task 'attempt_local537170479_0001_m_000000_0' done.
2025-04-29 15:29:15.329 INFO mapred.Task: Final Counters for attempt_local537170479_0001_m_000000_0: Counters: 23
File System Counters
  FILE: Number of bytes read=4156
  FILE: Number of bytes written=42471
  FILE: Number of read operations=0
  FILE: Number of large read operations=0
  FILE: Number of write operations=0
  HDFS: Number of bytes read=50
  HDFS: Number of bytes written=0
  HDFS: Number of read operations=5
  HDFS: Number of large read operations=0
  HDFS: Number of write operations=1
  HDFS: Number of bytes read erasure-coded=0
Map-Reduce Framework
  Map input records=6
  Map output records=20
  Map output bytes=169
  Map output materialized bytes=215
  Input split bytes=0
  Combine input records=0
  Spilled Records=20
  Failed Shuffles=0
  Merged Map outputs=0
  GC time elapsed (ms)=0
  Total committed heap usage (bytes)=633339904
File Input Format Counters
  Bytes Read=90
2025-04-29 15:29:15.329 INFO mapred.LocalJobRunner: Finishing task: attempt_local537170479_0001_m_000000_0
2025-04-29 15:29:15.330 INFO mapred.LocalJobRunner: map task executor complete.
2025-04-29 15:29:15.331 INFO mapred.LocalJobRunner: Waiting for reduce tasks
2025-04-29 15:29:15.331 INFO mapred.LocalJobRunner: Starting task: attempt_local537170479_0001_r_000000_0
2025-04-29 15:29:15.335 INFO output.FileOutputCommitter: File Output Committer Algorithm version is 2
2025-04-29 15:29:15.335 INFO output.FileOutputCommitter: FileOutputCommitter skip cleanup_temporary folders under output directory:false, ignore cleanup failures: false
2025-04-29 15:29:15.336 INFO mapred.ReduceTask: Using ShuffleConsumerPlugin: org.apache.hadoop.mapreduce.task.reduce.Shuffle27a9203f
2025-04-29 15:29:15.337 WARN Impl.MetricsSystemImpl: JobTracker metrics system already initialized!
2025-04-29 15:29:15.344 INFO reduce.MergeManagerImpl: MergerManager: memoryLimit=5827985408, maxSingleShuffleLimit=1456996352, mergeThreshold=3846470400, ioSortFactor=10, memToMemMergeOutputsThreshold=10
2025-04-29 15:29:15.346 INFO reduce.LocalFetcher: localFetcher#1 about to shuffle output of map attempt_local537170479_0001_m_000000_0 decomp: 211 len: 215 to MEMORY
2025-04-29 15:29:15.357 INFO reduce.InMemoryMapOutput: Read 211 bytes from map-output for attempt_local537170479_0001_m_000000_0
2025-04-29 15:29:15.358 INFO reduce.MergeManagerImpl: closeInMemoryFile -> map-output of size: 211, inMemoryMapOutputs.size() -> 1, commitMemory -> 0, usedMemory ->211
2025-04-29 15:29:15.359 INFO reduce.EventFetcher: EventFetcher is interrupted.. Returning
2025-04-29 15:29:15.359 INFO mapred.LocalJobRunner: 1 / 1 copied.
2025-04-29 15:29:15.359 INFO reduce.MergeManagerImpl: finalMerge called with 1 in-memory map-outputs and 0 on-disk map-outputs
2025-04-29 15:29:15.362 INFO mapred.Merger: Merging 1 sorted segments
2025-04-29 15:29:15.362 INFO mapred.Merger: Down to the last merge-pass, with 1 segments left of total size: 205 bytes
2025-04-29 15:29:15.362 INFO reduce.MergeManagerImpl: Merged 1 segments, 211 bytes to disk to satisfy reduce memory limit
2025-04-29 15:29:15.362 INFO reduce.MergeManagerImpl: Merging 1 files, 215 bytes from disk
2025-04-29 15:29:15.363 INFO reduce.MergeManagerImpl: Merging 0 segments, 0 bytes from memory into reduce
```

```
Activities Terminal Apr 29 15:33
hadoop@bmsccse-HP-Elite-Tower-800-G9-Desktop-PC: ~
2025-04-29 15:29:15.329 INFO mapred.LocalJobRunner: Finishing task: attempt_local537170479_0001_m_000000_0
2025-04-29 15:29:15.330 INFO mapred.LocalJobRunner: map task executor complete.
2025-04-29 15:29:15.331 INFO mapred.LocalJobRunner: Waiting for reduce tasks
2025-04-29 15:29:15.331 INFO mapred.LocalJobRunner: Starting task: attempt_local537170479_0001_r_000000_0
2025-04-29 15:29:15.335 INFO output.FileOutputCommitter: FileOutputCommitter Algorithm version is 2
2025-04-29 15:29:15.335 INFO mapred.ReduceTask: Using ResourceCalculatorProcessTree: [ ]
2025-04-29 15:29:15.336 INFO mapred.ReduceTask: Using ShuffleConsumerPlugin: org.apache.hadoop.mapreduce.task.reduce.Shuffle27a9203f
2025-04-29 15:29:15.337 WARN Impl.MetricsSystemImpl: JobTracker metrics system already initialized!
2025-04-29 15:29:15.344 INFO reduce.EventFetcher: attempt_local537170479_0001_r_000000_0 Thread started: EventFetcher for fetching Map Completion Events
2025-04-29 15:29:15.346 INFO reduce.LocalFetcher: localFetcher#1 about to shuffle output of map attempt_local537170479_0001_m_000000_0 decomp: 211 len: 215 to MEMORY
2025-04-29 15:29:15.357 INFO reduce.InMemoryMapOutput: Read 211 bytes from map-output for attempt_local537170479_0001_m_000000_0
2025-04-29 15:29:15.358 INFO reduce.MergeManagerImpl: closeInMemoryFile -> map-output of size: 211, inMemoryMapOutputs.size() -> 1, commitMemory -> 0, usedMemory ->211
2025-04-29 15:29:15.359 INFO reduce.EventFetcher: EventFetcher is interrupted.. Returning
2025-04-29 15:29:15.359 INFO mapred.LocalJobRunner: 1 / 1 copied.
2025-04-29 15:29:15.359 INFO reduce.MergeManagerImpl: finalMerge called with 1 in-memory map-outputs and 0 on-disk map-outputs
2025-04-29 15:29:15.362 INFO mapred.Merger: Merging 1 sorted segments
2025-04-29 15:29:15.362 INFO mapred.Merger: Down to the last merge-pass, with 1 segments left of total size: 205 bytes
2025-04-29 15:29:15.362 INFO reduce.MergeManagerImpl: Merged 1 segments, 211 bytes to disk to satisfy reduce memory limit
2025-04-29 15:29:15.362 INFO reduce.MergeManagerImpl: Merging 1 files, 215 bytes from disk
2025-04-29 15:29:15.363 INFO reduce.MergeManagerImpl: Merging 0 segments, 0 bytes from memory into reduce
2025-04-29 15:29:15.363 INFO mapred.LocalJobRunner: 1 / 1 copied.
2025-04-29 15:29:15.415 INFO mapred.Task: Task:attempt_local537170479_0001_r_000000_0 is done. And is in the process of committing
2025-04-29 15:29:15.416 INFO mapred.LocalJobRunner: 1 / 1 copied.
2025-04-29 15:29:15.416 INFO mapred.Task: Task attempt_local537170479_0001_r_000000_0 is allowed to commit now
2025-04-29 15:29:15.423 INFO output.FileOutputCommitter: Saved output of task 'attempt_local537170479_0001_r_000000_0' to hdfs://localhost:9000/output
2025-04-29 15:29:15.424 INFO mapred.LocalJobRunner: reduce > reduce
2025-04-29 15:29:15.424 INFO mapred.Task: Task 'attempt_local537170479_0001_r_000000_0' done.
2025-04-29 15:29:15.424 INFO mapred.Task: Final Counters for attempt_local537170479_0001_r_000000_0: Counters: 30
File System Counters
  FILE: Number of bytes read=4618
  FILE: Number of bytes written=42686
  FILE: Number of read operations=0
  FILE: Number of large read operations=0
  FILE: Number of write operations=0
  HDFS: Number of bytes read=50
  HDFS: Number of bytes written=69
  HDFS: Number of read operations=10
  HDFS: Number of large read operations=0
  HDFS: Number of write operations=3
  HDFS: Number of bytes read erasure-coded=0
Map-Reduce Framework
  Combine input records=0
  Combine output records=0
  Reduce input groups=10
  Reduce shuffle bytes=215
  Reduce input records=20
  Reduce output records=10
  Spilled Records=20
  Shuffled Maps =1
  Failed Shuffles=0
  Merged Map outputs=1
```

```
Activities Terminal Apr 29 15:33
hadoop@bmsccse-HP-Elite-Tower-800-G9-Desktop-PC: ~

File System Counters
  FILE: Number of bytes read=4618
  FILE: Number of bytes written=842686
  FILE: Number of read operations=0
  FILE: Number of large read operations=0
  FILE: Number of write operations=0
  HDFS: Number of bytes read=59
  HDFS: Number of bytes written=69
  HDFS: Number of read operations=10
  HDFS: Number of large read operations=0
  HDFS: Number of write operations=1
  HDFS: Number of bytes read erasure-coded=0
Map-Reduce Framework
  Combine input records=0
  Combine output records=0
  Reduce input groups=10
  Reduce shuffle bytes=215
  Reduce input records=20
  Reduce output records=10
  Spilled Records=20
  Shuffled Maps =1
  Failed Shuffles=0
  Merged Map outputs=1
  GC time elapsed (ms)=0
  Total committed heap usage (bytes)=633339904
Shuffle Errors
  BAD_ID=0
  CONNECTION=0
  IO_ERROR=0
  WRONG_LENGTH=0
  WRONG_MAP=0
  WRONG_REDUCE=0
File Output Format Counters
  Bytes Written=69
2025-04-29 15:29:15,424 INFO mapred.LocalJobRunner: Finishing task: attempt_local537170479_0001_r_000000_0
2025-04-29 15:29:15,424 INFO mapred.LocalJobRunner: reduce task executor complete.
2025-04-29 15:29:16,185 INFO mapreduce.Job: Job job_local537170479_0001 running in uber mode : false
2025-04-29 15:29:16,186 INFO mapreduce.Job: map 100% reduce 100%
2025-04-29 15:29:16,186 INFO mapreduce.Job: Job job_local537170479_0001 completed successfully
2025-04-29 15:29:16,190 INFO mapreduce.Job: Counters: 36
File System Counters
  FILE: Number of bytes read=8774
  FILE: Number of bytes written=1285157
  FILE: Number of read operations=0
  FILE: Number of large read operations=0
  FILE: Number of write operations=0
  HDFS: Number of bytes read=180
  HDFS: Number of bytes written=69
  HDFS: Number of read operations=15
  HDFS: Number of large read operations=0
  HDFS: Number of write operations=4
  HDFS: Number of bytes read erasure-coded=0
Map-Reduce Framework
  Map input records=6
  Map output records=20
  Map output bytes=169
```

```
Activities Terminal Apr 29 15:34
hadoop@bmsccse-HP-Elite-Tower-800-G9-Desktop-PC: ~

2025-04-29 15:29:15,424 INFO mapred.LocalJobRunner: reduce task executor complete.
2025-04-29 15:29:16,185 INFO mapreduce.Job: Job job_local537170479_0001 running in uber mode : false
2025-04-29 15:29:16,186 INFO mapreduce.Job: map 100% reduce 100%
2025-04-29 15:29:16,186 INFO mapreduce.Job: Job job_local537170479_0001 completed successfully
2025-04-29 15:29:16,190 INFO mapreduce.Job: Counters: 36
File System Counters
  FILE: Number of bytes read=8774
  FILE: Number of bytes written=1285157
  FILE: Number of read operations=0
  FILE: Number of large read operations=0
  FILE: Number of write operations=0
  HDFS: Number of bytes read=180
  HDFS: Number of bytes written=69
  HDFS: Number of read operations=15
  HDFS: Number of large read operations=0
  HDFS: Number of write operations=4
  HDFS: Number of bytes read erasure-coded=0
Map-Reduce Framework
  Map input records=6
  Map output records=20
  Map output bytes=169
  Map output materialized bytes=215
  Input split bytes=86
  Combine input records=0
  Combine output records=0
  Reduce input groups=10
  Reduce shuffle bytes=215
  Reduce input records=20
  Reduce output records=10
  Spilled Records=40
  Shuffled Maps =1
  Failed Shuffles=0
  Merged Map outputs=1
  GC time elapsed (ms)=0
  Total committed heap usage (bytes)=1266679808
Shuffle Errors
  BAD_ID=0
  CONNECTION=0
  IO_ERROR=0
  WRONG_LENGTH=0
  WRONG_MAP=0
  WRONG_REDUCE=0
File Input Format Counters
  Bytes Read=0
File Output Format Counters
  Bytes Written=69
0
hadoop@bmsccse-HP-Elite-Tower-800-G9-Desktop-PC: ~$ hadoop fs -ls /output
Found 2 items
-rw-r--r-- 1 hadoop supergroup 0 2025-04-29 15:29 /output/_SUCCESS
-rw-r--r-- 1 hadoop supergroup 69 2025-04-29 15:29 /output/part-00000
hadoop@bmsccse-HP-Elite-Tower-800-G9-Desktop-PC: ~$ hadoop fs -cat /output/part-00000
are 1
brother 1
family 1
hi 1
```

```
Activities Terminal Apr 29 15:34
hadoop@bmsccse-HP-Elite-Tower-800-G9-Desktop-PC: ~

FILE: Number of read operations=0
FILE: Number of large read operations=0
FILE: Number of write operations=0
HDFS: Number of bytes read=180
HDFS: Number of bytes written=69
HDFS: Number of read operations=15
HDFS: Number of large read operations=0
HDFS: Number of write operations=4
HDFS: Number of bytes read erasure-coded=0

Map-Reduce Framework
  Map input records=6
  Map output records=20
  Map output bytes=169
  Map output materialized bytes=215
  Input split bytes=40
  Combine input records=0
  Combine output records=0
  Reduce input groups=10
  Reduce shuffle bytes=215
  Reduce input records=20
  Reduce output records=10
  Spilled Records=40
  Shuffled Maps =1
  Failed Shuffles=0
  Merged Map outputs=1
  GC time elapsed (ms)=0
  Total committed heap usage (bytes)=1266679808

Shuffle Errors
  BAD_ID=0
  CONNECTION=0
  IO_ERROR=0
  WRONG_LENGTH=0
  WRONG_MAP=0
  WRONG_REDUCE=0

File Input Format Counters
  Bytes Read=90
File Output Format Counters
  Bytes Written=69

0
hadoop@bmsccse-HP-Elite-Tower-800-G9-Desktop-PC: ~$ hadoop fs -ls /output
Found 2 items
-rw-r--r-- 1 hadoop supergroup 0 2025-04-29 15:29 /output/_SUCCESS
-rw-r--r-- 1 hadoop supergroup 69 2025-04-29 15:29 /output/part-00000
hadoop@bmsccse-HP-Elite-Tower-800-G9-Desktop-PC: ~$ hadoop fs -cat /output/part-00000
are 1
brother 1
family 1
ht 1
how 5
is 4
job 1
sister 1
you 1
your 4
hadoop@bmsccse-HP-Elite-Tower-800-G9-Desktop-PC: ~$
```

```
Activities Eclipse Apr 29 15:34
eclipse-workspace - WordCount/src/MyDriver.java - Eclipse IDE

File Edit Source Refactor Navigate Search Project Run Window Help

Package Explorer x
  JRE System Library [JVMSE-1.8]
  src
    MyDriver.java
    MyMapper.java
    MyReducer.java
  Referenced Libraries
  .settings
  .classpath
  .project

Task List x
Find
All Activate...

Outline x
MyDriver

Problems x Javadoc Declaration Install Java 24 Support x
Java 24 has been released.
Would you like to install Eclipse Java Development Tools support for Java 24 into your IDE?
Install Java 24 Support Remind Me Later Don't Ask Again
Would you like to download a Java 24 JDK?
Download Temurin JDK 24 Download OpenJDK 24

1: import java.io.IOException;
2: import org.apache.hadoop.conf.Configuration;
3: import org.apache.hadoop.fs.Path;
4: import org.apache.hadoop.io.IntWritable;
5: import org.apache.hadoop.io.Text;
6: import org.apache.hadoop.mapred.FileInputFormat;
7: import org.apache.hadoop.mapred.FileOutputFormat;
8: import org.apache.hadoop.mapred.JobClient;
9: import org.apache.hadoop.mapred.JobConf;
10: import org.apache.hadoop.util.Tool;
11: import org.apache.hadoop.util.ToolRunner;
12:
13: public class MyDriver extends Configuration implements Tool {
14:     public int run(String[] args) throws IOException {
15:         if (args.length < 2) {
16:             System.out.println("Please provide input and output paths");
17:             return -1;
18:         }
19:         JobConf conf = new JobConf(MyDriver.class);
20:         FileInputFormat.setInputPaths(conf, new Path(args[0]));
21:         FileOutputFormat.setOutputPath(conf, new Path(args[1]));
22:         conf.setMapperClass(MyMapper.class);
23:         conf.setReducerClass(MyReducer.class);
24:         conf.setMapOutputKeyClass(Text.class);
25:         conf.setMapOutputValueClass(IntWritable.class);
26:         conf.setOutputKeyClass(Text.class);
27:         conf.setOutputValueClass(IntWritable.class);
28:         JobClient.runJob(conf);
29:         return 0;
30:     }
31:
32:     public static void main(String[] args) throws Exception {
33:         int exitCode = ToolRunner.run(new MyDriver(), args);
34:         System.out.println(exitCode);
35:     }
36: }
37:
```







